



# DECISION TREE

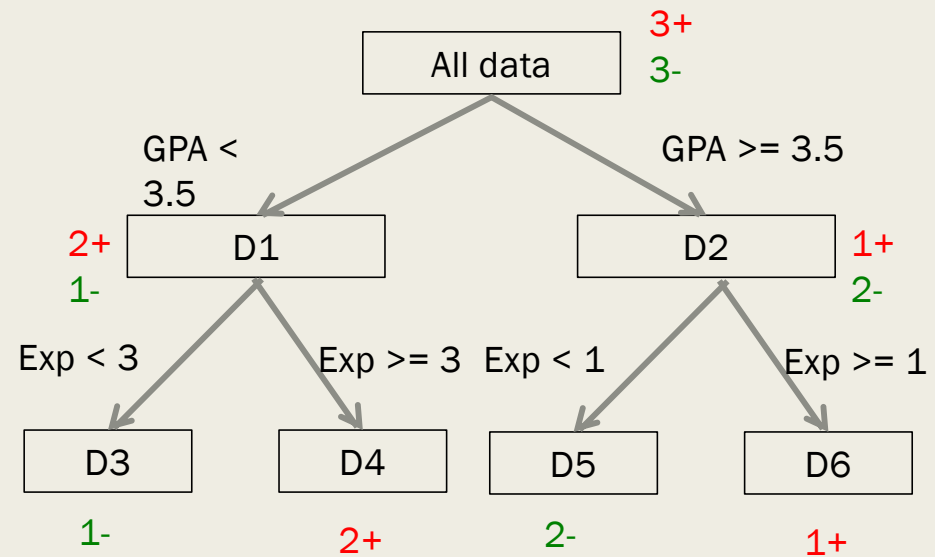
Anurag Nagar

# What is a DT?

# Decision Tree

## ■ 2-D case

GPA	Years Exp	Internship
4.0	0	0
3.0	1	0
3.4	4	1
3.6	0	0
3.8	4	1
2.5	3	1



Of course, you could start with Exp as the first sorting or splitting criteria and get a different tree.

# Decision Tree - Representation

- In this class, DT is a way to represent concepts & hypothesis about a target concept
- Can be written in form of **rules**

IF exp > 3 and GPA > 3.5 THEN internship = 1

- **Leaf nodes** decide values of output variable
- **Internal nodes & edges** represent splitting (sorting) criteria.
- We will consider Boolean attributes and output.
- Given instances with n Boolean attributes i.e. each  $X^i$  is of the form:


$X^i = (x_1, x_2, \dots, x_n)$  e.g.  $X^i = (0, 1, \dots, 0)$

How will you represent one hypothesis

=> A binary tree


# Properties of Binary Trees

- A complete binary tree of height  $h$  has  $2^{h+1} - 1$  nodes
- Number of nodes at depth  $d$  is  $2^d$



# Properties of Binary Trees

- A complete binary tree of height  $h$  has  $2^{h+1} - 1$  nodes
- Number of nodes at depth  $d$  is  $2^d$



$$\# \text{ nodes} = 2^0 + 2^1 + \dots + 2^h = 2^{h+1} - 1$$

(c) Anurag Nagar - Do Not Distribute


# Learning a DT

# Use of DT in learning

- Use the training examples and their labels to construct decision tree
- For example,  $(X^1, y^1)$  could be  $((0, 0, 0), 1)$
- You can use DT to model knowledge from training data.

## A Decision tree for

F: <Outlook, Humidity, Wind, Temp>  $\rightarrow$  PlayTennis?



Each internal node: test one discrete-valued attribute  $X_i$


Each branch from a node: selects one value for  $X_i$

Each leaf node: predict  $Y$  (or  $P(Y|X \in \text{leaf})$ )



# Using DT to represent hypotheses

- **Internal nodes** test the value of particular features  $x_j$  and branch according to the results of the test.
- **Leaf nodes** specify the class  $h(\mathbf{x})$ .




Suppose the features are **Outlook** ( $x_1$ ), **Temperature** ( $x_2$ ), **Humidity** ( $x_3$ ), and **Wind** ( $x_4$ ). Then the feature vector  $\mathbf{x} = (\text{Sunny}, \text{Hot}, \text{High}, \text{Strong})$  will be classified as **No**. The **Temperature** feature is irrelevant.


(c) Anurag Nagar - Do Not Distribute

# Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the  $K$  classes.





(c) Anurag Nagar - Do Not Distribute





# Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the  $K$  classes.





# Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the  $K$  classes.





# Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the  $K$  classes.



# Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the  $K$  classes.



# Hypothesis Space of DT

As the number of nodes (or depth) of tree increases, the hypothesis space grows

- **depth 1** (“decision stump”) can represent any boolean function of one feature.
- **depth 2** Any boolean function of two features; some boolean functions involving three features (e.g.,  $(x_1 \wedge x_2) \vee (\neg x_1 \wedge \neg x_3)$ )
- **etc.**

# Finding the best split (also called sort)




# DT – How to find best sorting?

- Which attribute should I sort (split) on first?
  - It DOES make a difference.
- Informally, we want that split that gives maximum purity at each node i.e. split such that all instances are of a single class (or close to it).

Which test is more informative?


**Split over whether  
Balance exceeds 50K**



Less or equal 50K

Over 50K

**Split over whether  
applicant is employed**



Unemployed

Employed

# Entropy

Entropy is a measure of Information Content (IC).

$$H(X) = \sum -p_i \log_2 p_i$$

where  $p_i$  is the probability of the  $i^{\text{th}}$  class.

If you think deeply, it is the **expected value** of  $-\log_2 p_i$  or  $\log(1/p_i)$ .


This quantity is also known as information of an attribute.

Also,  $H(x)$  can be thought of as the number of bits needed to encode a dataset

- Entropy =  $\sum_i -p_i \log_2 p_i$

$p_i$  is the probability of class  $i$

Compute it as the proportion of class  $i$  in the set.



16/30 are green circles; 14/30 are pink crosses

$\log_2(16/30) = -.9$ ;  $\log_2(14/30) = -1.1$

Entropy =  $-(16/30)(-.9) - (14/30)(-1.1) = .99$

- Entropy comes from information theory. The higher the entropy the more the information content.

What does that mean for learning from examples?

- What is the entropy of a group in which all examples belong to the same class?

–  $\text{entropy} = -1 \log_2 1 = 0$


not a good training set for learning

- What is the entropy of a group with 50% in either class?


–  $\text{entropy} = -0.5 \log_2 0.5 - 0.5 \log_2 0.5 = 1$

good training set for learning


**Minimum  
impurity**



**Maximum  
impurity**



## Entropy of a binary random variable



- Entropy is maximum at  $p=0.5$
- Entropy is zero and  $p=0$  or  $p=1$ .

(c) Anurag Nagar - Do Not Distribute

# Information Gain

- We want to determine **which attribute** in a given set of training feature vectors is **most useful** for discriminating between the classes to be learned.
- **Information gain** tells us how important a given attribute of the feature vectors is.
- We will use it to decide the ordering of attributes in the nodes of a decision tree.

# Information Gain

- Suppose you just know the class labels initially.
- Then you know one of the attributes.

=> Does it really help you?

=> Do you get any information gain or reduction in entropy ?

=> Do you get any increase in purity of the classes by knowing an attribute?

Mean the same thing

# Information Gain


- Information Gain (IG) is defined as:

IG = Parent Entropy – Average Entropy of children

- Average entropy of children:  
Probability weighted entropy of all the child nodes.

$$= \frac{C_1}{P} \times H_1 + \frac{C_2}{P} \times H_2 + \dots + \frac{C_n}{P} \times H_n$$

- Information Gain (IG) =  
$$= H - \left( \frac{C_1}{P} \times H_1 + \frac{C_2}{P} \times H_2 + \dots + \frac{C_n}{P} \times H_n \right)$$



Parent has P data instances and entropy of H

After splitting into n child nodes:

Child<sub>1</sub> has C<sub>1</sub> data instances and entropy of H<sub>1</sub>

Child<sub>2</sub> has C<sub>2</sub> data instances and entropy of H<sub>2</sub>

....

Child<sub>n</sub> has C<sub>n</sub> data instances and entropy of H<sub>n</sub>

# Example of IG

Predicting credit risk

<2 years at current job?	missed payments?	defaulted?
N	N	N
Y	N	Y
N	N	N
N	N	N
N	Y	Y
Y	N	N
N	Y	N
N	Y	Y
Y	N	N
Y	N	N

(c) Anurag Nagar - Do Not Distribute

Class attribute is **defaulted?**

Independent Attributes - the **first two**

How many bits does it take to specify the attribute of 'defaulted?'

- $P(\text{defaulted} = Y) = 3/10$
- $P(\text{defaulted} = N) = 7/10$

$$\begin{aligned} H(Y) &= - \sum_{i=Y,N} P(Y = y_i) \log_2 P(Y = y_i) \\ &= -0.3 \log_2 0.3 - 0.7 \log_2 0.7 \\ &= 0.8813 \end{aligned}$$

Can you do better than this by knowing another attribute?



## Back to the credit risk example

$$\begin{aligned}
 H(Y|X) &\equiv - \sum_x P(x) \sum_y P(y|x) \log_2 P(y|x) \\
 &= - \sum_x P(x) \sum_y P(Y = y|X = x) \log_2 P(Y = y|X = x) \\
 &= - \sum_x P(x) H(Y|X = x)
 \end{aligned}$$

$$H(\text{defaulted} | < 2\text{years} = N) = -\frac{4}{4+2} \log_2 \frac{4}{4+2} - \frac{2}{6} \log_2 \frac{2}{6} = 0.9183$$

$$H(\text{defaulted} | < 2\text{years} = Y) = -\frac{3}{4} \log_2 \frac{3}{4} - \frac{1}{4} \log_2 \frac{1}{4} = 0.8133$$

$$H(\text{defaulted} | < 2\text{years}) = \frac{6}{10} 0.9183 + \frac{4}{10} 0.8133 = 0.8763$$

$$H(\text{defaulted} | \text{missed} = N) = -\frac{6}{7} \log_2 \frac{6}{7} - \frac{1}{7} \log_2 \frac{1}{7} = 0.5917$$

$$H(\text{defaulted} | \text{missed} = Y) = -\frac{1}{3} \log_2 \frac{1}{3} - \frac{2}{3} \log_2 \frac{2}{3} = 0.9183$$

$$H(\text{defaulted} | \text{missed}) = \frac{7}{10} 0.5917 + \frac{3}{10} 0.9183 = 0.6897$$

(c) Anurag Nagar - Do Not Distribute

Predicting credit risk

<2 yrs	missed	def?
N	N	N
Y	N	Y
N	N	N
N	N	N
N	Y	Y
Y	N	N
N	Y	N
N	Y	Y
Y	N	N
Y	N	N

Average entropy given  
value of "<2years"  
attribute

Average entropy given  
value of "missed"  
attribute

- We now have the entropy - the minimal number of bits required to specify the target attribute:

$$H(Y) = \sum_y P(y) \log_2 P(y)$$

- The conditional entropy - the remaining entropy of Y knowing X

$$H(Y|X) = - \sum_x P(x) \sum_y P(y|x) \log_2 P(y|x)$$


Attribute X has different values represented by x  
y represents the possible class labels

- So we can now define the reduction of the entropy after learning Y.
- This is known as the *mutual information* between Y and X

$$I(Y; X) = H(Y) - H(Y|X)$$

Original	Average entropy after
Entropy	sorting on attribute X

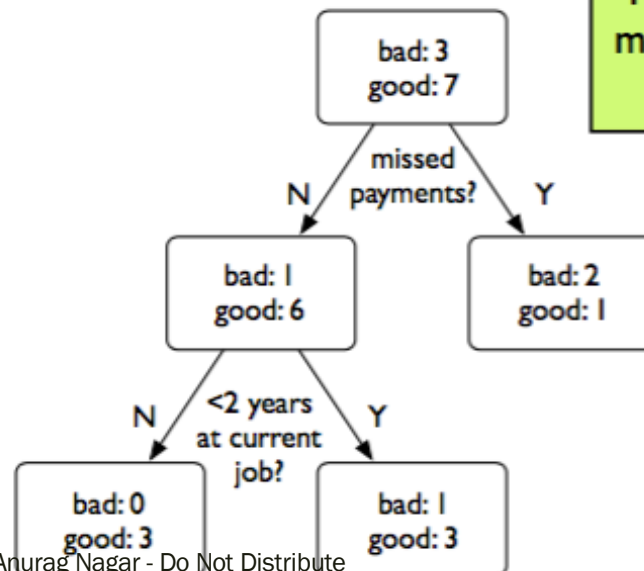
# So... which attribute should I split on?



$$\begin{aligned} H(\text{defaulted}) &- H(\text{defaulted} | < 2 \text{ years}) \\ 0.8813 &- 0.8763 = 0.0050 \end{aligned}$$

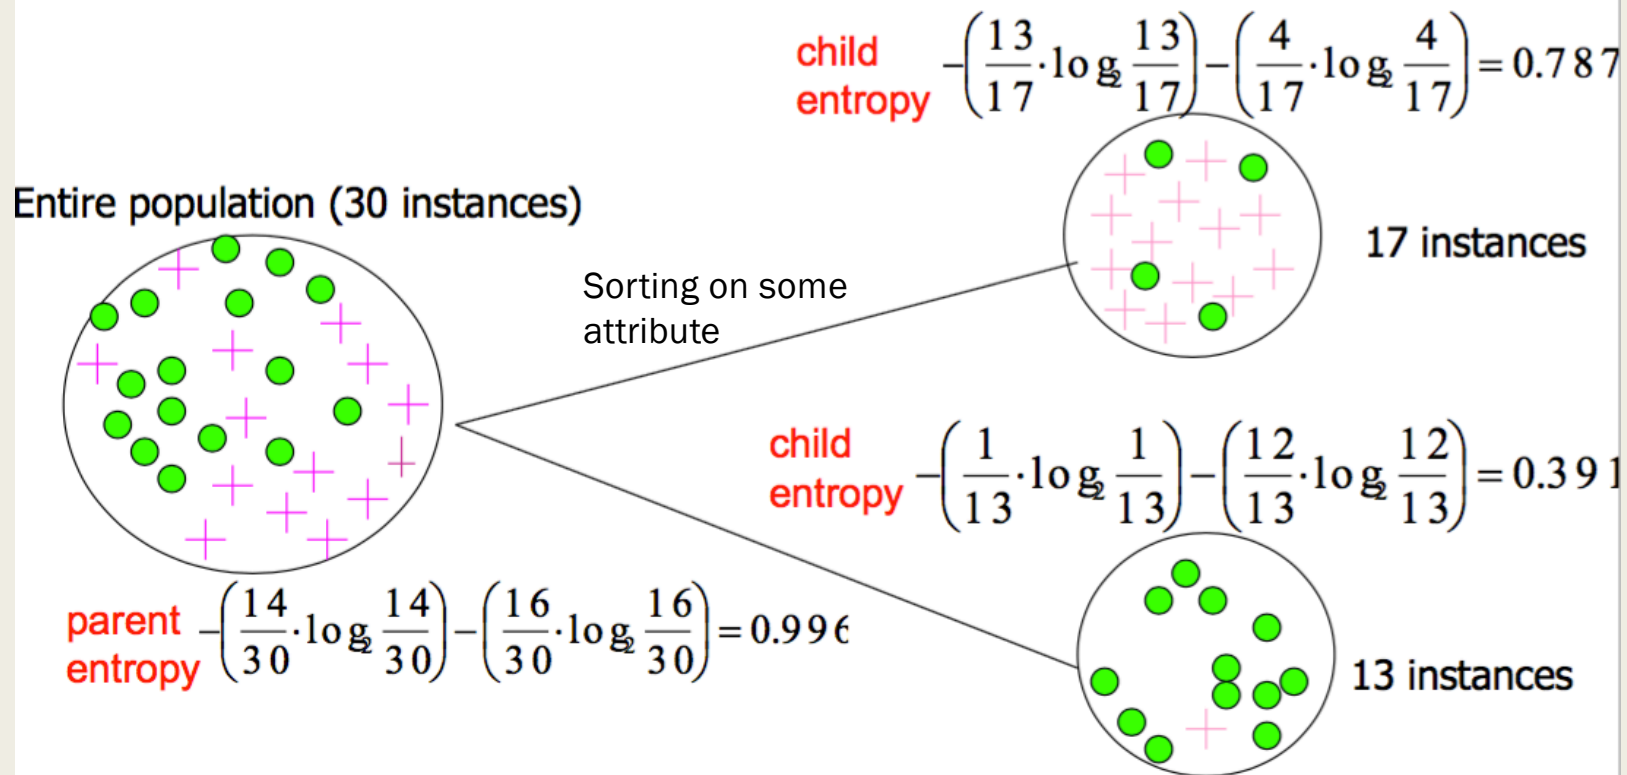
$$\begin{aligned} H(\text{defaulted}) &- H(\text{defaulted} | \text{missed}) \\ 0.8813 &- 0.6897 = 0.1916 \end{aligned}$$

Missed payments are the most informative attribute about defaulting.



## Calculating Information Gain

$$\text{Information Gain} = \text{entropy}(\text{parent}) - [\text{average entropy}(\text{children})]$$



$$(\text{Weighted}) \text{ Average Entropy of Children} = \left(\frac{17}{30} \cdot 0.787\right) + \left(\frac{13}{30} \cdot 0.391\right) = 0.615$$

$$\text{Information Gain} = 0.996 - 0.615 = 0.38 \quad \text{for this split}$$

# Calculating IG

$$E(S) = E(29, 35)$$

$$E(X1) = E(21, 5)$$



$$E(X2) = E(30, 8)$$

$$IG = E(S) - [26/64 * E(X1) + 38/64 * E(X2)]$$

## Information Gain

$Gain(S, A) =$  expected reduction in entropy due to sorting on  $A$


$$Gain(S, A) \equiv Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$



# Use IG to build a tree

- Start at the root of the tree (N1) and choose the best attribute according to IG. Let's say it is  $a_1$
- Create its children according to possible values of  $a_1$ . If  $a_1$  is Boolean, it will have 2 values 0 and 1.
- These two nodes will be the leaf nodes for now -> since we haven't constructed whole tree yet.
- Also, these two nodes will have only less attribute to choose from. Their available choices will be  $\{a_2, \dots, a_n\}$
- Similarly create the rest of the tree, choosing the best attribute at each step.
- Terminating Condition:
  - The node is pure in terms of class label.
  - There are no more attributes remaining.

Available attributes =  $\{a_1, a_2, \dots, a_n\}$



# ID3 algorithm

## Top-Down Induction of Decision Trees

[ID3, C4.5, Quinlan]

*node* = Root

Main loop:

1.  $A \leftarrow$  the “best” decision attribute for next *node*

one that gives  
the best IG

2. Assign  $A$  as decision attribute for *node*

3. For each value of  $A$ , create new descendant of  
*node*

4. Sort training examples to leaf nodes

Note: These newly created  
nodes will be leaf nodes for now

5. If training examples perfectly classified, Then  
STOP, Else iterate over new leaf nodes

# ID3 algorithm

## ■ Simpler version

ID3 (node, {attributes})

1. Let A = best attribute among {attributes} according to IG
2. Sort according to A and create child nodes (which will be leaf nodes for now)
3. For all the child nodes
  - If it is pure or {attributes} is empty,  
STOP
  - else  
call ID3(childNode, {attributes} - A) /\* Recursive Call \*/




# Worked out example

- Please see the handout/notes for a worked out example using ID3
- Remember:
  - For each node, you have to find the best attribute
  - You can only use an attribute once along a path. So, a node needs to inherit a list of attributes from its parent class
    - > You have to program this. 😊
  - At the leaf node, find the majority class (by count). Use that for the prediction rule.

# Does it really matter which attribute comes first?

- ID3 helps us in selecting the shortest i.e. most compact tree



(c) Anurag Nagar - Do Not Distribute

- ID3 performs heuristic search through space of decision trees
- It stops at smallest acceptable tree. Why?

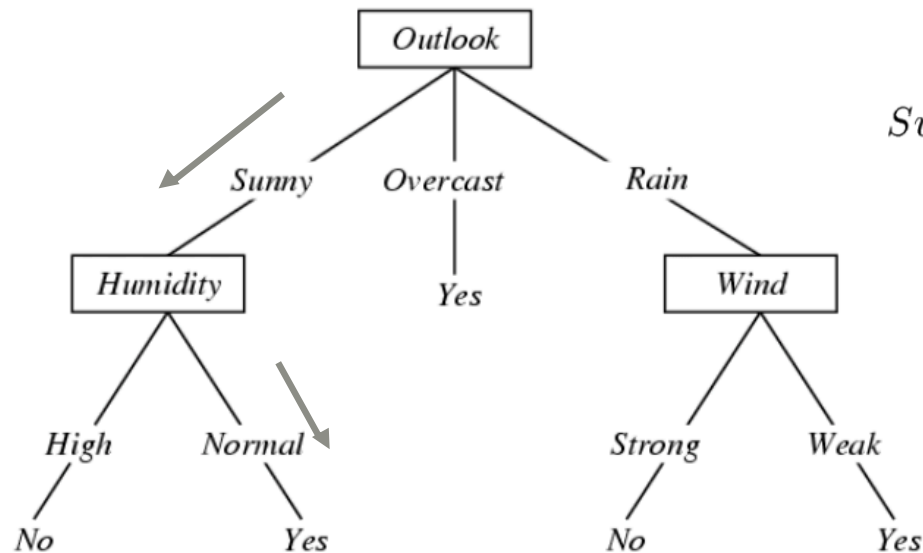
Occam's razor: prefer the simplest hypothesis that fits the data

ID3 is a greedy algorithm  
Top-down induction of trees

# Problem with DT

# Problem with DT?

- Over-zealous learner -> learns all features
- What if there is noise?
- DT will try to change everything??
- Consider tree below:



What would happen if you get **noisy** data point  
*Sunny, Hot, Normal, Strong, PlayTennis = No*

# Overfitting

- You train on the training dataset

The data that the learner trains with.

=> It is possible to design a DT that gives 100% accuracy on training data. Think how??

e.g. each instance gets its own leaf node

- But is that a good thing?

=> NO! Because you are in fact memorizing (rote learning) the training data

=> No room for generalization, it's a case of Overfitting

- So, you have to find a balance between underfitting (learning very little) and overfitting.

- Notation:

Training error of hypothesis  $h = e_{\text{train}}(h)$

True error (on unseen data) of hypothesis  $h = e_{\text{true}}(h)$

# Overfitting

Consider a hypothesis  $h$  and its

- Error rate over training data:  $error_{train}(h)$
- True error rate over all data:  $error_{true}(h)$


We say  $h$  overfits the training data if

$$error_{true}(h) > error_{train}(h)$$

Amount of overfitting =

$$error_{true}(h) - error_{train}(h)$$

# Overfitting in Decision Tree Learning



(c) Anurag Nagar - Do Not Distribute

# How to avoid overfitting?

1. Stop growing when splits are not statistically significant (TOUGHER PROBLEM)

OR

2. Grow full tree then post-prune i.e. remove nodes and see if true error decreases (EASIER PROBLEM)



# How to avoid overfitting?

- Keep another dataset -> validation dataset
- Build model on training, test accuracy on validation
- Learn model from training dataset.
- Randomly remove nodes and see if validation accuracy improves

# Reduced-Error Pruning

---

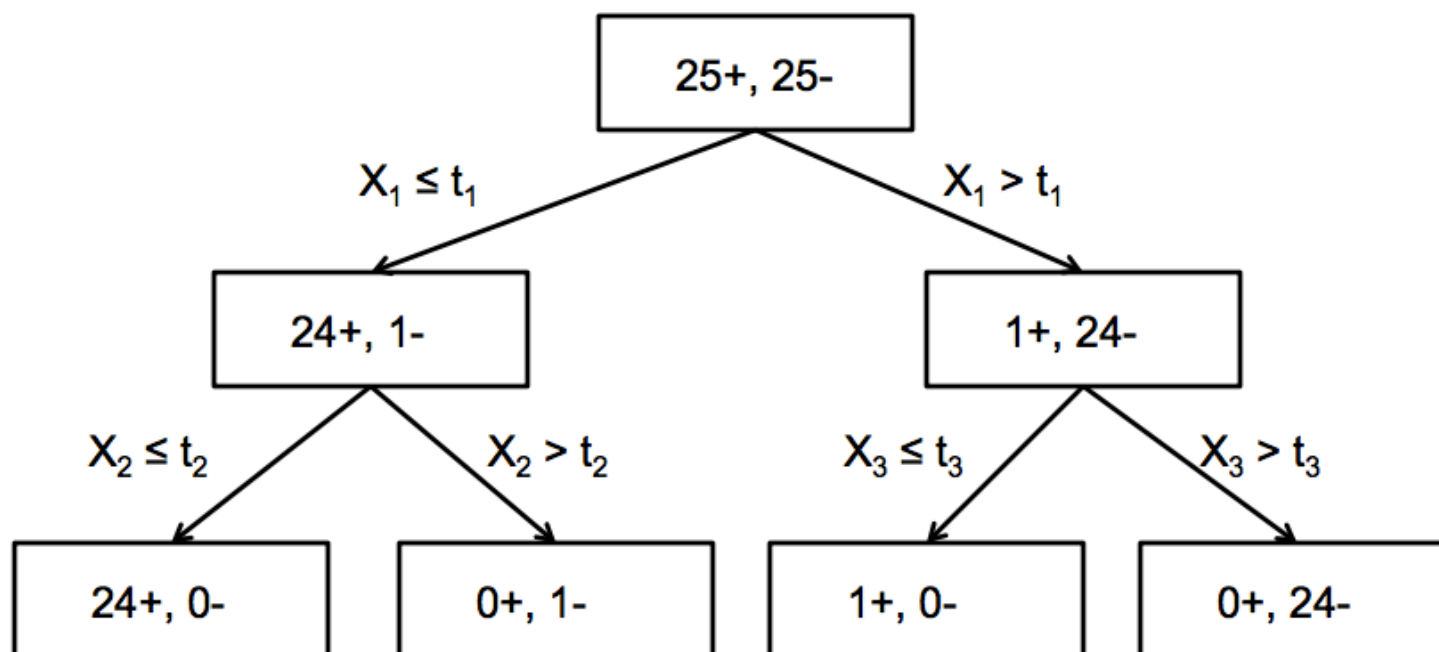
Split data into *training* and *validation* set

Create tree that classifies *training* set correctly

Do until further pruning is harmful:

1. Evaluate impact on *validation* set of pruning each possible node (plus those below it)
2. Greedily remove the one that most improves *validation* set accuracy

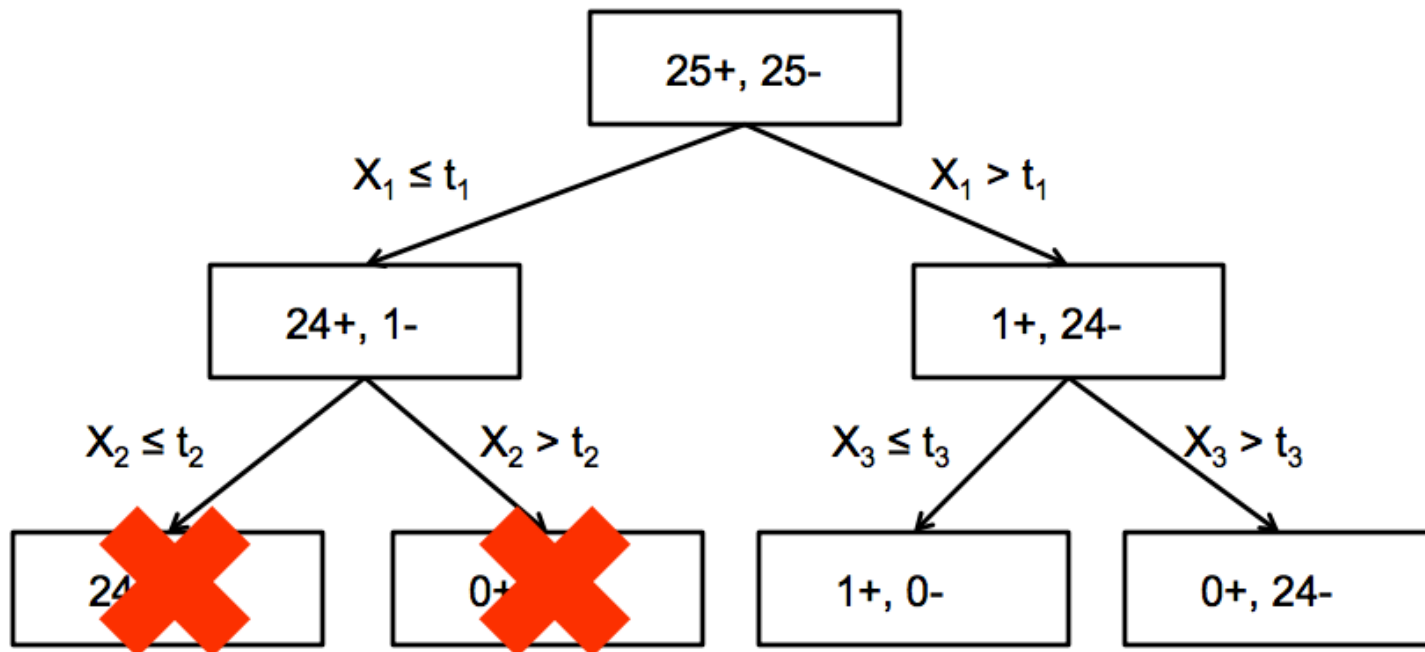
# Post-pruning



Validation Set Accuracy: 80%

(c) Anurag Nagar. Do Not Distribute

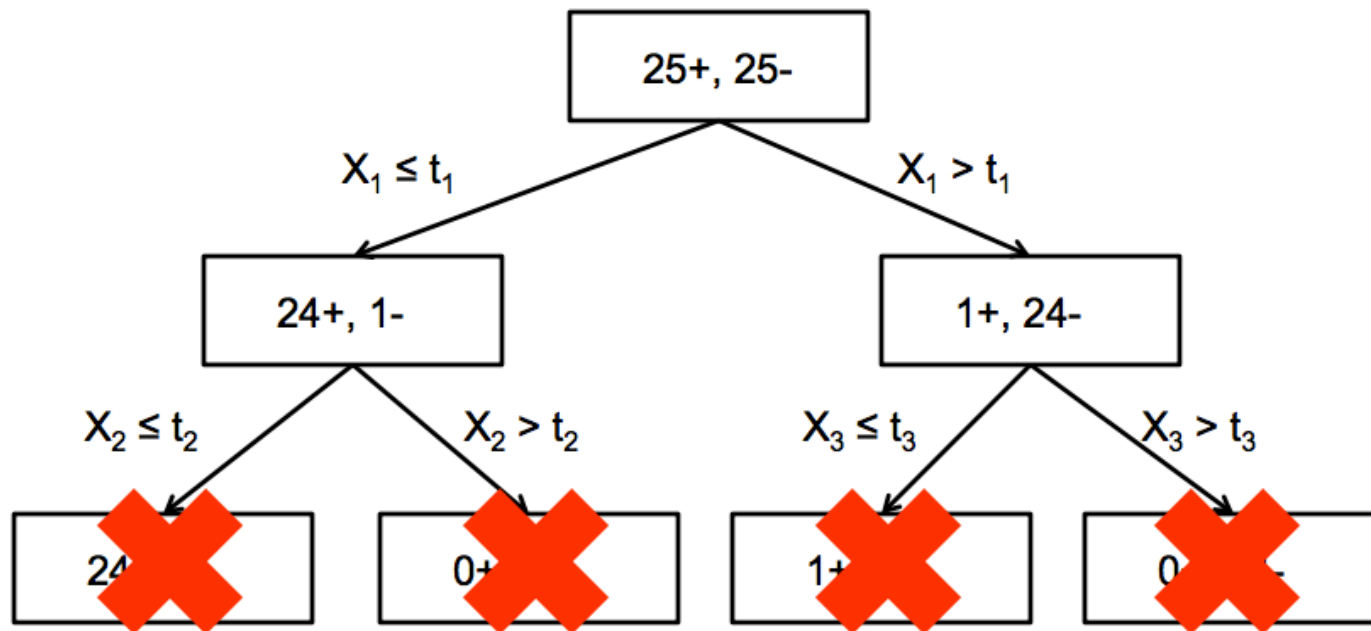
# Post-pruning



**Validation Set Accuracy: 85%**

(c) Anurag Nagar - Do Not Distribute

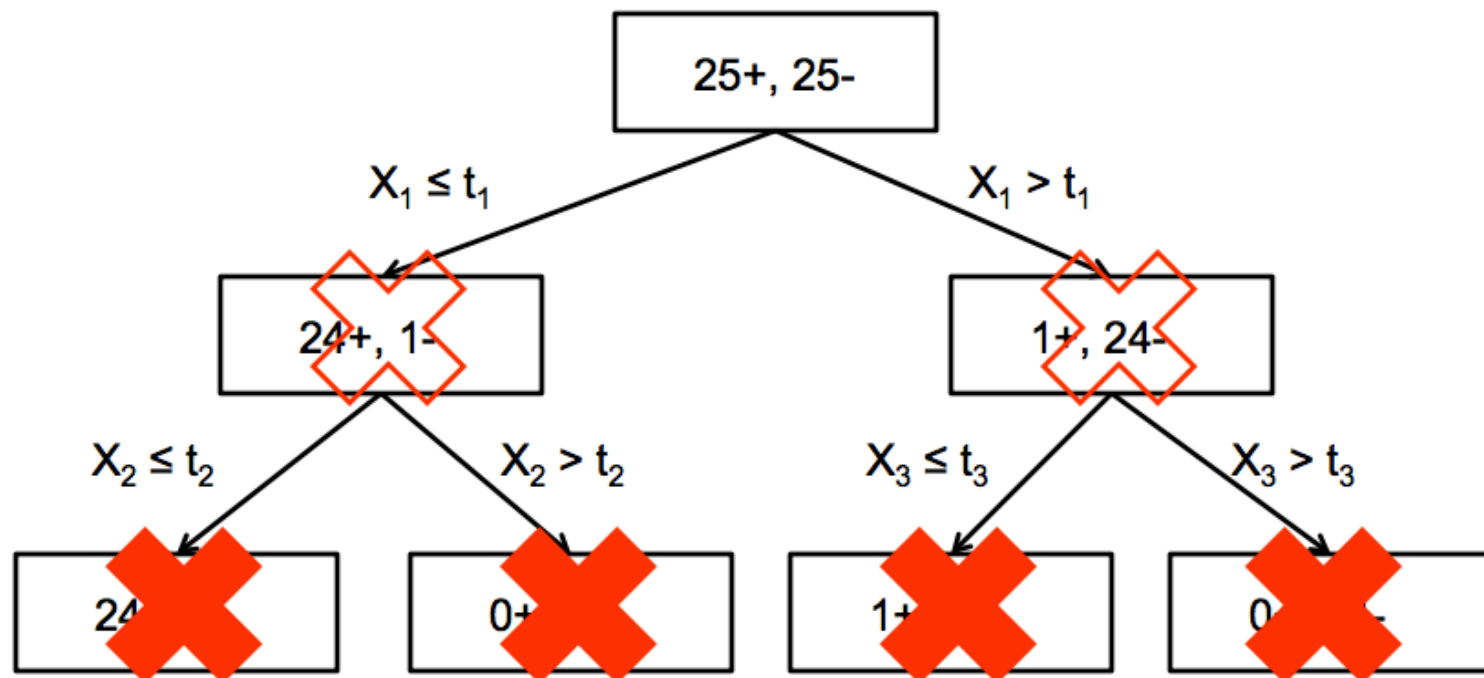
# Post-pruning



Validation Set Accuracy: 90%

(c) Anurag Nagar - Do Not Distribute


# Post-pruning



Validation Set Accuracy: 50%

(c) Anurag Nagar - Do Not Distribute

# Post-pruning



Final Decision Tree

# What have we learnt?

- Idea of DT
- Number of instances (leaf nodes) and hypotheses
- How to choose best sorting attribute for each node
- How to induce top-down tree using ID3
- What is overfitting
- Avoiding overfitting