

Logistic Regression

Generative vs Discriminative Classifiers

- Classification task is to find a function $h: X \rightarrow Y$ from training sample.
- h is an approximation to the real classification function f
- Generative Classifiers: approximate the function by $P(Y|X)$ based on:
$$P(Y|X) \propto P(X|Y) P(Y)$$
 - X is a vector of features
 - First term $P(X|Y)$ is **likelihood** and
Second term $P(Y)$ is **prior**
 - How do we get their values \rightarrow from training data.
 - For likelihood term, we need **joint probability distribution**

Generative vs Discriminative Classifiers

- What is **joint probability distribution**?

Suppose you have 3 Boolean attributes

$\mathbf{X} = (X1, X2, X3)$

and two classes $Y = 0$ and $Y = 1$

- For each class, you need the complete table filled out.

How many entries do you need? $2^3 - 1 = 7$

X1	X2	X3	P (Y=i)
0	0	0	?
0	0	1	?
..	?
..	?
1	1	1	?

- For two classes, total number of values (parameters) needed = $2 * (2^3 - 1)$
- For n Boolean variables, parameters = $2 * (2^n - 1)$
- That's too many parameters to estimate. Can we do better?




Generative vs Discriminative Classifiers

- How does conditional independence help
 $P(\mathbf{X} \mid Y) = P(X_1 \mid Y) * P(X_2 \mid Y) * P(X_3 \mid Y)$

e.g. for class $Y = 1$, if I know $P(X_1 = 1)$, I know $P(X_1 = 0)$. So, one parameter for each attribute per class.

- In this case, we just need 3 parameters for each class. We needed 7 without conditional independence (CI) assumption.
- For n Boolean attributes and 2 classes, we need $2n$ parameters, if we make the conditional independence assumption. We needed $2 * (2^n - 1)$ without conditional independence (CI) assumption.



Generative vs Discriminative Classifiers

- Well, what if we make a functional model for **probability**:

$$P(Y=1 \mid \mathbf{X}) = h(\mathbf{X})$$

- Sounds like a good idea!
- What should be the form of h ? How do we learn it from data?
- This is the focus of **discriminative classifiers**.
- Logistic regression is an example of discriminative classifiers.

Review of Types of Classifiers

3 types of classifiers:

1. Create a model for y (output) as a function of attributes. e.g. Perceptron, ANN, SVM

$$y = \text{sign}(w^T x)$$


2. Probabilistic (Generative) classifiers

$$P(Y | X) \propto P(X | Y) * P(Y)$$

We estimate likelihood and prior from training data.

3. Discriminative classifiers –

Create a model for $P(Y|X)$ – Logistic Regression



Today's topic

Discriminative Classifiers

- Discriminative classifiers assume a functional form for $P(Y|X)$.
- It can be shown that for discrete-valued Y and discrete or continuous X , naïve Bayes equation is equivalent to the following functional form for $P(Y | X)$:

$$P(Y = y_k | X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

- Details of derivation are in the updated chapter of Tom Mitchell's book section 3.1.

Logistic
Regression
Learning
Scenario
for
continuous
attributes:

- Consider learning $f: X \rightarrow Y$, where
 - X is a vector of real-valued features, $\langle X_1 \dots X_n \rangle$
 - Y is boolean
 - assume all X_i are conditionally independent given Y
 - model $P(X_i | Y = y_k)$ as Gaussian $N(\mu_{ik}, \sigma_i)$
 - model $P(Y)$ as Bernoulli (π)
- What does that imply about the form of $P(Y|X)$?

You can see the text for a complete derivation of this equation using above assumptions.

$$P(Y = 1 | X = \langle X_1, \dots, X_n \rangle) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

Logistic Regression

Note: Sometimes we may swap the equations for $P(Y=1|X)$ and $P(Y=0|X)$.

- Form of LR for Boolean Y:

$$P(Y = 1 | X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

- This is similar to the logistic function if we let: $-z = w_0 + \sum_{i=1}^n w_i X_i$

$$P(Y = 1 | z) = \frac{1}{1 + \exp(-z)}$$

- Since $P(Y=1 | X) + P(Y=0 | X) = 1$

$$P(Y = 0 | X) = \frac{\exp(w_0 + \sum_{i=1}^n w_i X_i)}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

Logit Function

- If we let $P(Y=1 | X) = p$, and take the ratio of $P(Y=0 | X)$ and $P(Y=1 | X)$,

$$\frac{p}{1-p} = \exp(z)$$

- Take log of both sides:

$$\log\left(\frac{p}{1-p}\right) = z = \beta_0 + \sum_i \beta_i x_i$$

- The LHS of the function is also called the **logit** function

$$\text{logit}(p) = \beta_0 + \sum_i \beta_i x_i$$

Commonly used by
most statistics
textbooks

Logistic Regression

- We assume a functional form for learning $P(Y|X)$
 - logistic or sigmoid function
- Plot shown on the right.
- We have encountered this function before as an activation function for perceptrons.





Figure 1: Form of the logistic function. In Logistic Regression, $P(Y|X)$ is assumed to follow this form.

Logistic Regression

- We are creating a model for $P(Y|X)$, but what does the decision boundary in the attribute (e.g. x_1, x_2) plane look like.

Remember for linear classification:

- The separating hyperplane has the equation: $w^T x + b = 0$
- The two classes are decided by whether $w^T x + b > 0$ or < 0
- Can we get something similar in logistic regression?



Functional Form: Two classes

$$P(Y = 1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

implies

$$P(Y = 0|X) = \frac{\exp(w_0 + \sum_{i=1}^n w_i X_i)}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

So, logistic regression is a linear classifier after all 😊

Classification Rule: Assign the label $Y=0$ if

$$1 < \frac{P(Y = 0|X)}{P(Y = 1|X)}$$

Take logs
and simplify:

$$0 < w_0 + \sum_{i=1}^n w_i X_i$$

linear classification
rule!

$Y=0$ if the RHS > 0

Another way to express it

Logistic Regression can also be expressed as:

$$h_{\theta(x)} = P(Y = 1|X) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$P(Y = 0|X) = 1 - h_{\theta(x)} = 1 - g(\theta^T x) = \frac{e^{-\theta^T x}}{1 + e^{-\theta^T x}}$$

To predict class = 1, $\frac{P(Y=1|X)}{P(Y=0|X)} > 1$ or $e^{\theta^T x} > 1$ or $\theta^T x > 0$

Logistic Regression

- The ratio of probability of success and failure is called **odds ratio** in statistics.
- In our case, it would be ratio of classes i.e. $P(Y=1 | X)$ divided by $P(Y=0 | X)$.
- Let's do some probability,

Statistical Insight – Odds Ratio

What is the odds (or odds ratio)?

Suppose you have following data:



What is the probability of choosing a red ball $P(F)$: $\frac{3}{12}$

What is the probability of NOT choosing a red ball $P(\neg F)$: $\frac{9}{12}$

We define **odds for (of) red** as: $\frac{\text{Favorable Outcomes}}{\text{Unfavorable Outcomes}} = \frac{P(F)}{P(\neg F)} = \frac{3}{9}$

Statistical Insight – Odds Ratio

What is the odds (or odds ratio)?

Suppose you have following data:



What is the probability of choosing a red ball : $\frac{3}{12}$

What is the probability of NOT choosing a red ball : $\frac{9}{12}$

We define **odds against red** as: $\frac{\text{Unfavorable Outcomes}}{\text{Favorable Outcomes}} = \frac{P(\neg F)}{P(F)} = \frac{9}{3}$

Statistical Insight – Odds Ratio

What are the odds of winning a game of roulette?

Choices are numbers:

$$1 - 36 + 0 + 00 = \text{Total } 38$$

Only 1 of them wins:

$$\text{Odds of winning: } \frac{1}{37}$$

What if I play on 1st / 2nd / 3rd 12:

$$\text{Odds of winning: } \frac{12}{26}$$

$$\text{Odd / Even Numbers: Odds of winning: } \frac{18}{20}$$



Statistical Insight – Expected Value

What is the probability of winning a game of roulette?

Choices are numbers:

$$1 - 35 + 0 + 00 = \text{Total } 38$$

Only 1 of them wins:

Probability of winning: $\frac{1}{38}$ Probability of losing: $\frac{37}{38}$



If I bet \$1 and my number wins, I get \$35, else I lose \$1.

What's expected value of win/loss?

$$E(W) = 35 * \frac{1}{38} - 1 * \frac{37}{38} = -0.0526 \text{ or } 5.26\% \text{ House Edge}$$

Back to Machine Learning

- Like in any other model, we need to find the parameters – **weights** in this case.
- How do we do that? What techniques do we know for parameter estimation?
 - Gradient Descent of error
 - Maximum Likelihood
 - MAP (naïve Bayes)

Learning the weights

- How do we learn the weights?
- Maximum likelihood to the rescue.
- Remember, we want to maximize the parameters **given the training data**.

Example: Suppose you observe following observations:

$\{ Y = 1 | \mathbf{X1}, Y = 0 | \mathbf{X2}, Y = 0 | \mathbf{X3}, Y = 1 | \mathbf{X4} \}$

We first evaluate how likely is this data in terms of parameters (Θ)

Likelihood (L) = $P(Y=1 | \mathbf{X1}, \Theta) * P(Y=0 | \mathbf{X2}, \Theta) * P(Y=0 | \mathbf{X3}, \Theta) * P(Y=1 | \mathbf{X4}, \Theta)$

Log Likelihood (LL) = $\log[P(Y=1 | \mathbf{X1}, \Theta)] + \log[P(Y=0 | \mathbf{X2}, \Theta)] + \log[P(Y=0 | \mathbf{X3}, \Theta)] + \log[P(Y=1 | \mathbf{X4}, \Theta)]$

=> Differentiate LL w.r.t. Θ , set to 0 and solve for Θ

We know how to get this value in Logistic Regression.

Learning the weights

- The value $P(Y^l|X^l, W)$ is known as conditional likelihood for a training example l .
- The optimization problem can be expressed as:

$$W \leftarrow \arg \max_W \prod_l P(Y^l|X^l, W)$$

$W = \langle w_0, w_1, \dots, w_n \rangle$

vector of parameters to be estimated

Y^l = observed value of Y in the l^{th} example

X^l = observed value of X in the l^{th} example

Learning the weights

- Remember the old trick:
take the log of the likelihood function
- Want to maximize: $W \leftarrow \arg \max_W \sum_l \ln P(Y^l | X^l, W)$
- This is called the **conditional data log likelihood**, called $l(W)$
- Remember Y^l can take only 2 values -> 0 and 1

$$l(W) = \sum_l Y^l \ln P(Y^l = 1 | X^l, W) + (1 - Y^l) \ln P(Y^l = 0 | X^l, W)$$

Not convinced?

Set $Y^l = 0$ and check and then

Set $Y^l = 1$ and check.



Learning the weights

- To proceed, we use the following functional forms:

$$P(Y = 0|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

I know we flipped from earlier definition 😊
But this is done for mathematical convenience


$$P(Y = 1|X) = \frac{\exp(w_0 + \sum_{i=1}^n w_i X_i)}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

- Using in equation for $l(W)$:

$$\begin{aligned} l(W) &= \sum_l Y^l \ln P(Y^l = 1|X^l, W) + (1 - Y^l) \ln P(Y^l = 0|X^l, W) \\ &= \sum_l Y^l \ln \frac{P(Y^l = 1|X^l, W)}{P(Y^l = 0|X^l, W)} + \ln P(Y^l = 0|X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i^n w_i X_i^l)) \end{aligned}$$

Learning the weights

- There is no closed form solution for above equation.
- So how do we proceed? **Gradient Ascent**
- Why ascent? Because we want **maximum** value.
- When we wanted to **minimize** a function e.g. error, we used **gradient descent**




Convex and Concave functions

- For error, we want a convex function.
- For maxima evaluation, concave function.
- What type of function is the Entropy function
 $E = \sum -p * \log_2(p)$ [assume range of p to be between 0 and 1]?


Hint: Use R code:

```
curve(-x*log2(x)-(1-x)*log2(1-x), 0, 1)
```


Convex $f = x^2x + y^2y$




Convex (degenerate) $f = x^2x$



Concave $f = -x^2x - y^2y$



Nonconvex $f = x^2y + 0.3y^2y$



Learning the weights

- We need the slope of the curve:


$$\frac{\partial l(W)}{\partial w_i} = \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

where $\hat{P}(Y^l | X^l, W)$ is the LR prediction using current set of weights W

What's the update rule for gradient ascent?

$$w^{new} = w^{old} \oplus \eta \frac{\partial l}{\partial w}$$

Note the + sign. We are climbing up the hill.



Learning the weights

- This leads to the gradient ascent weight update rule:

$$w_i \leftarrow w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

- This is what most software packages use as a starting point for Logistic Regression.

Learning the weights - Regularization

- Large weights are a sign of overfitting. Let's penalize them:

$$W \leftarrow \arg \max_W \sum_l \ln P(Y^l | X^l, W) - \frac{\lambda}{2} ||W||^2$$

This is called **regularization**.

- If we repeat the steps using this new objective function, we get the following update rule:

$$w_i \leftarrow w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W)) - \eta \lambda w_i$$

where η a small constant that determines step size

Self Study Material:

- The textbook shows a nice derivation of Logistic Regression. Go over it and understand the steps.

<https://www.cs.cmu.edu/~tom/mlbook/NBayesLogReg.pdf>

Summary

- Logistic Regression directly learns the parameters of the model for $P(Y|X)$.
- Naïve Bayes learns the parameters for $P(X|Y)$ and $P(Y)$ and then uses the naïve Bayes equation.
- If we assume that for each value y_k of Y , the distribution of each continuous X_i is Gaussian, then this is called Gaussian Naïve Bayes (GNB).
- The two approaches have been shown as converging towards each other in the textbook and the famous paper:

Jordan, A. (2002). On discriminative vs. generative classifiers: A comparison of logistic regression and naive bayes. *Advances in neural information processing systems*, 14, 841.

Can be downloaded from:

<http://ai.stanford.edu/~ang/papers/nips01-discriminativegenerative.pdf>

R Resources for Logistic Regression

- Good resource with examples

https://cran.r-project.org/web/packages/HSAUR/vignettes/Ch_logistic_regression_glm.pdf

- Another Tutorial

<https://ww2.coastal.edu/kingw/statistics/R-tutorials/logistic.html>

- Tutorial that explains glm, family, and link

<http://data.princeton.edu/R/glms.html>