

# Artificial Intelligence

CS4365 --- Fall 2022

Adversarial Search

Instructor: Yunhui Guo

# Game Playing

## An AI Favorite

- Structured task
- Not initially thought to require large amounts of knowledge
- Focus on games of perfect information



# Game Playing: State-of-the-Art

- **Checkers:** 1950: First computer player. 1994: First computer champion: **Chinook** ended 40-year-reign of human champion Marion Tinsley using complete 8-piece endgame. 2007: Checkers solved!
- **Chess:** 1997: **Deep Blue** defeats human champion Gary Kasparov in a six-game match. Deep Blue examined 200M positions per second, used very sophisticated evaluation and undisclosed methods for extending some lines of search up to 40 ply.
- **Go:** In 2016, **AlphaGo** defeats the human champion




# Types of Games

- Many different kinds of games!
  - Stochastic vs. Deterministic
  - One, two or more players
  - Zero sum?
  - Perfect information?
- Want algorithms for calculating a strategy which recommends a move from each state

# Deterministic Games

- One possible formulation
  - **States**:  $S$  (start at  $s_0$ )
  - **Players**:  $P = \{1 \dots N\}$  (usually take turns)
  - **Actions**:  $A$  (may depend on player / state)
  - **Transition Function**:  $S \times A \rightarrow S$
  - **Terminal Test**:  $S \rightarrow \{\text{True}, \text{False}\}$
  - **Terminal Utilities**:  $S \times P \rightarrow R$
- Solution for a player is a **policy**:  $S \rightarrow A$



# Zero-Sum Games

- Agents have **opposite** utilities (values on outcomes)
- Lets us think of a single value that one (**MAX**) maximizes and the other minimizes
- Adversarial, pure competition

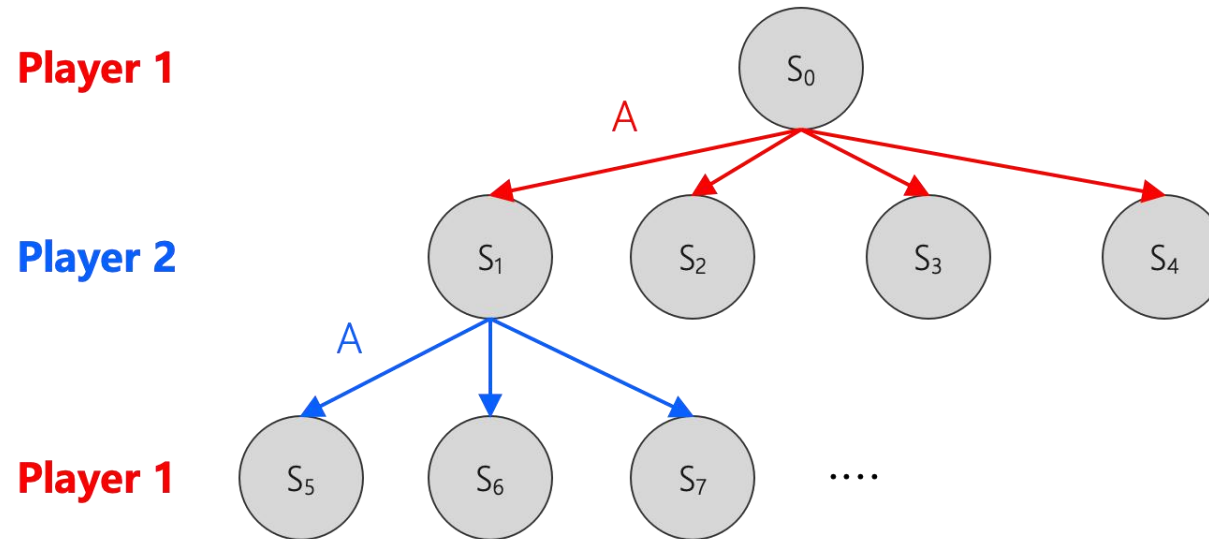


# Games VS. Search Problems

- Unpredictable opponent
  - specifying a move for **every possible** opponent reply
- Time limits
  - unlikely to find goal, must approximate

# Game Playing as Search


- We can list all the possible **actions** and **states**
- In each step, play 1 searches for an **action** which leads to the maximum **utility**






# Search Tree for Tic-Tac-Toe


MAX (x)




# Search Tree for Tic-Tac-Toe




# Search Tree for Tic-Tac-Toe



# Search Tree for Tic-Tac-Toe



# Search Tree for Tic-Tac-Toe




# Tic-Tac-Toe

- High values are **good** for **MAX** and **bad** for **MIN**. It is MAX's job to use the search tree and utility values to determine the best move.
- Root is initial position. Next level are all moves player 1 (MAX) can make; tree is from Max's viewpoint. Next level are all possible responses from player 2 (MIN).
- Max has to find a strategy that will lead to a winning terminal state **regardless of what Min does**. Strategy has to include the correct move for Max for each possible move by Min.

# Pac-Man Trees


- Each step costs 1 point
- The dot worths 10 points



# Value of a State

## Value of a state:

The best  
achievable  
outcome (utility)  
from that state



Non-Terminal States:

$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

$V(s) = \text{known}$



# Adversarial Search Tree


Initial



Pac-Man




Ghost




-20 -8

...




-18 -5

...



-10 +4



-20 +8


# Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:


$$V(s) = \text{known}$$

# Minimax Search

- Why do we take the min value every other level of the tree?
- These nodes represent the opponent's choice of move.
- The computer assumes that the human will choose that move that is of least value to the computer.

# Minimax Search

- **Deterministic, zero-sum games:**
  - Tic-tac-toe, chess, checkers
  - One player maximizes result
  - The other minimizes result
- Compute each node's **minimax value**:
  - the best achievable utility against a rational (optimal) adversary




# Simplified Minimax Algorithm

1. Expand the **entire tree** below the root
2. Evaluate the **terminal nodes** as wins for the minimizer or maximimizer
3. Select an unlabeled node,  $n$ , all of whose children have been assigned values. If there is no such node, we're done — return the value assigned to the root.
4. If  $n$  is a **minimizer** move, assign it a value that is **the minimum of the values of its children**. If  $n$  is a **maximizer** move, assign it a value that is the **maximum of the values of its children**. Return to Step 3.

# Another Example

Max

Min



# Summary

- In game tree search, **a move is a pair of actions**. one player's action is a ply. 2-ply = one move.
- Called a minimax decision because it maximizes the utility under the assumption that the opponent will play perfectly to minimize it.
- **Time complexity:**
  - $O(b^m)$  (m plies and b branching.) Impractical for e.g. chess ( $b \approx 30$  to 40).  $1000^k$  for k moves.
- **Space complexity:**  $O(bm)$


# Size of Search Space

- Chess:
  - branching factor  $b \approx 35$
  - game length  $m \approx 100$
  - search space  $b^m = 35^{100} \approx 10^{154}$
- The Universe
  - number of atoms  $\approx 10^{78}$
- Exact search is infeasible



# The Need for Imperfect Decisions

- Problem:
  - **Minimax** assumes the program has time to search to the terminal nodes.
- In realistic games, cannot search to leaves!
- Solution: Cut off search earlier and apply a heuristic evaluation function to the leaves
- Guarantee of optimal play is gone



# Static Evaluation Functions

- **Minimax** depends on the translation of board quality into a single, summarizing number. Difficult. Expensive
- **Evaluation functions** score non-terminals in depth-limited search
  - Do you control the center of the board?
  - How well protected is your king?
  - Add up values of pieces each player has (weighted by importance of piece).
  - Mobility
- Strategies change as the game proceeds.

# Design Issues for Heuristic Minimax

Evaluation Function:


What **features** should we evaluate and how should we use them?

An evaluation function should:

1. Match utility function on terminal states
2. Not take too long
3. Accurately reflect the chance of winning



# Evaluation Functions for Tic-Tac-Toe

- Let  $p$  be a position in the game
- Define the utility function  $f(p)$  by
  - count the number of lines where  $X$  can win
  - subtract number of lines where  $O$  can win




# Evaluation Functions for Tic-Tac-Toe

- $f(p)$  = the number of lines where X can win - the number of lines where O can win




$$8 - 8 = 0$$



$$8 - 4 = 4$$



$$6 - 4 = 2$$



$$6 - 2 = 4$$

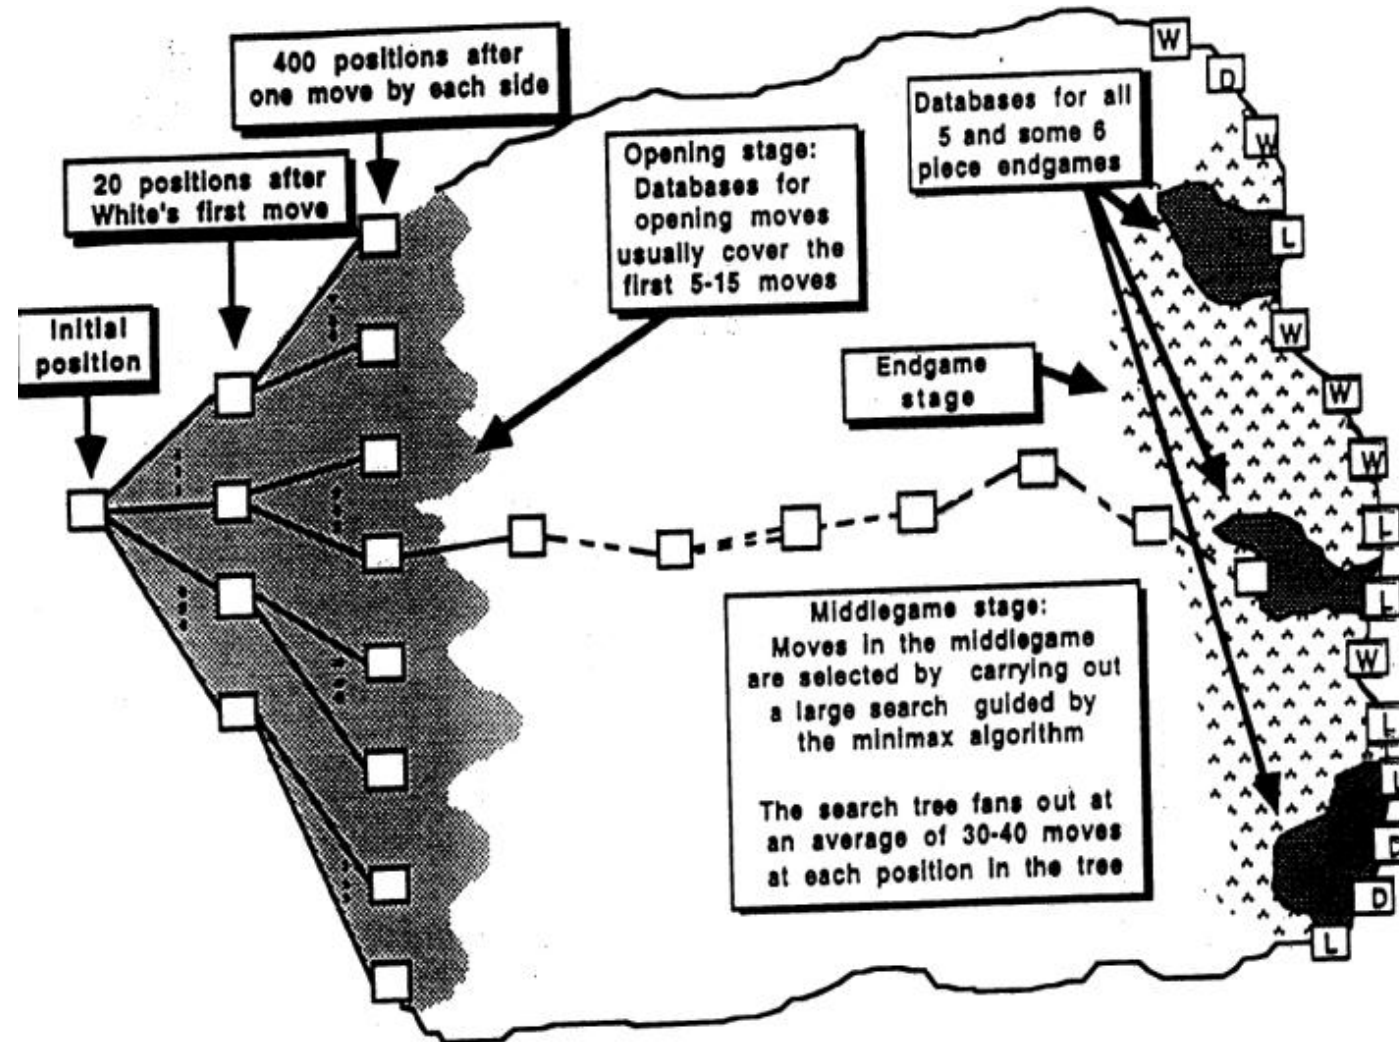
# Linear Evaluation Functions

- Let  $f_i$  be **features** and  $w_i$  be weights
- Linear evaluation function:  $w_1f_1 + w_2f_2 + \dots + w_nf_n$   
This is what most game playing programs use
- For example:  $f = 6 \cdot \text{material} + 4 \cdot \text{mobility} + \text{center control}$

# Linear Evaluation Functions

- Let  $f_i$  be **features** and  $w_i$  be weights
- Linear evaluation function:  $w_1f_1 + w_2f_2 + \dots + w_nf_n$   
This is what most game playing programs use
- Steps in designing an evaluation function:
  1. Pick informative features
  2. Find the weights that make the program play well
- Deep Blue used ~6,000 different features!

# Minimax Search





# Design Issues for Heuristic Minimax


- Depth-limited search:
  - search to a constant depth
- Problems:
  - Some portions of the game tree may be **less stable** than the others
  - Horizon effect

# Unstable States

- Unstable state: drastic change from one level to the next
  - A chess evaluation function that counts material gains may evaluate a given state poorly even the play can capture a queen in the next move.
  - Are you about to lose an important piece?
- Evaluation functions can only be trusted when applied to **stable board states**
- One solution is to extend the normal search to look for stable states

# The Horizon Effect


- You may incorrectly estimate the value of a state by overlooking an event that is just **beyond the depth limit**
  - Opponent moves, move is significant damage and cannot not be avoided
- Fixed-depth searches can be **mislead** by the fact that these damaging moves can be delayed
  - The damage is beyond the search horizon and so is not seen



# The Horizon Effect

- The **negative** horizon effect
  - MAX may try to avoid a bad situation which is actually **inevitable**.
- For example, MAX tries to avoid losing the queen and appears to be able to do so using a lookahead tree of depth 6, but a little deeper it becomes obvious that the queen is going to be lost.

Piece	Value
Pawn	1
Knight	3
Bishop	3
Rook	5
Queen	9



# The Horizon Effect

The **positive** horizon effect:


MAX may not realize that something good is going to be achievable.

For example, MAX would like to take MIN's queen and that can happen  
- but the restricted horizon prevents MAX from making the right choices to realize this possibility

# $\alpha$ - $\beta$ pruning

Max

Min




# Improving Minimax -- $\alpha$ - $\beta$ pruning

- The number of game states is **exponential** in the depth of the tree.
- It is possible to prune large parts of the tree from consideration.
- Intuition: prune away branches that **cannot** possibly influence the value of the state.

# Improving Minimax -- $\alpha$ - $\beta$ pruning

Max

Min

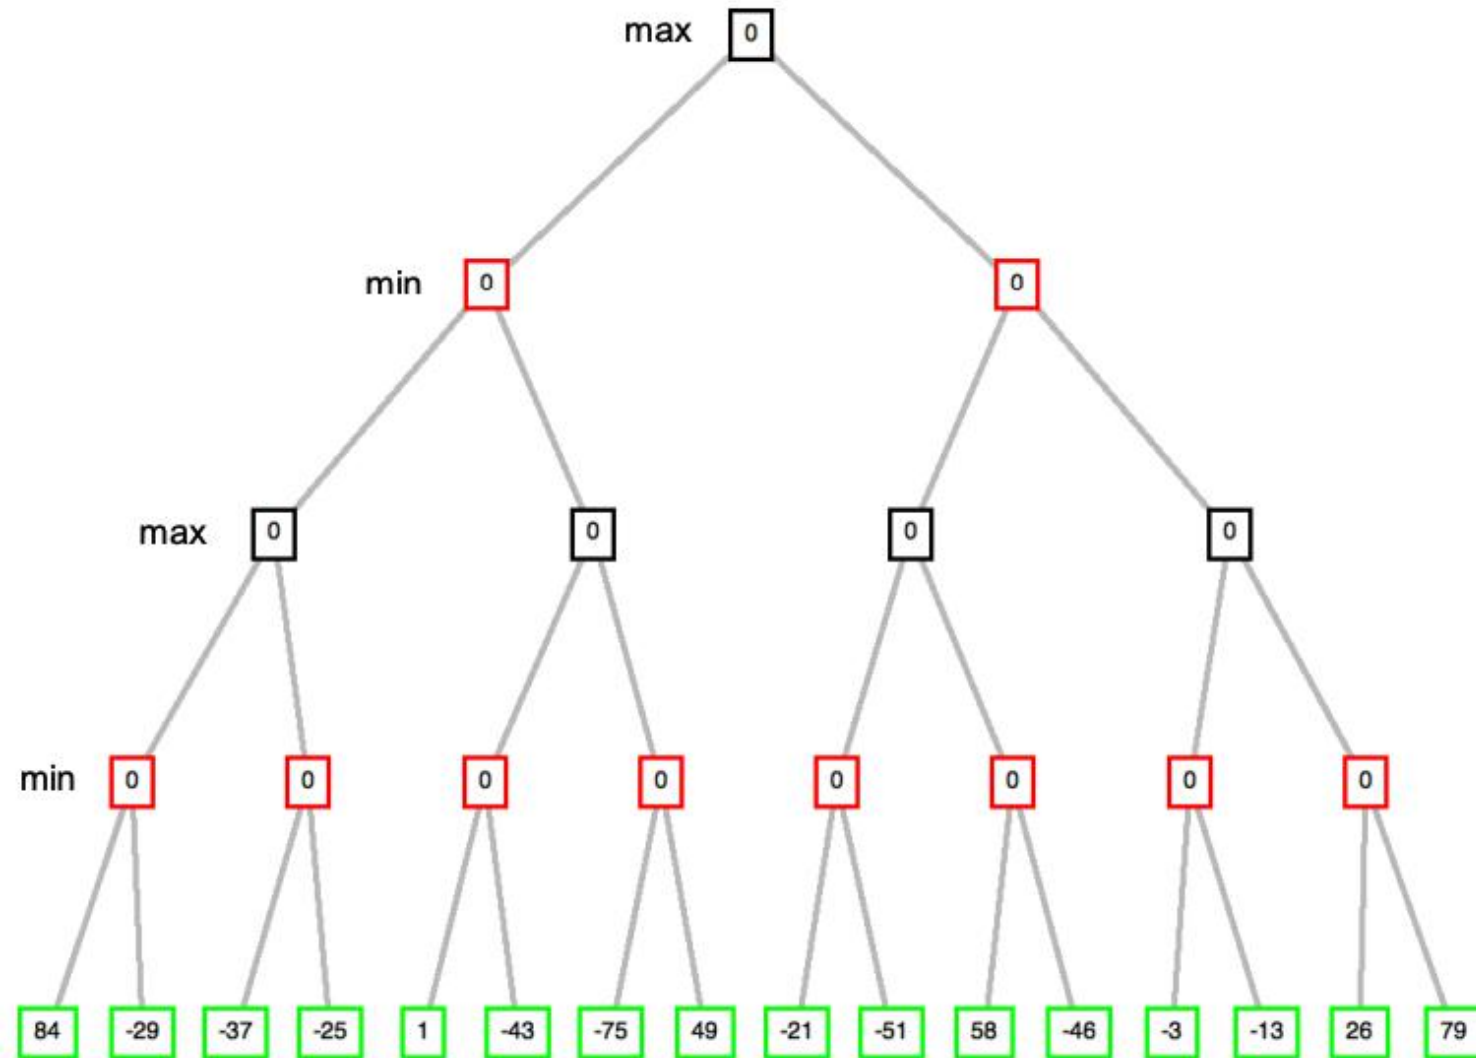




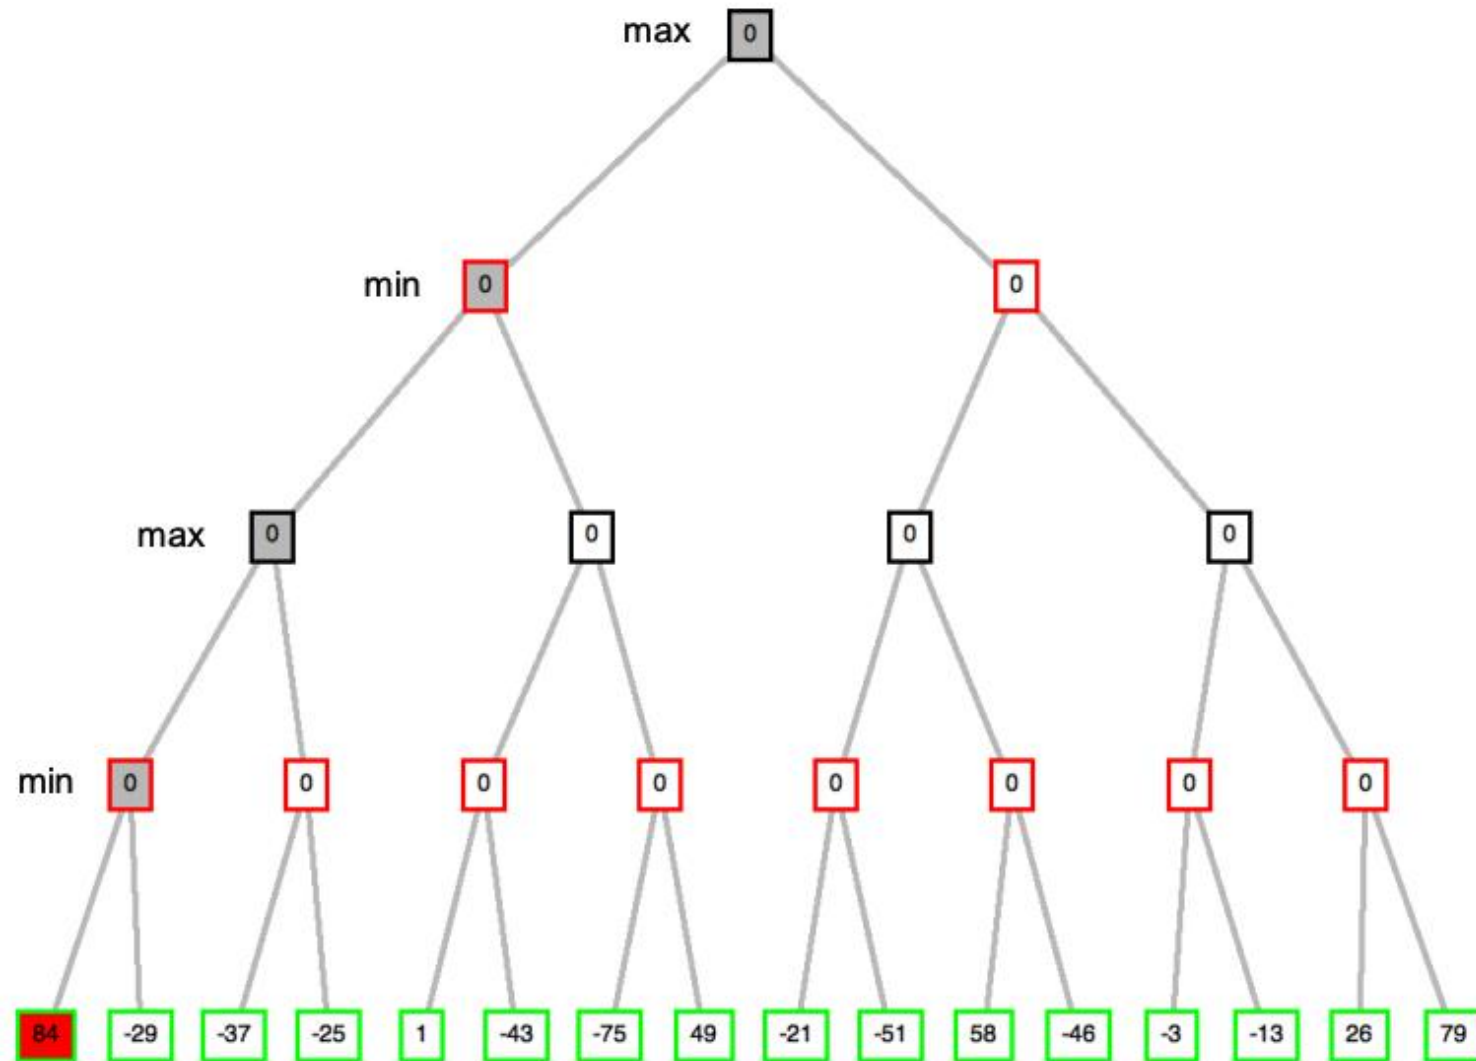
# Algebraic Solution

- $\text{Minimax}(\text{root}) = \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2))$   
     $= \max(3, \min(2, x, y), 2)$   
     $= \max(3, z, 2)$       where  $z = \min(2, x, y) \leq 2$   
     $= 3$
- The value of the root is independent of the values of  $x$  and  $y$

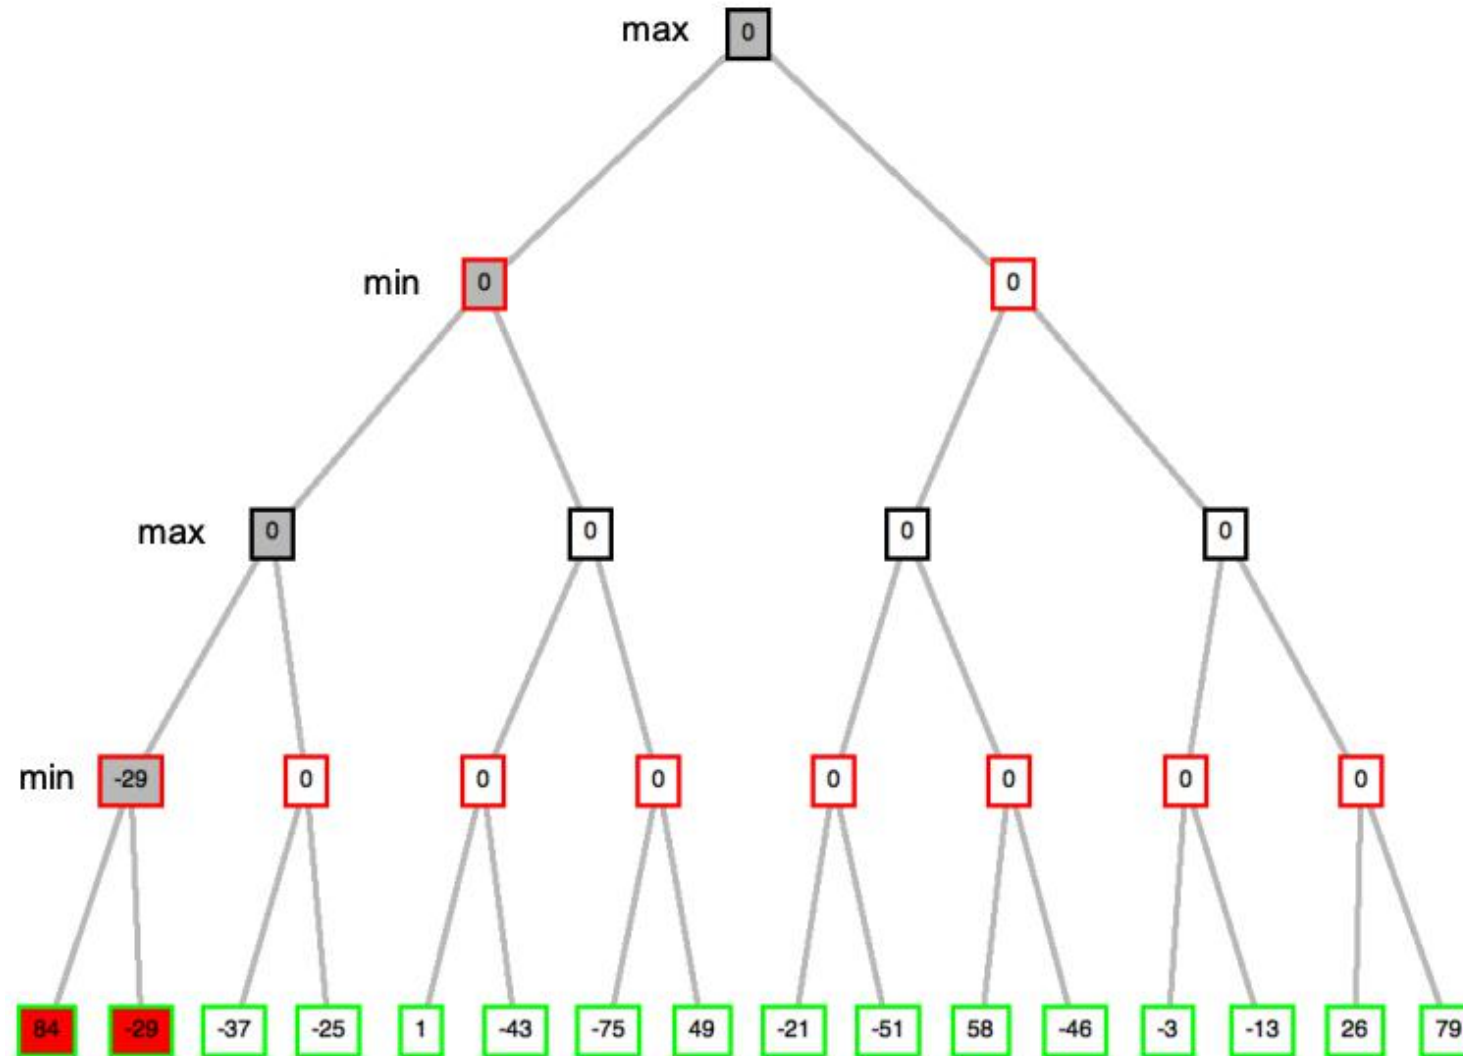
# Another Example



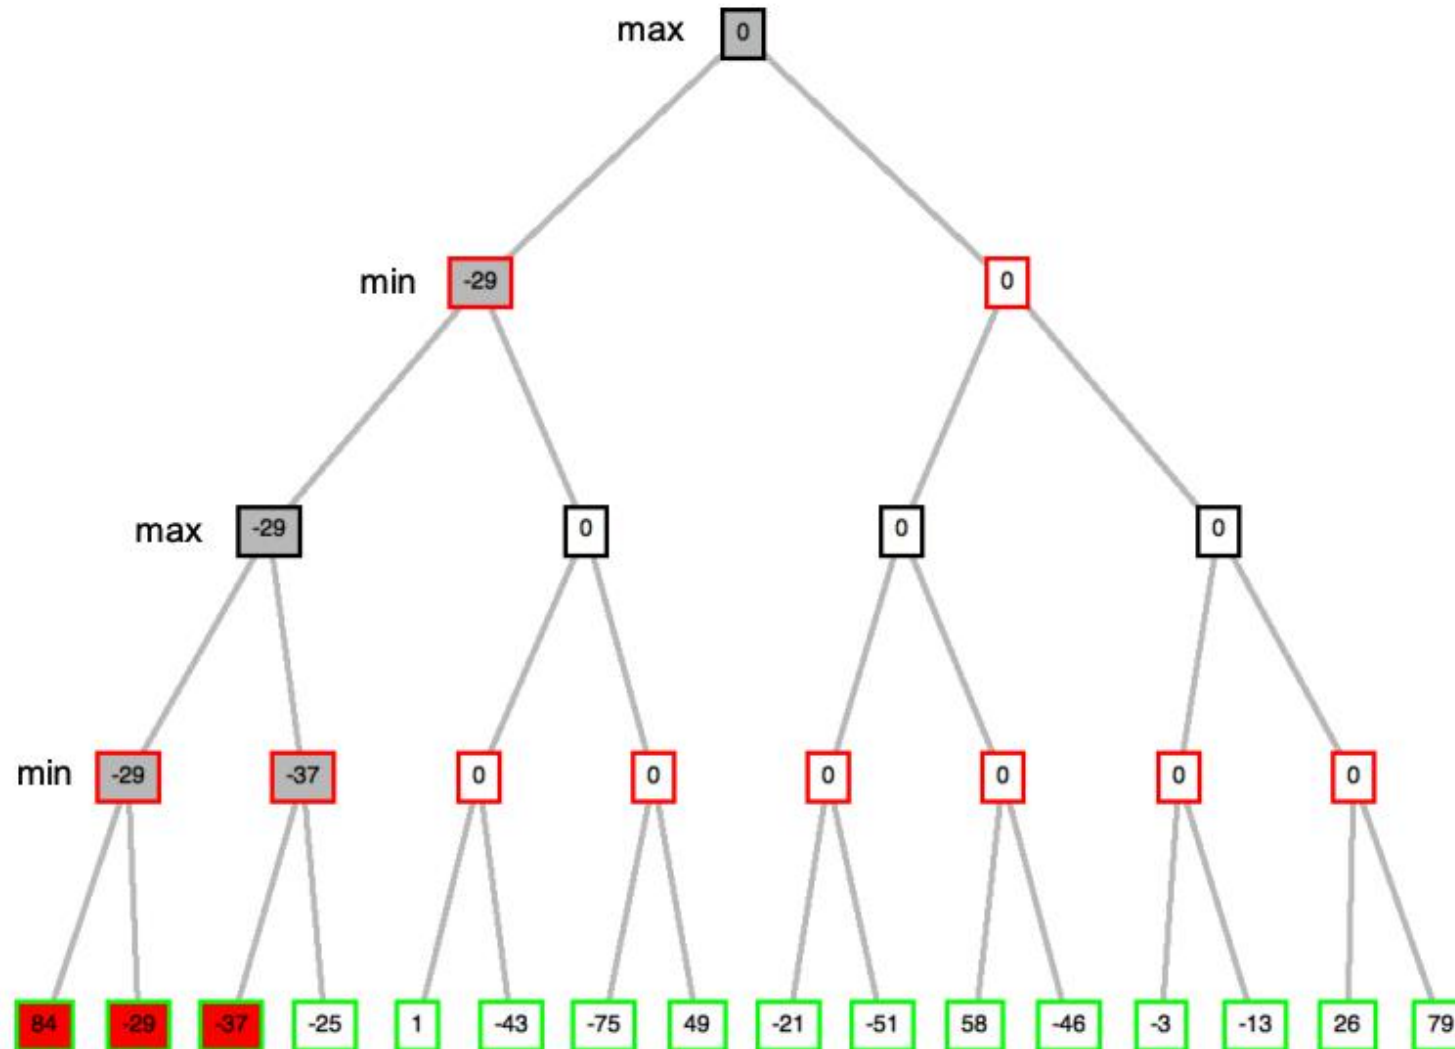
# Another Example



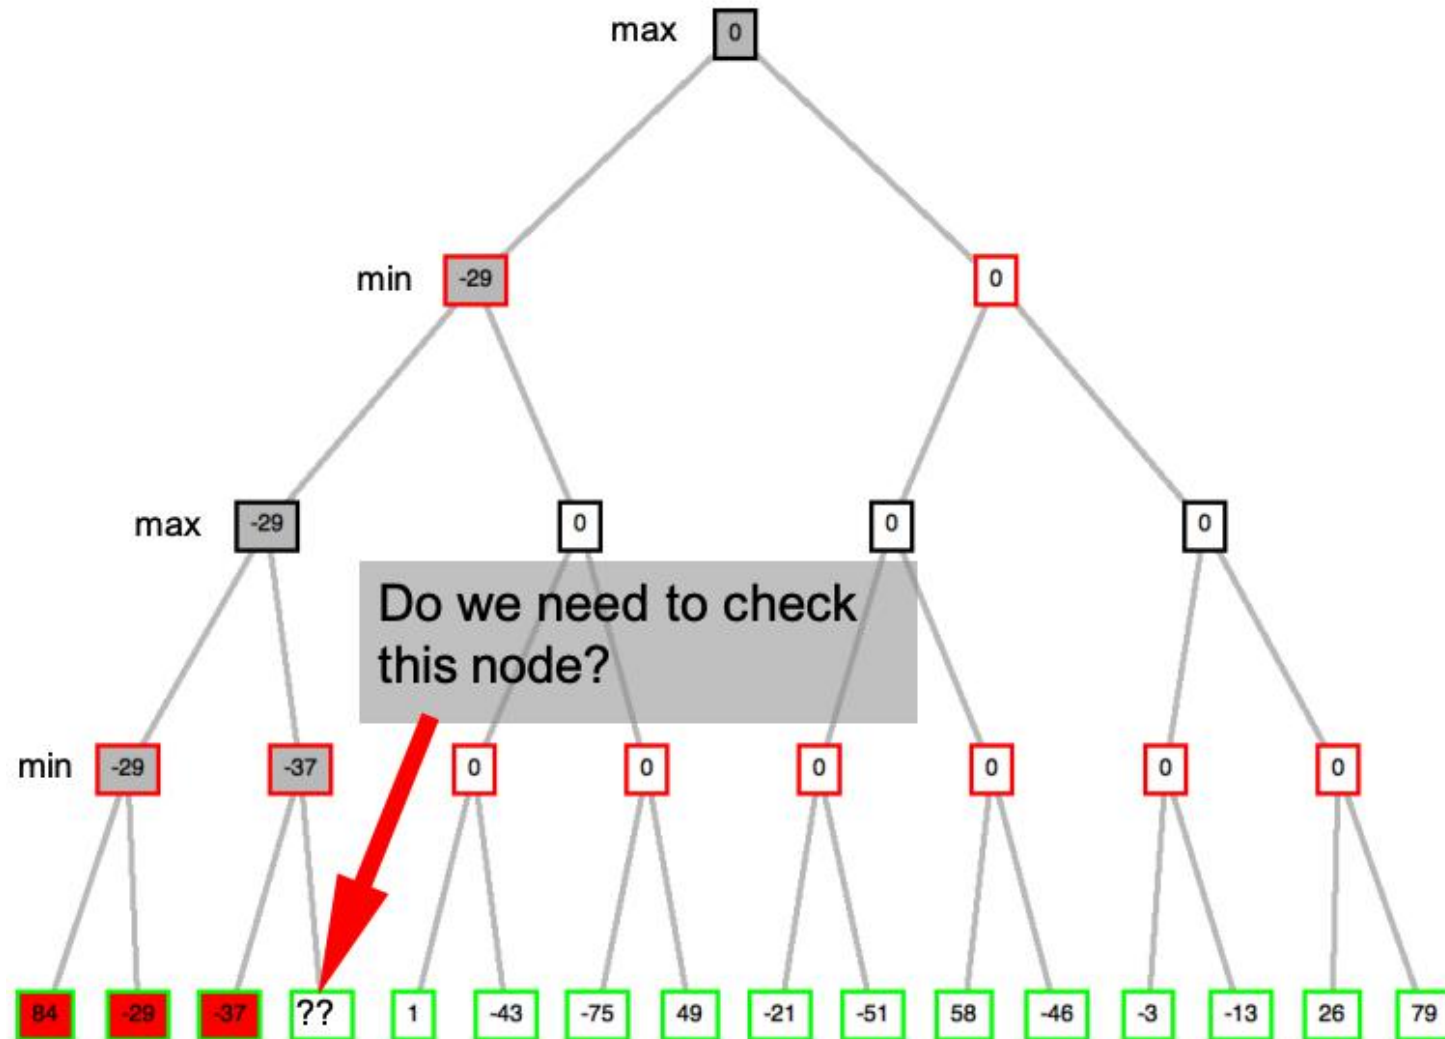
# Another Example



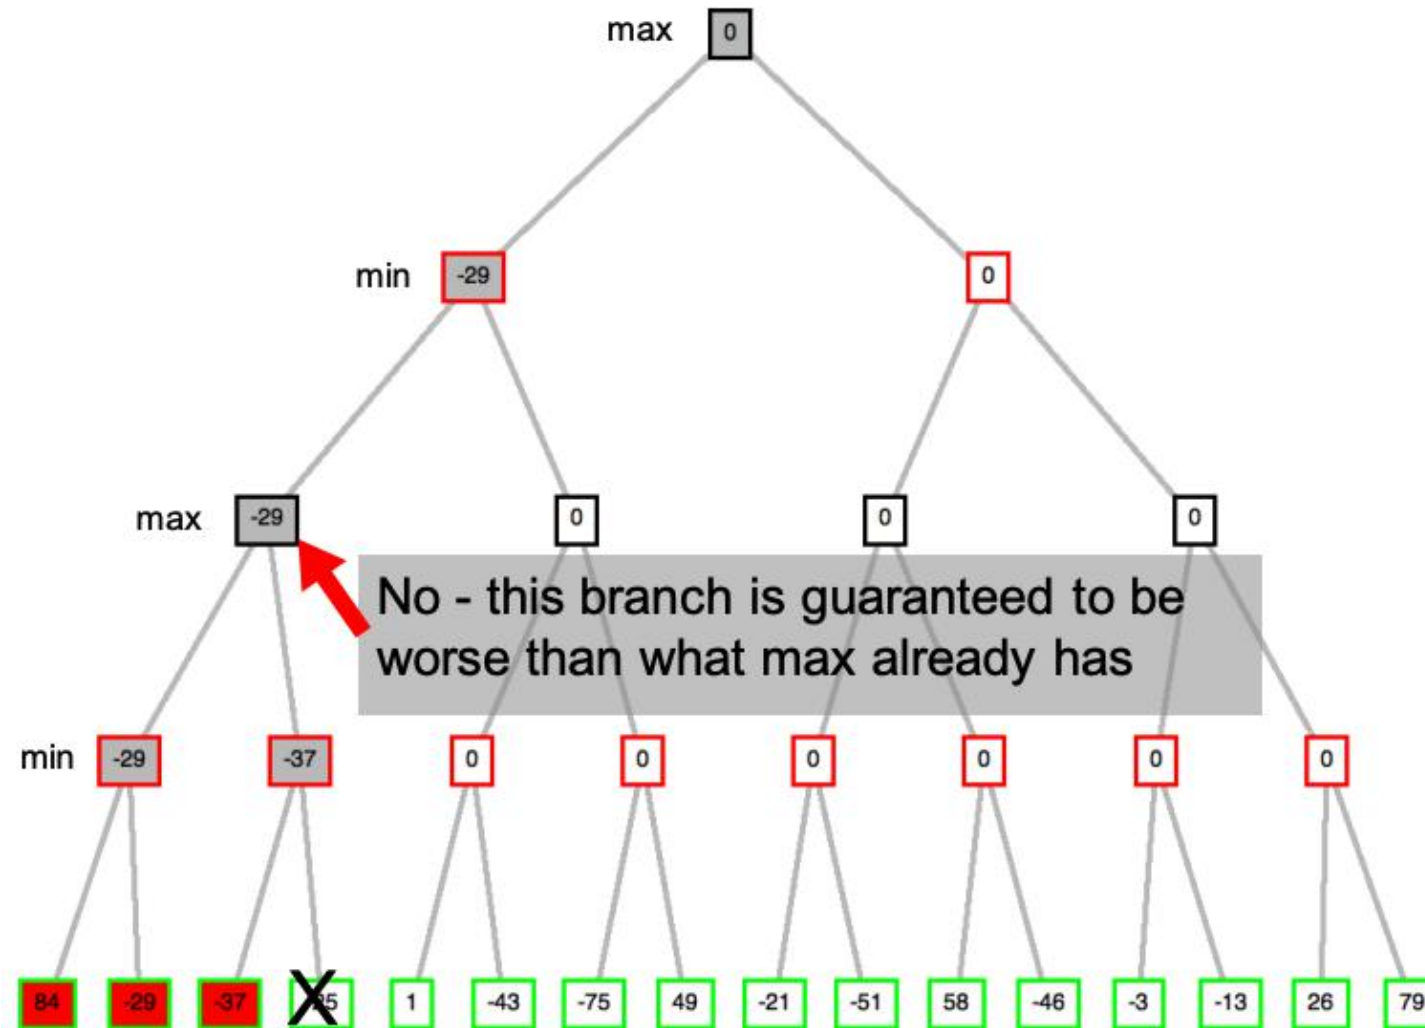
# Another Example



# Another Example



# Another Example



# Improving Minimax -- $\alpha$ - $\beta$ pruning

- The alpha-beta procedure can speed up a depth-first minimax search.
- **Alpha**: a **lower bound** on the value that a **max** node may ultimately be assigned at that level or above

$$v \geq \alpha$$


- **Beta**: an **upper bound** on the value that a **min** node may ultimately be assigned at that level or above

$$v \leq \beta$$



# $\alpha - \beta$ Pruning

- General configuration (MIN version)
  - We're computing the MIN-VALUE at some node  $n$
  - We're looping over  $n$ 's children
  - $n$ 's estimate of the childrens' min is dropping
  - Who cares about  $n$ 's value? MAX
  - Let  $a$  be the best value that MAX can get at any choice point along the current path from the root
  - If  $n$  becomes worse than  $a$ , MAX will avoid it, so we can stop considering  $n$ 's other children  
(it's already bad enough that it won't be played)



# $\alpha - \beta$ Pruning

**Beta:** an upper bound on the value that a **min** node may ultimately be assigned at that level or above

**Alpha:** a lower bound on the value that a **max** node may ultimately be assigned at that level or above

Min


$\beta$



# $\alpha$ - $\beta$ Pruning

**Beta:** an upper bound on the value that a **min** node may ultimately be assigned at that level or above


**Alpha:** a lower bound on the value that a **max** node may ultimately be assigned at that level or above



# $\alpha$ - $\beta$ Pruning

**Beta:** an **upper bound** on the value that a **min** node may ultimately be assigned at that level or above


**Alpha:** a **lower bound** on the value that a **max** node may ultimately be assigned at that level or above



# $\alpha$ - $\beta$ Pruning

**Beta:** an **upper bound** on the value that a **min** node may ultimately be assigned at that level or above


**Alpha:** a **lower bound** on the value that a **max** node may ultimately be assigned at that level or above



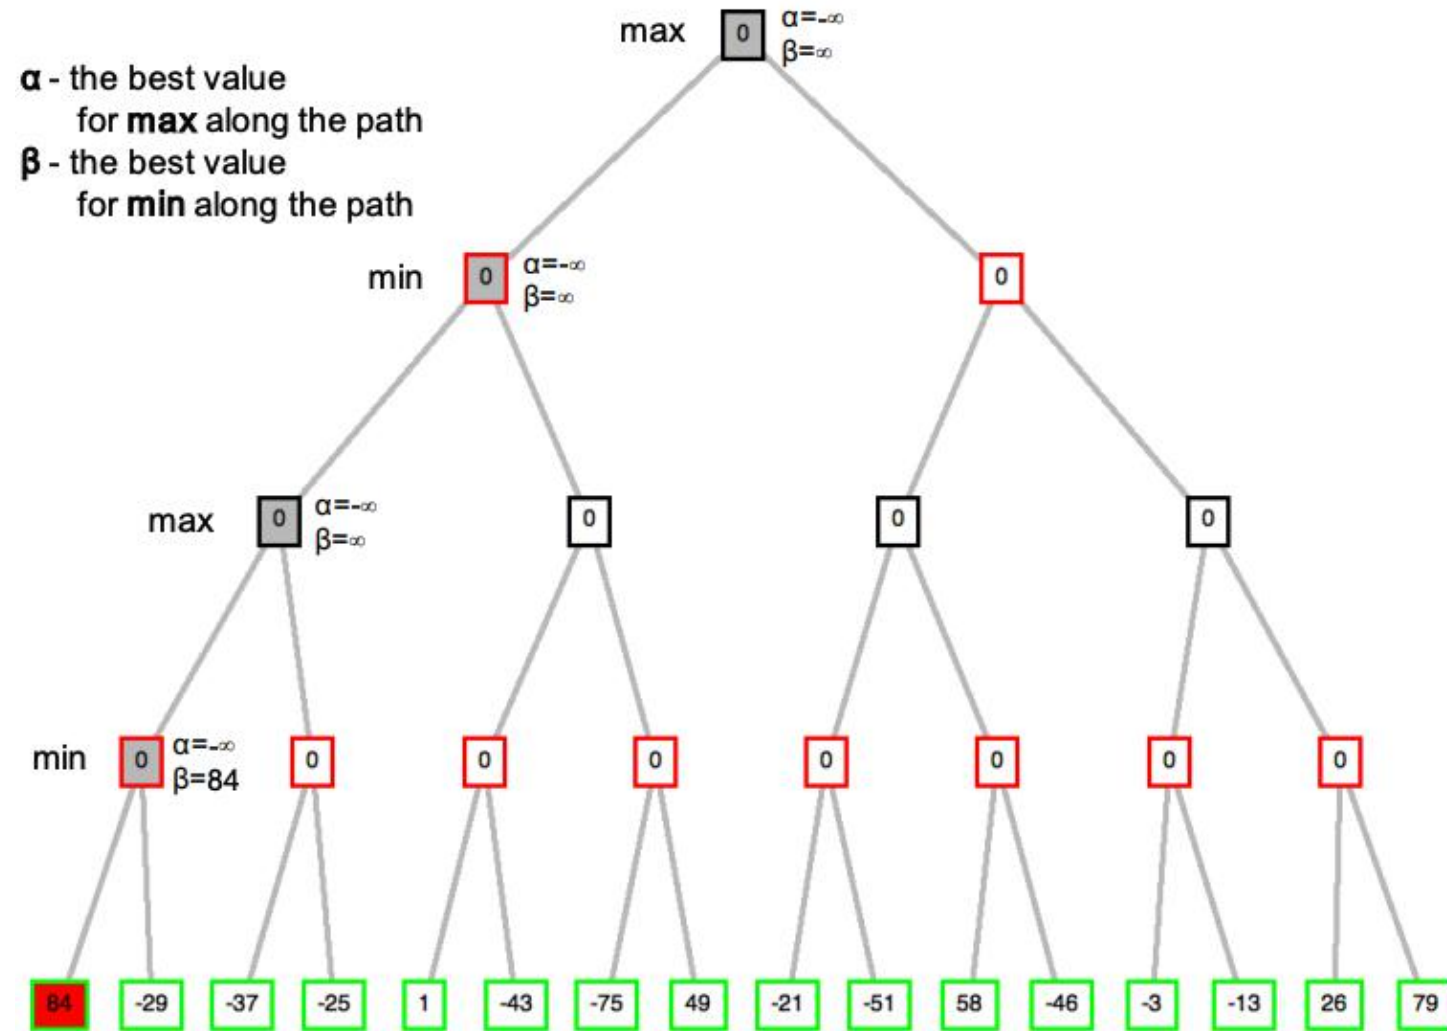
# $\alpha - \beta$ Pruning

**Beta:** an **upper bound** on the value that a **min** node may ultimately be assigned at that level or above

**Alpha:** a **lower bound** on the value that a **max** node may ultimately be assigned at that level or above

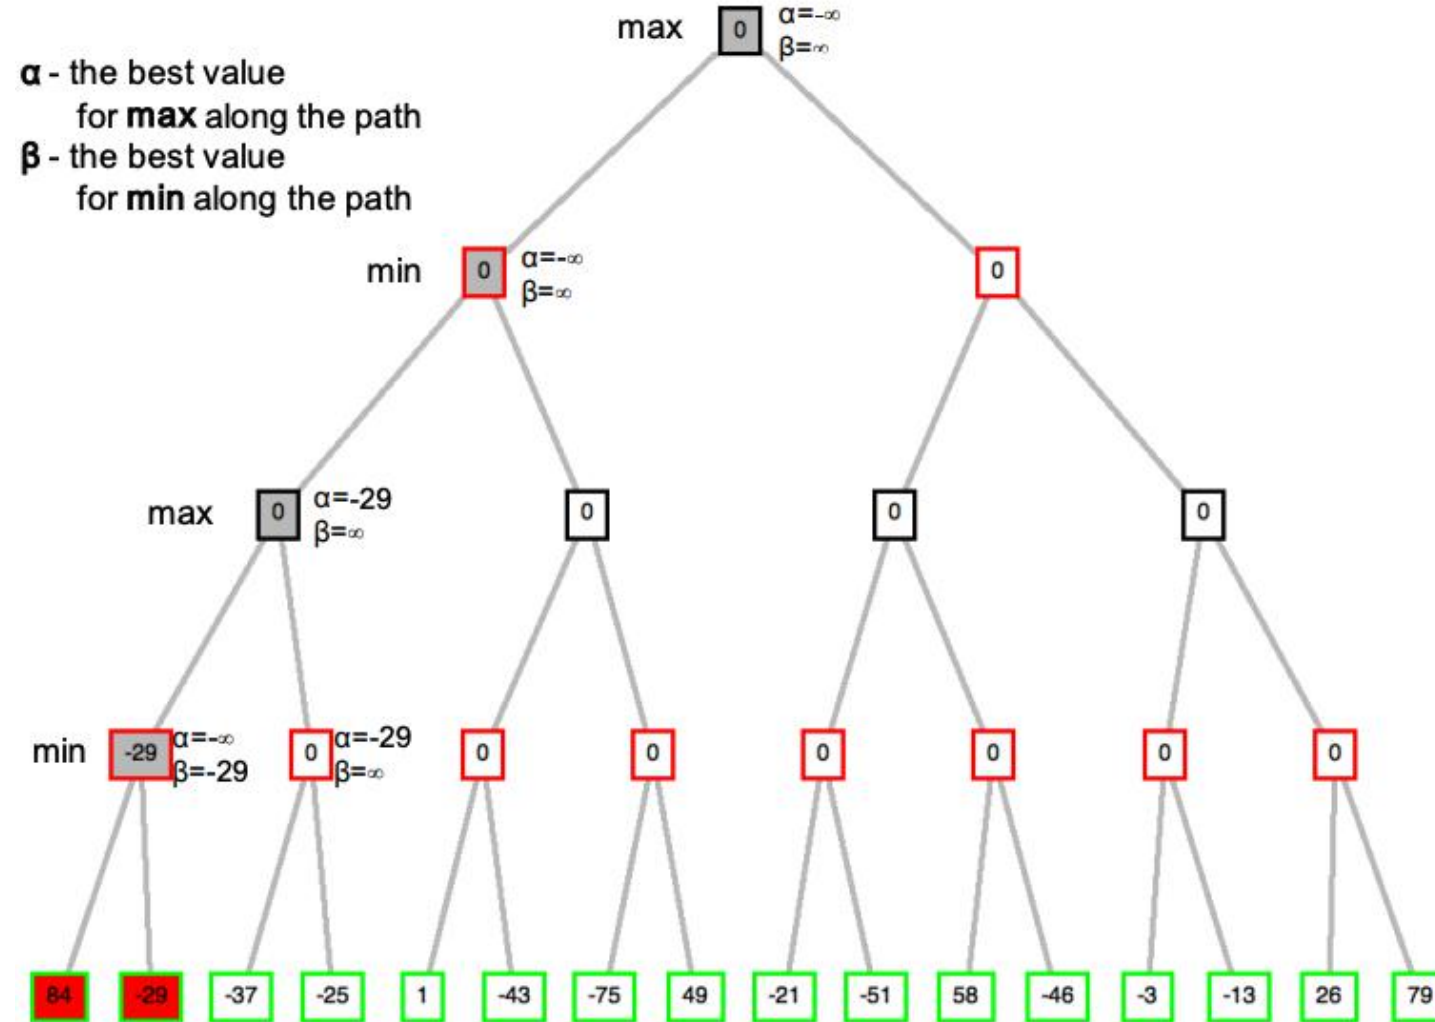


# Another Example



update  $\beta$ ,  
use  $\alpha$  to stop

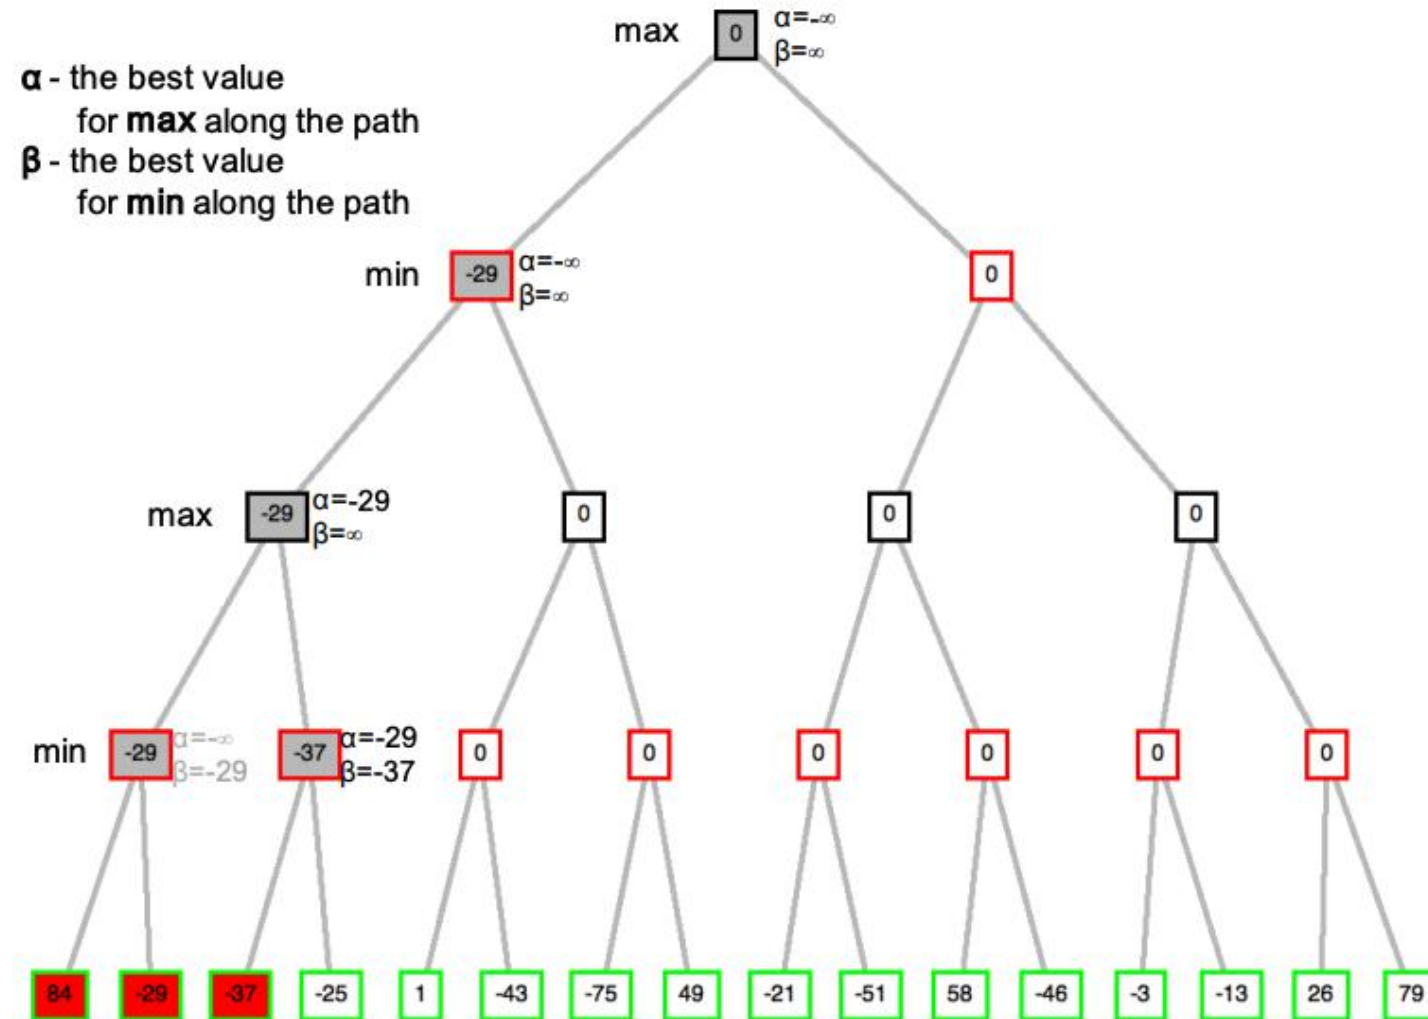
# Another Example



update  $\beta$ ,  
use  $\alpha$  to stop




# Another Example



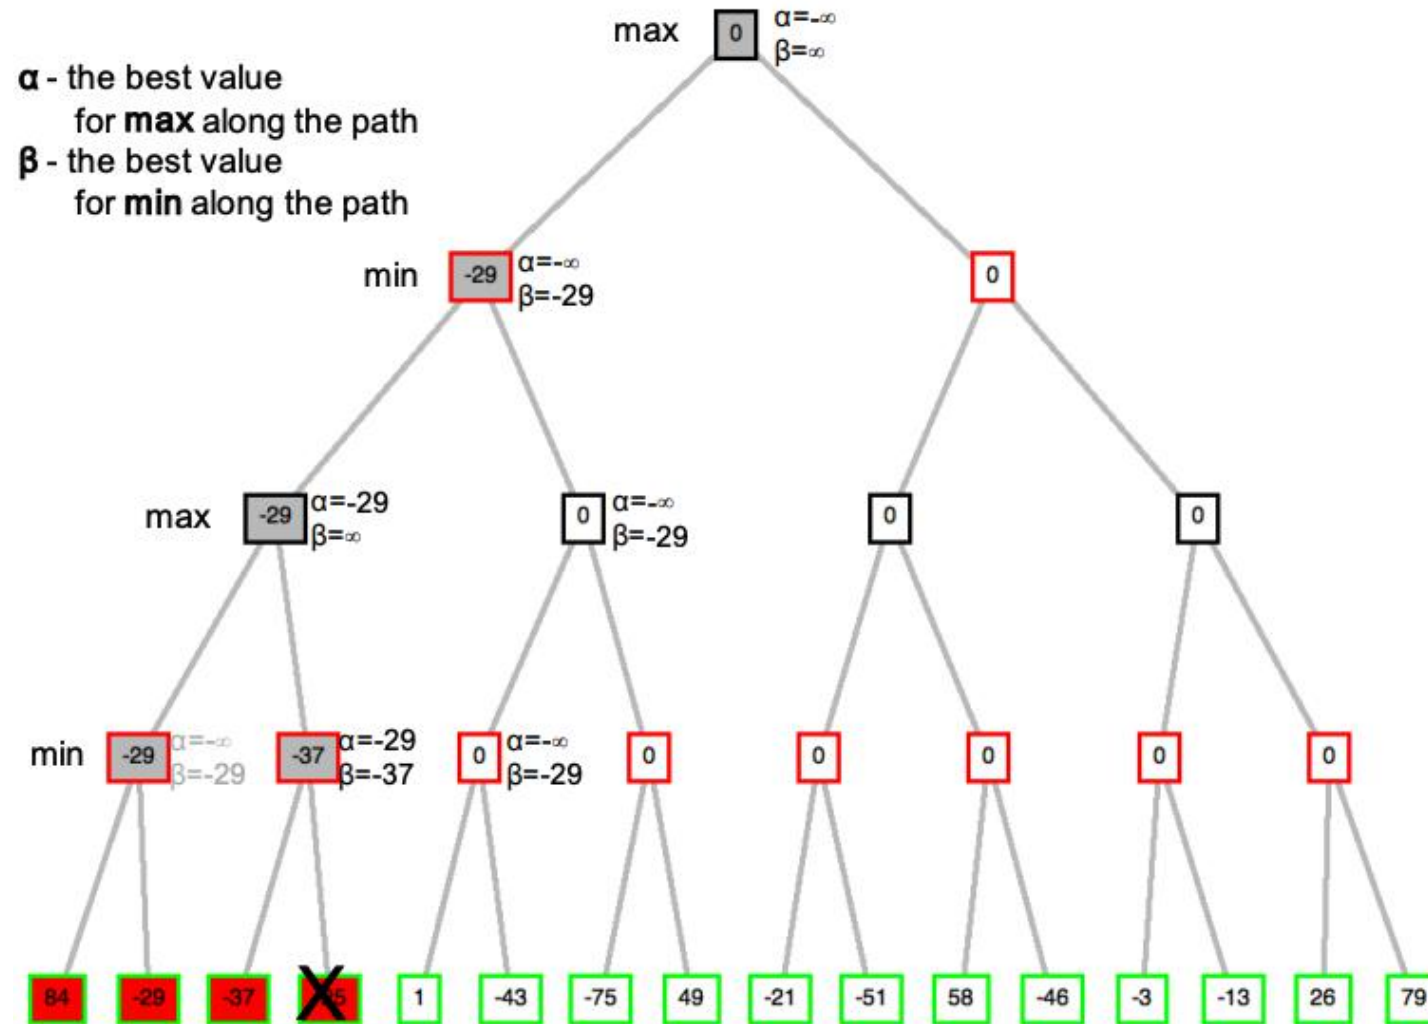
update  $\beta$ ,  
use  $\alpha$  to stop

# Another Example



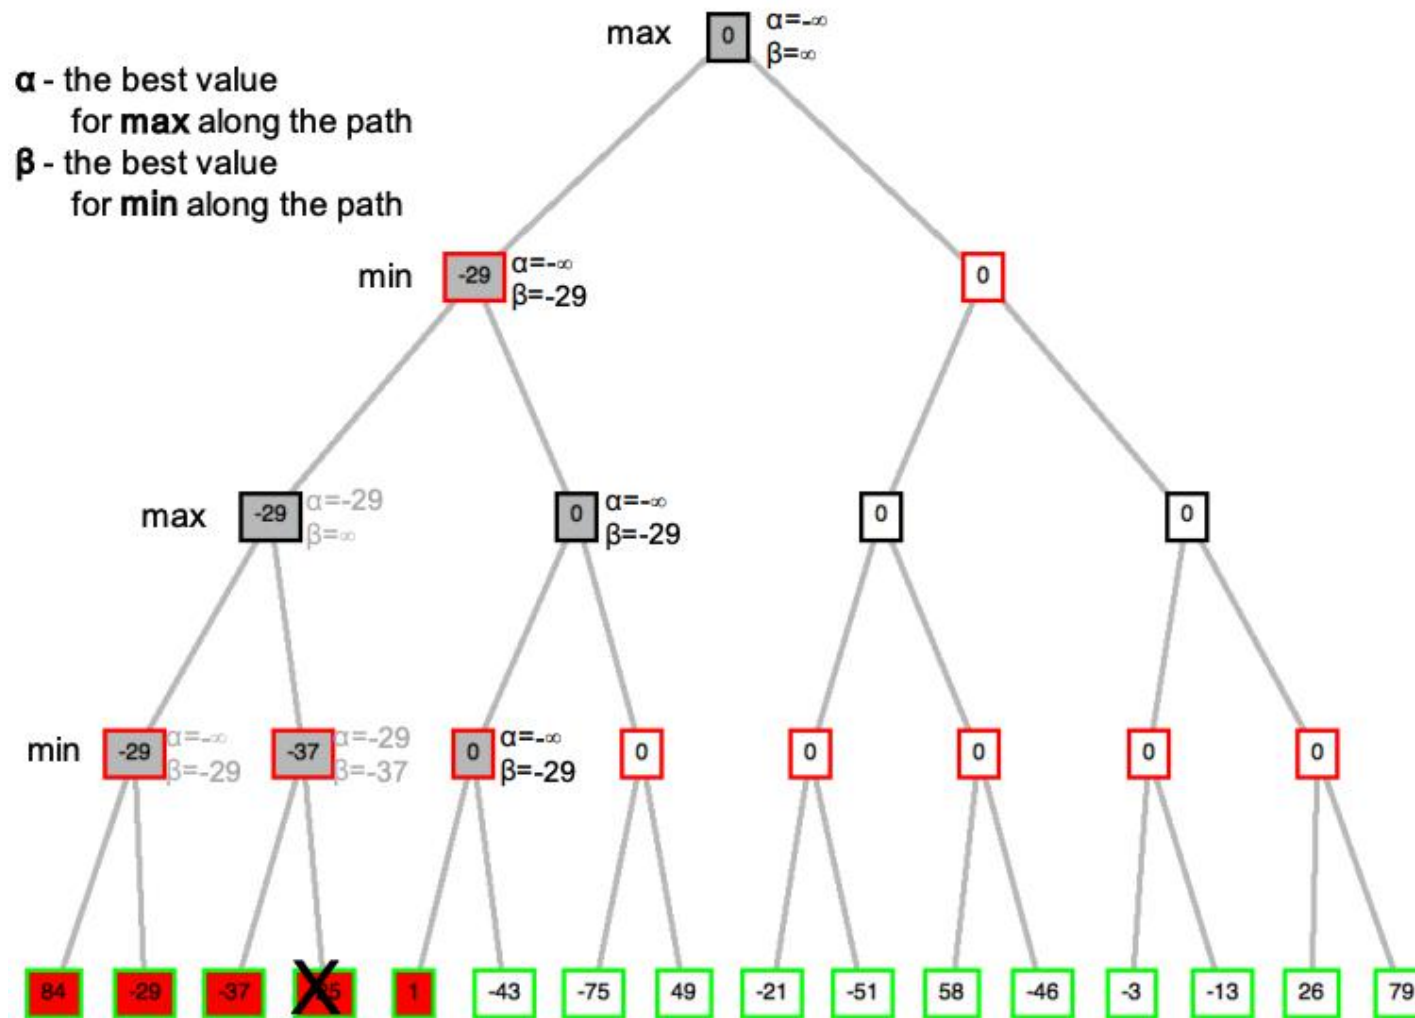
update  $\beta$ ,  
use  $\alpha$  to stop

# Another Example



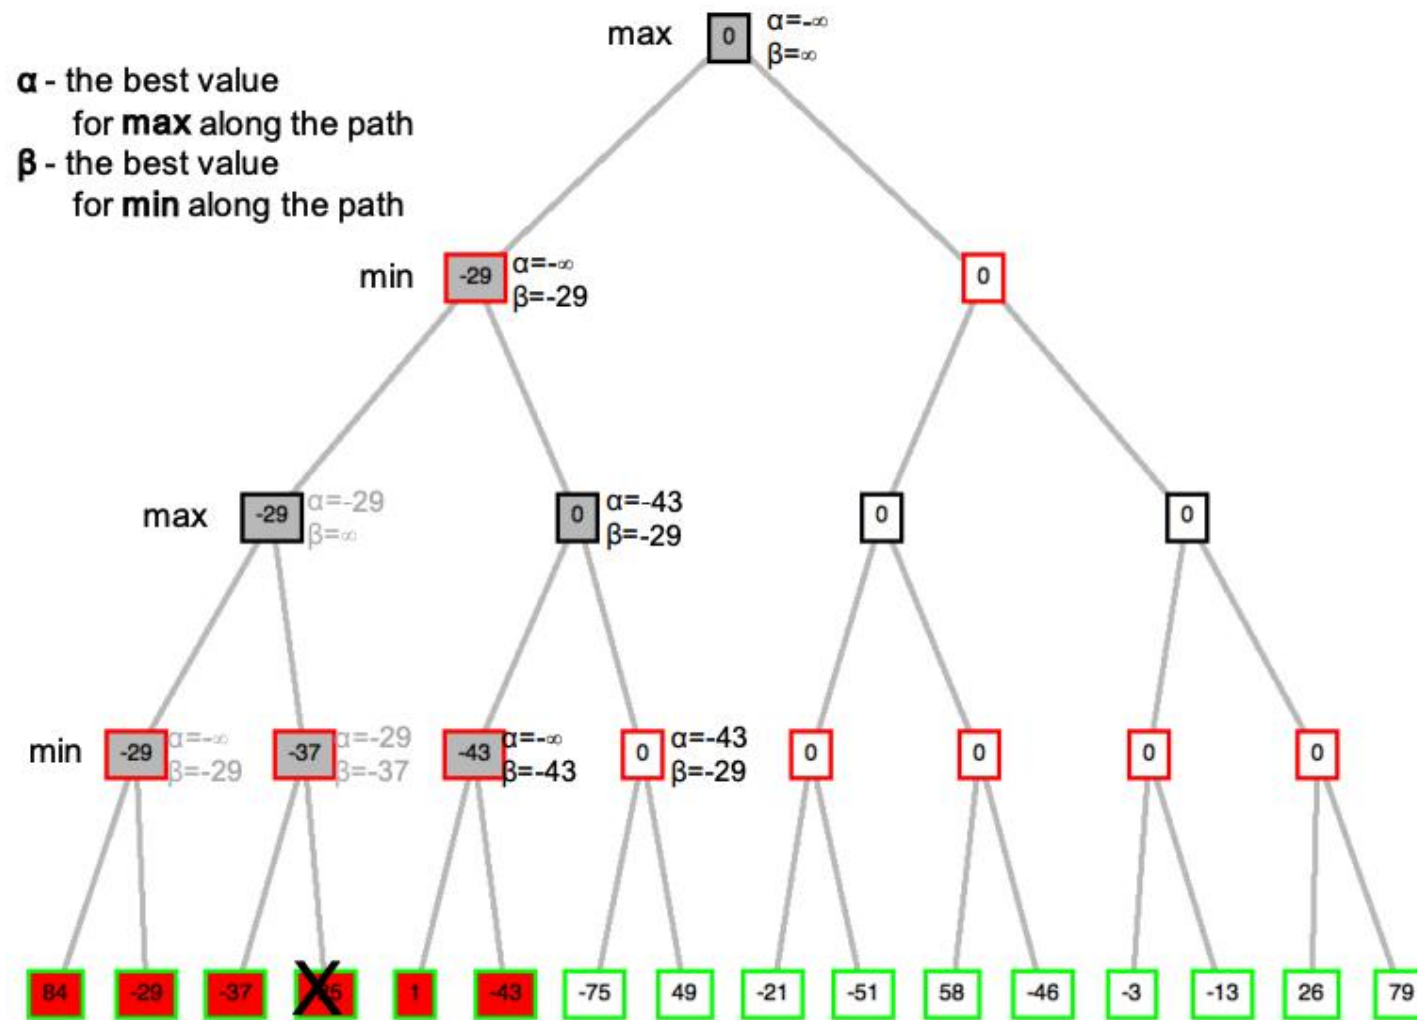
update  $\beta$ ,  
use  $\alpha$  to stop

# Another Example



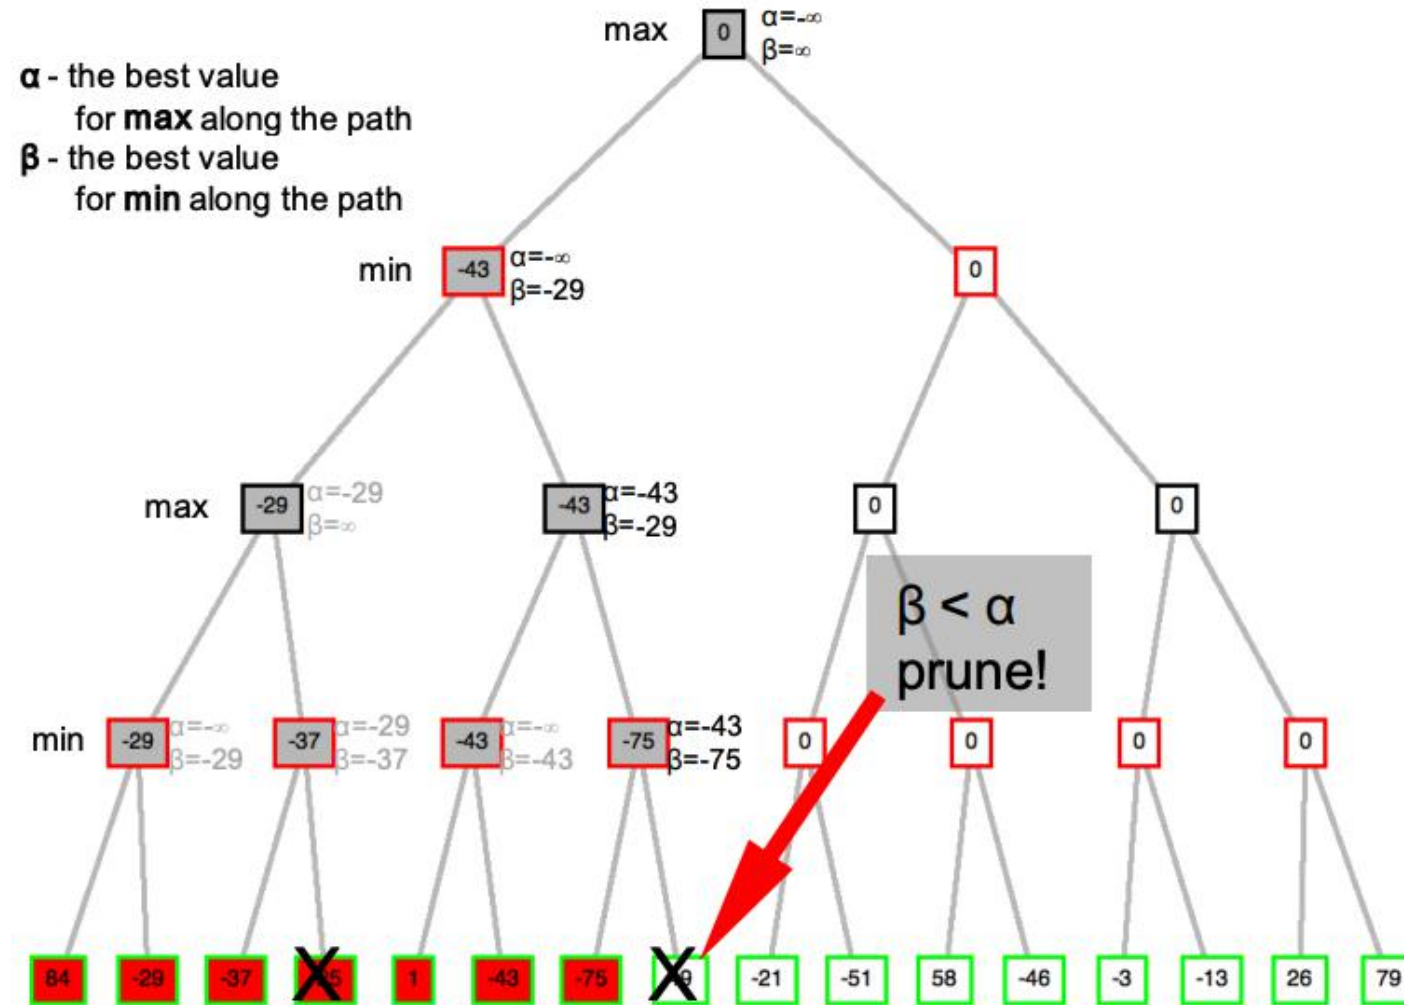
update  $\beta$ ,  
use  $\alpha$  to stop

# Another Example



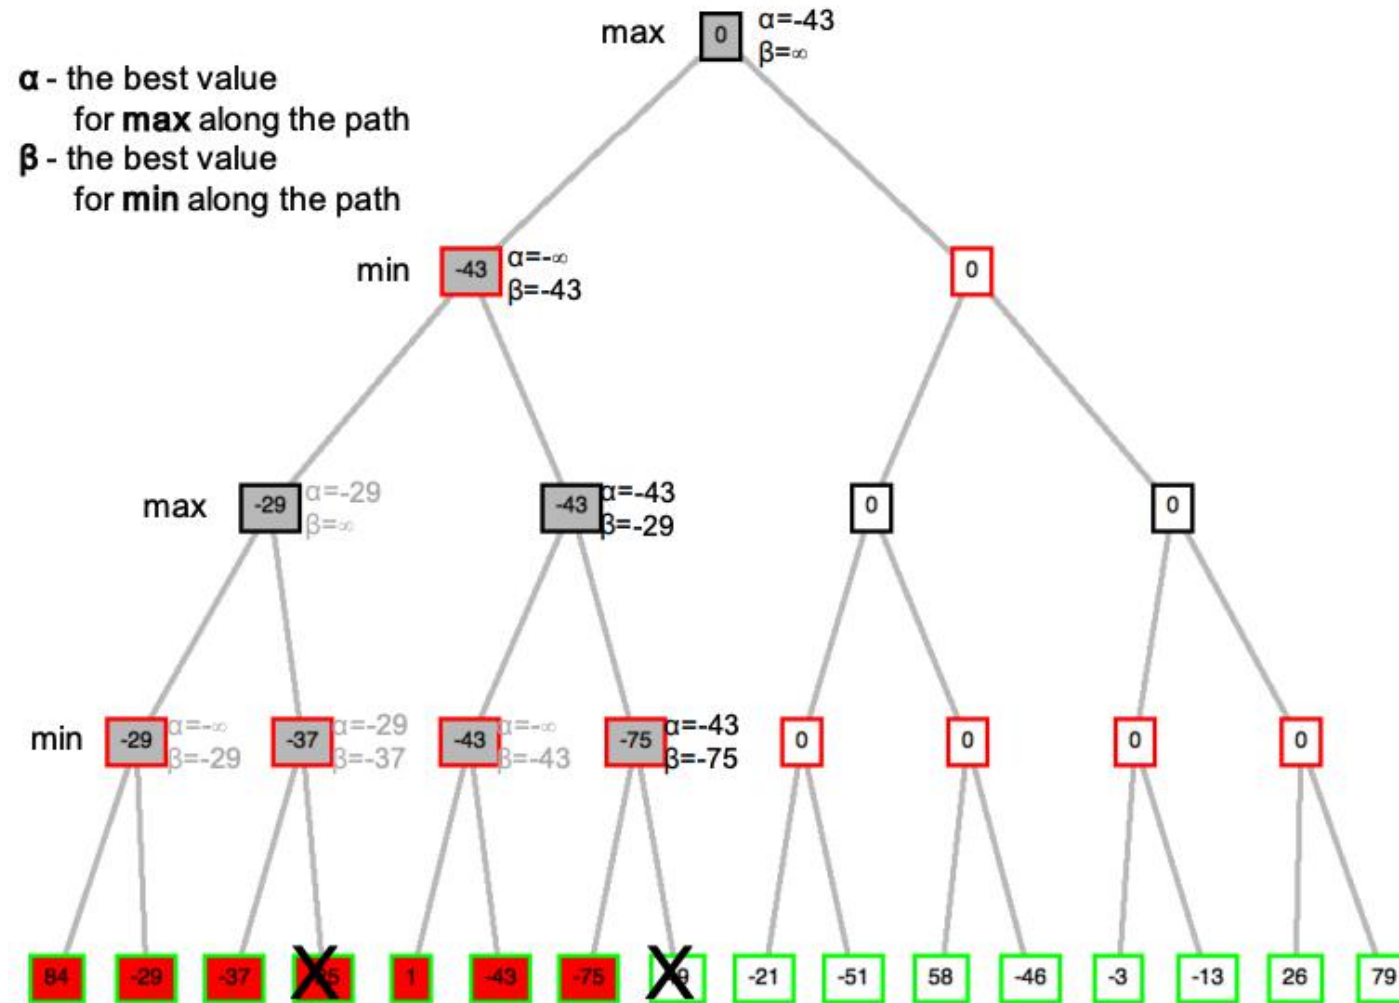
update  $\beta$ ,  
use  $\alpha$  to stop

# Another Example



update  $\beta$ ,  
use  $\alpha$  to stop

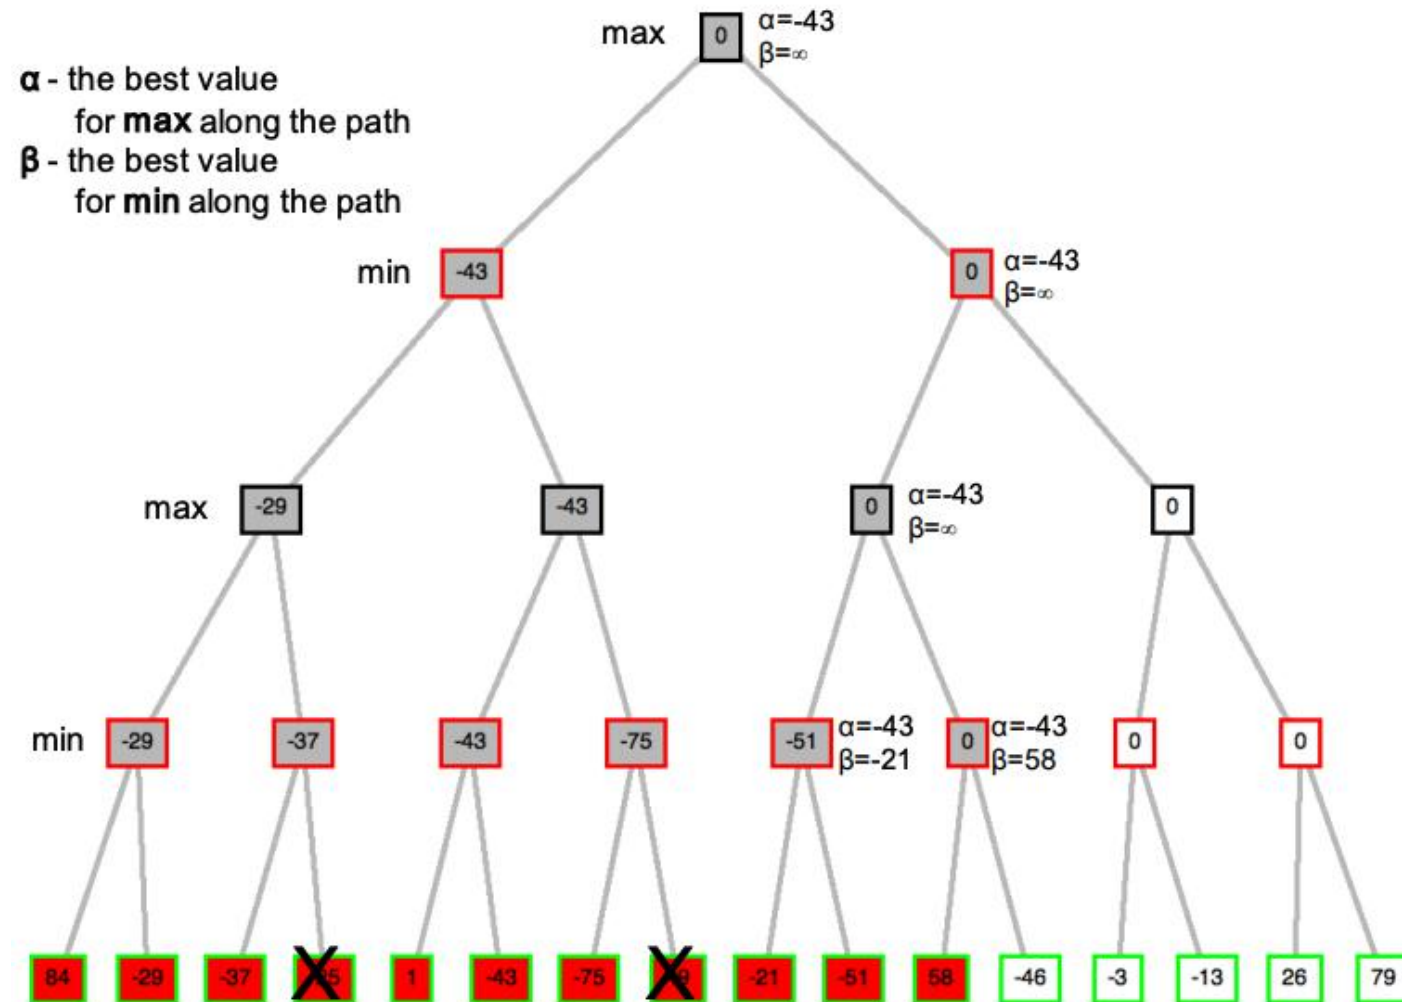
# Another Example



update  $\beta$ ,  
use  $\alpha$  to stop



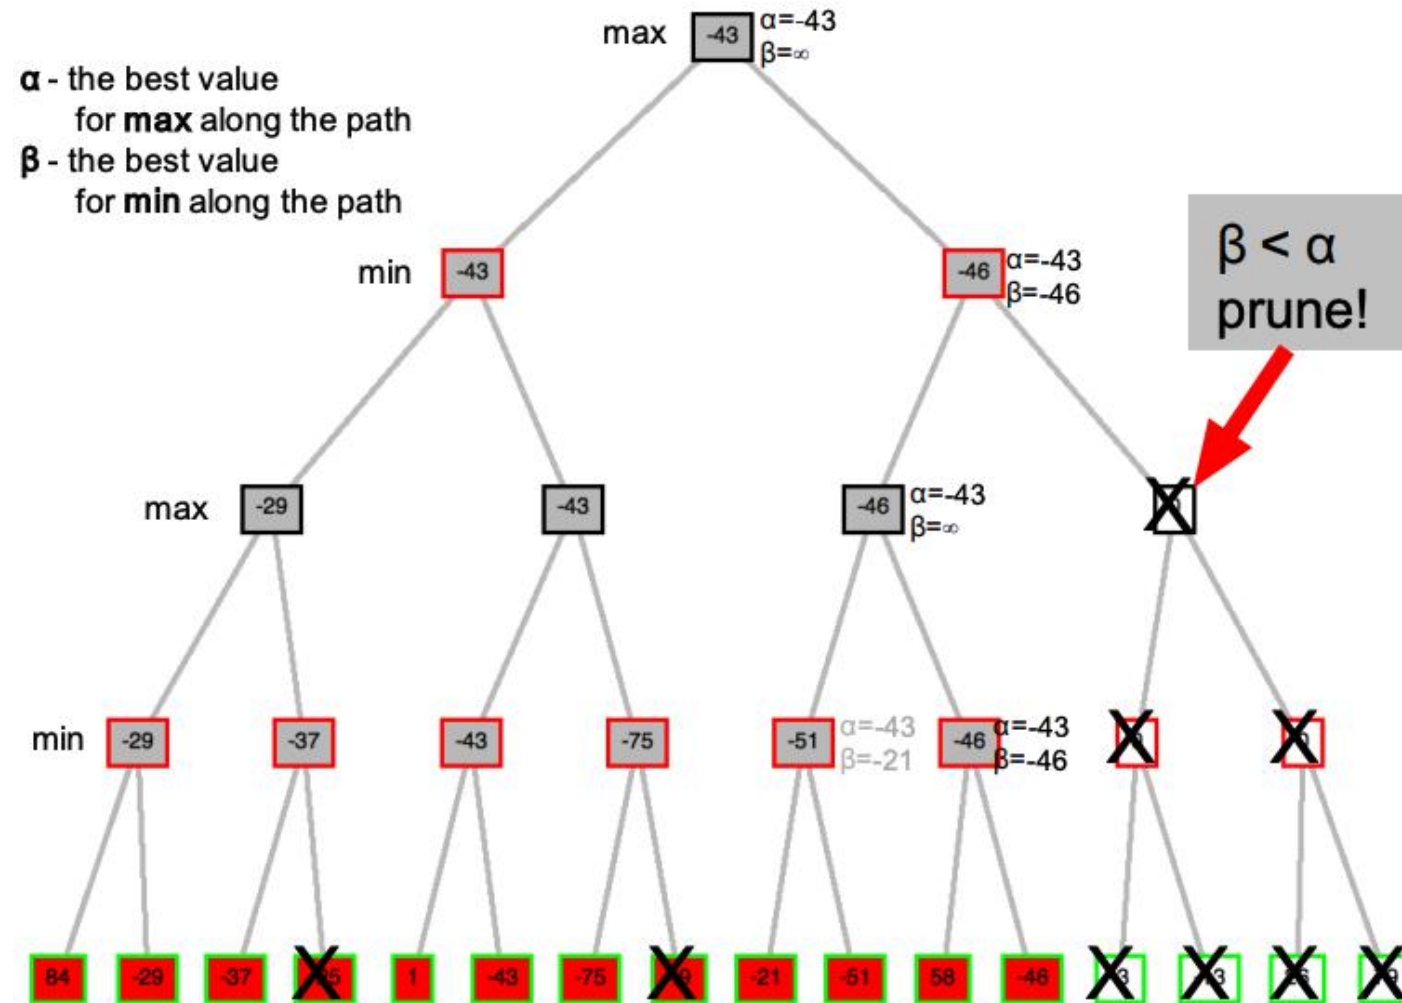
# Another Example



update  $\beta$ ,  
use  $\alpha$  to stop



# Another Example



# $\alpha - \beta$ Search

- $c$  = search cutoff
- $\alpha$  = **lower bound** on Max's outcome; initially set to  $-\infty$
- $\beta$  = **upper bound** on Min's outcome; initially set to  $+\infty$
- We'll call  $\alpha - \beta$  procedure recursively with a **narrowing range** between  $\alpha$  and  $\beta$
- Maximizing levels may reset  $\alpha$  to a **higher** value;
- Minimizing levels may reset  $\beta$  to a **lower** value.

# $\alpha - \beta$ Search Algorithm


$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```


```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

# Alpha-Beta Pruning Properties


- This pruning has **no effect** on minimax value computed for the **root**!
- Values of **intermediate nodes** might be wrong
  - Important: children of the root may have the wrong value
  - So the most naïve version won't let you do **action selection**
- Good child ordering improves **effectiveness** of pruning
- With “perfect ordering”:
  - Time complexity drops to  $O(b^{m/2})$
  - Doubles solvable depth!
  - Full search of, e.g. chess, is still hopeless...



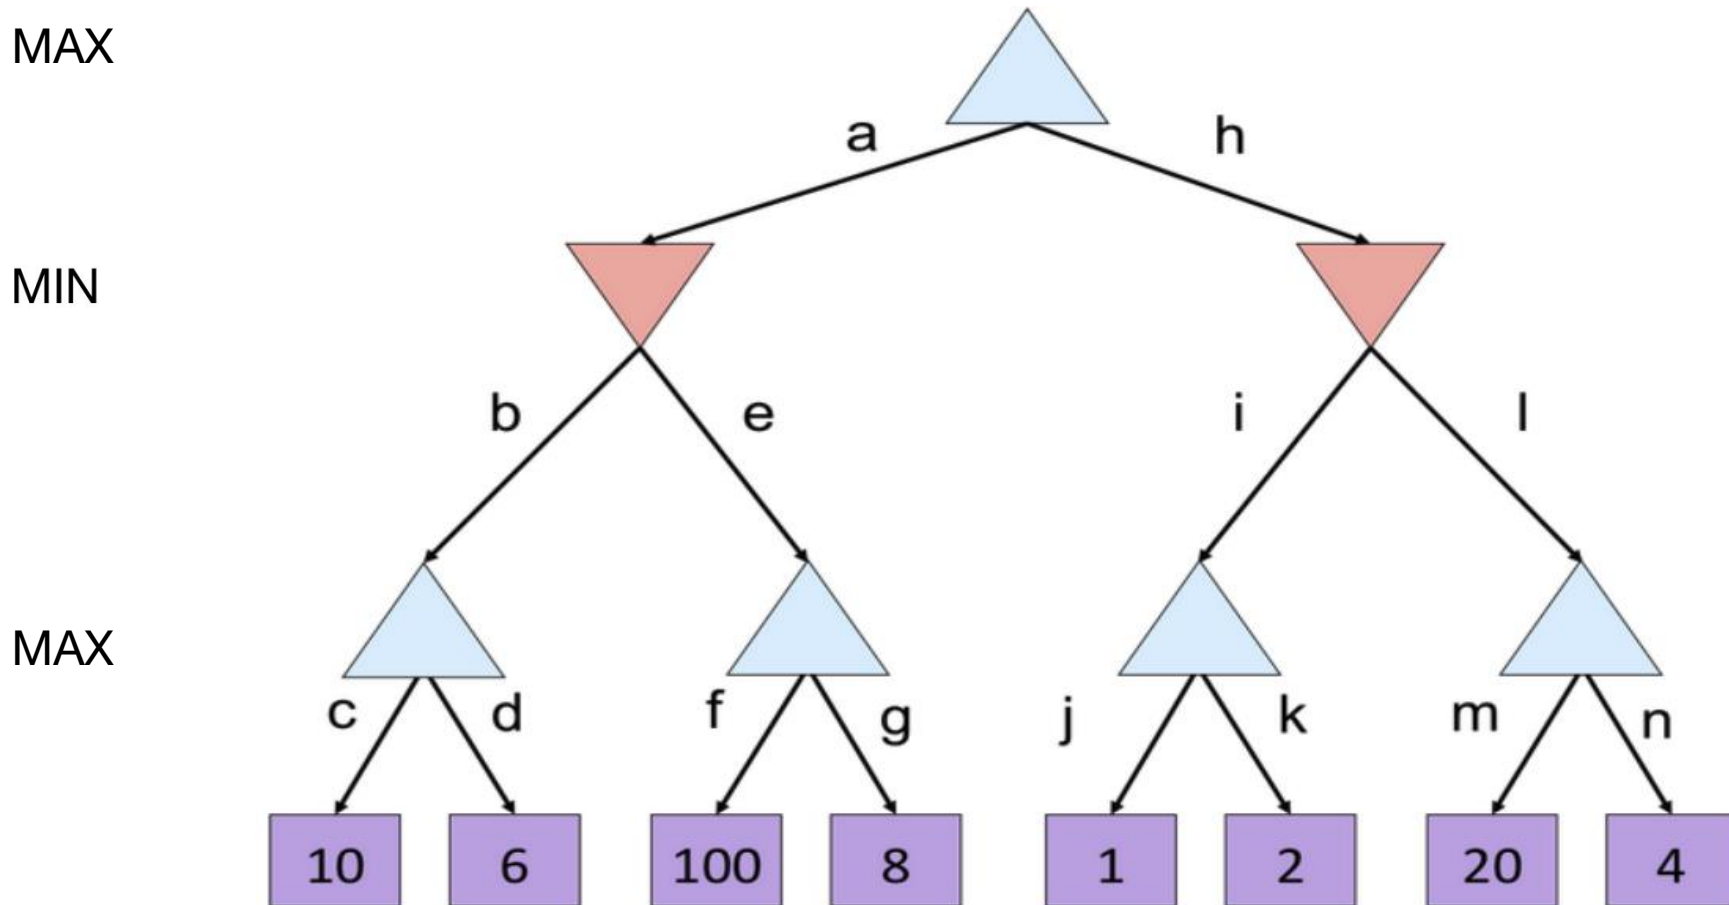
# $\alpha$ - $\beta$ Search Quiz



# $\alpha$ - $\beta$ Search Quiz



# $\alpha$ - $\beta$ Search Quiz



# $\alpha$ - $\beta$ Search Quiz

