

Artificial Intelligence

CS4365 --- Fall 2022

Adversarial Search

Instructor: Yunhui Guo

Game Playing

An AI Favorite

- Structured task
- Not initially thought to require large amounts of knowledge
- Focus on games of perfect information



Game Playing: State-of-the-Art

- **Checkers:** 1950: First computer player. 1994: First computer champion: **Chinook** ended 40-year-reign of human champion Marion Tinsley using complete 8-piece endgame. 2007: Checkers solved!
- **Chess:** 1997: **Deep Blue** defeats human champion Gary Kasparov in a six-game match. Deep Blue examined 200M positions per second, used very sophisticated evaluation and undisclosed methods for extending some lines of search up to 40 ply.
- **Go:** In 2016, **AlphaGo** defeats the human champion




Types of Games

- Many different kinds of games!
 - Stochastic vs. Deterministic
 - One, two or more players
 - Zero sum?
 - Perfect information?
- Want algorithms for calculating a strategy which recommends a move from each state

Deterministic Games

- One possible formulation
 - **States**: S (start at s_0)
 - **Players**: $P = \{1 \dots N\}$ (usually take turns)
 - **Actions**: A (may depend on player / state)
 - **Transition Function**: $S \times A \rightarrow S$
 - **Terminal Test**: $S \rightarrow \{\text{True}, \text{False}\}$
 - **Terminal Utilities**: $S \times P \rightarrow R$
- Solution for a player is a **policy**: $S \rightarrow A$



Zero-Sum Games

- Agents have **opposite** utilities (values on outcomes)
- Lets us think of a single value that one (**MAX**) maximizes and the other minimizes
- Adversarial, pure competition

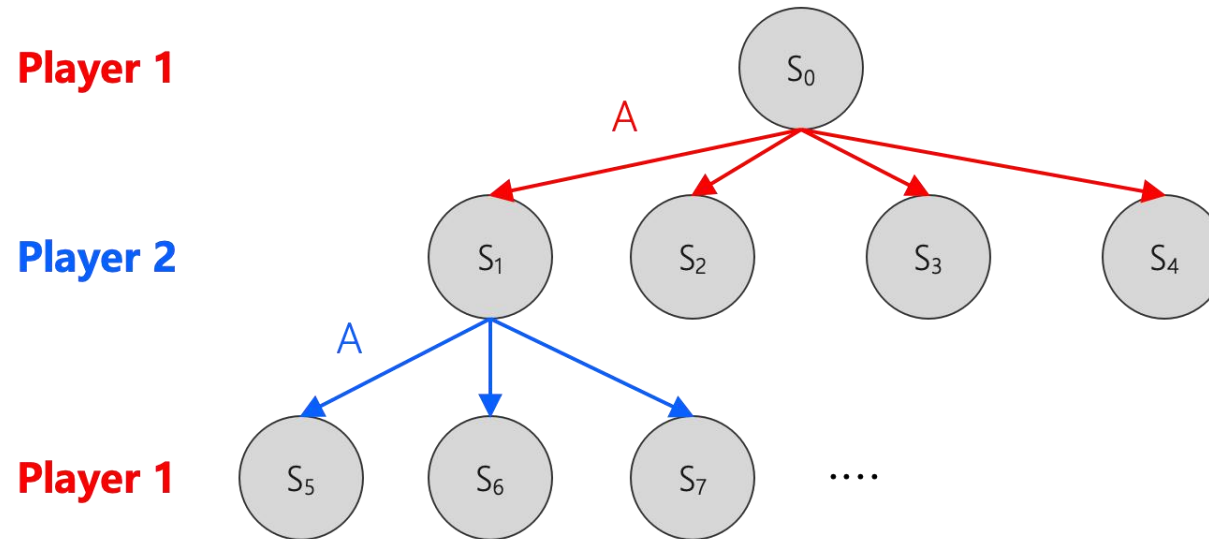


Games VS. Search Problems

- Unpredictable opponent
 - specifying a move for **every possible** opponent reply
- Time limits
 - unlikely to find goal, must approximate


Game Playing as Search

- We can list all the possible **actions** and **states**
- In each step, play 1 searches for an **action** which leads to the maximum **utility**




Search Tree for Tic-Tac-Toe


MAX (x)




Search Tree for Tic-Tac-Toe




Search Tree for Tic-Tac-Toe



Search Tree for Tic-Tac-Toe



Search Tree for Tic-Tac-Toe




Tic-Tac-Toe

- High values are **good** for **MAX** and **bad** for **MIN**. It is MAX's job to use the search tree and utility values to determine the best move.
- Root is initial position. Next level are all moves player 1 (MAX) can make; tree is from Max's viewpoint. Next level are all possible responses from player 2 (MIN).
- Max has to find a strategy that will lead to a winning terminal state **regardless of what Min does**. Strategy has to include the correct move for Max for each possible move by Min.

Pac-Man Trees


- Each step costs 1 point
- The dot worths 10 points



Value of a State

Value of a state:

The best
achievable
outcome (utility)
from that state



Non-Terminal States:

$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

$V(s) = \text{known}$

Adversarial Search Tree


Initial




Pac-Man




Ghost




-20 -8



... -18 -5



... -10 +4



-20 +8


Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:


$$V(s) = \text{known}$$

Minimax Search

- Why do we take the min value every other level of the tree?
- These nodes represent the opponent's choice of move.
- The computer assumes that the human will choose that move that is of least value to the computer.

Minimax Search

- **Deterministic, zero-sum games:**
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- Compute each node's **minimax value**:
 - the best achievable utility against a rational (optimal) adversary




Simplified Minimax Algorithm

1. Expand the **entire tree** below the root
2. Evaluate the **terminal nodes** as wins for the minimizer or maximimizer
3. Select an unlabeled node, n , all of whose children have been assigned values. If there is no such node, we're done — return the value assigned to the root.
4. If n is a **minimizer** move, assign it a value that is **the minimum of the values of its children**. If n is a **maximizer** move, assign it a value that is the **maximum of the values of its children**. Return to Step 3.

Another Example

Max

Min



Summary


- In game tree search, **a move is a pair of actions**. one player's action is a ply. 2-ply = one move.
- Called a minimax decision because it maximizes the utility under the assumption that the opponent will play perfectly to minimize it.
- **Time complexity:**
 - $O(b^m)$ (m plies and b branching.) Impractical for e.g. chess ($b \approx 30$ to 40). 1000^k for k moves.
- **Space complexity:** $O(bm)$

Size of Search Space

- Chess:
 - branching factor $b \approx 35$
 - game length $m \approx 100$
 - search space $b^m = 35^{100} \approx 10^{154}$
- The Universe
 - number of atoms $\approx 10^{78}$
- Exact search is infeasible

The Need for Imperfect Decisions

- Problem:
 - **Minimax** assumes the program has time to search to the terminal nodes.
- In realistic games, cannot search to leaves!
- Solution: Cut off search earlier and apply a heuristic evaluation function to the leaves
- Guarantee of optimal play is gone



Static Evaluation Functions

- **Minimax** depends on the translation of board quality into a single, summarizing number. Difficult. Expensive
- **Evaluation functions** score non-terminals in depth-limited search
 - Do you control the center of the board?
 - How well protected is your king?
 - Add up values of pieces each player has (weighted by importance of piece).
 - Mobility
- Strategies change as the game proceeds.

Design Issues for Heuristic Minimax

Evaluation Function:


What **features** should we evaluate and how should we use them?

An evaluation function should:

1. Match utility function on terminal states
2. Not take too long
3. Accurately reflect the chance of winning



Evaluation Functions for Tic-Tac-Toe

- Let p be a position in the game
- Define the utility function $f(p)$ by
 - count the number of lines where X can win
 - subtract number of lines where O can win




Evaluation Functions for Tic-Tac-Toe

- $f(p)$ = the number of lines where X can win - the number of lines where O can win



$$8-8 = 0$$



$$8-4 = 4$$



$$6-4 = 2$$



$$6-2 = 4$$

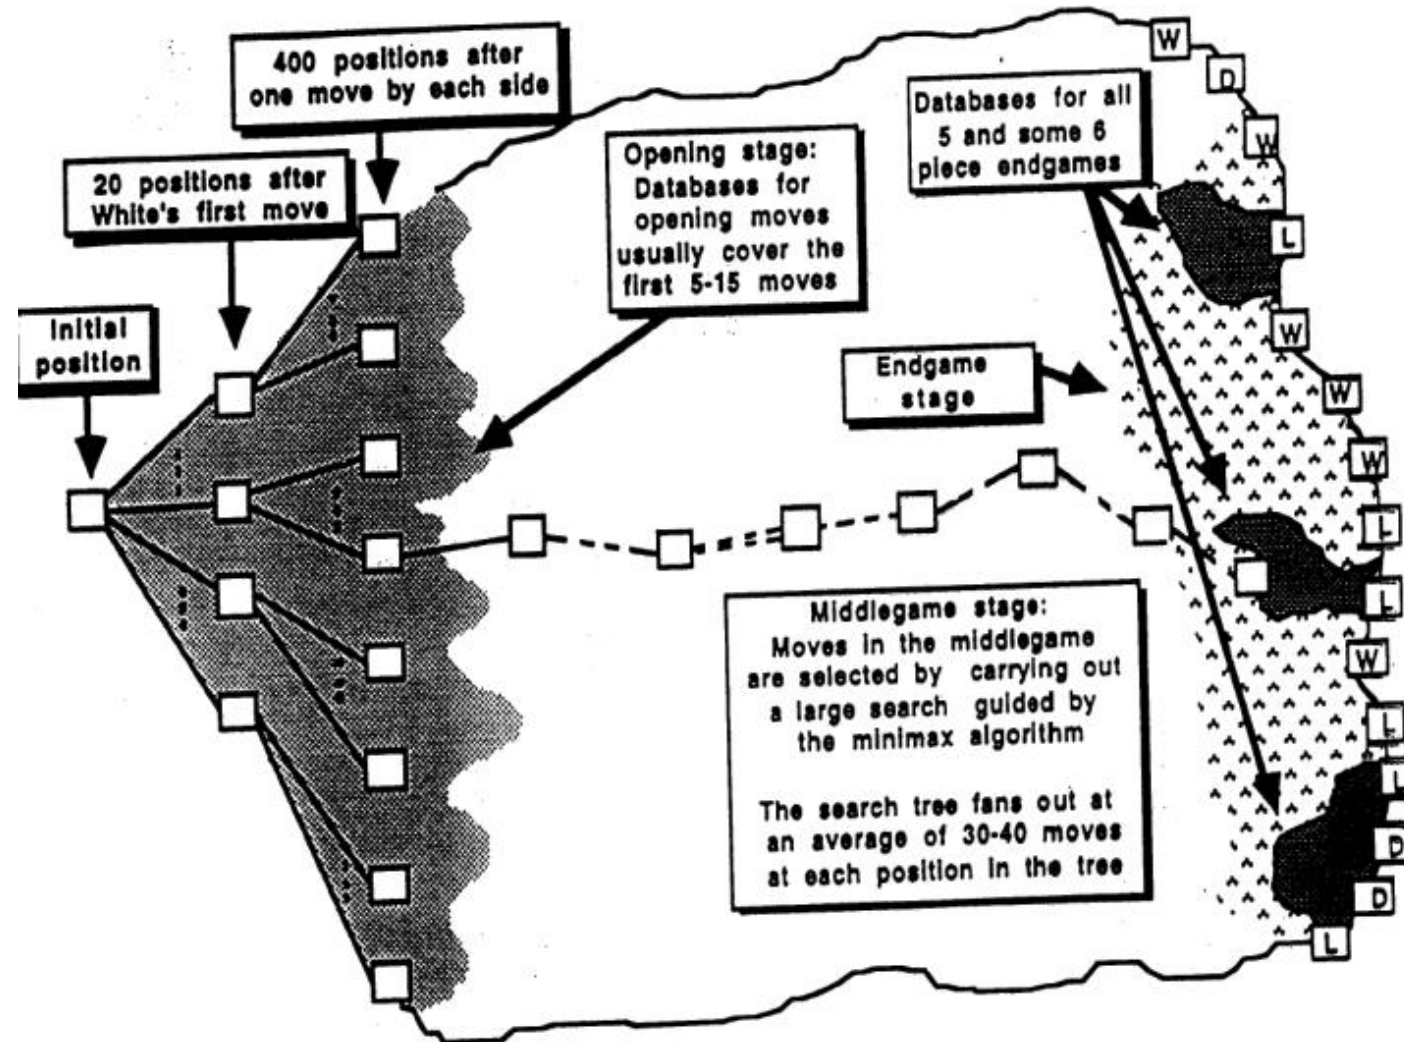
Linear Evaluation Functions

- Let f_i be **features** and w_i be weights
- Linear evaluation function: $w_1f_1 + w_2f_2 + \dots + w_nf_n$
This is what most game playing programs use
- For example: $f = 6 \cdot \text{material} + 4 \cdot \text{mobility} + \text{center control}$

Linear Evaluation Functions

- Let f_i be **features** and w_i be weights
- Linear evaluation function: $w_1f_1 + w_2f_2 + \dots + w_nf_n$
This is what most game playing programs use
- Steps in designing an evaluation function:
 1. Pick informative features
 2. Find the weights that make the program play well
- Deep Blue used ~6,000 different features!

Minimax Search



Design Issues for Heuristic Minimax


- Depth-limited search:
 - search to a constant depth
- Problems:
 - Some portions of the game tree may be **less stable** than the others
 - Horizon effect

Unstable States

- Unstable state: drastic change from one level to the next
 - A chess evaluation function that counts material gains may evaluate a given state poorly even the play can capture a queen in the next move.
 - Are you about to lose an important piece?
- Evaluation functions can only be trusted when applied to **stable board states**
- One solution is to extend the normal search to look for stable states

The Horizon Effect


- You may incorrectly estimate the value of a state by overlooking an event that is just **beyond the depth limit**
 - Opponent moves, move is significant damage and cannot not be avoided
- Fixed-depth searches can be **mislead** by the fact that these damaging moves can be delayed
 - The damage is beyond the search horizon and so is not seen



The Horizon Effect

- The **negative** horizon effect
 - MAX may try to avoid a bad situation which is actually **inevitable**.
- For example, MAX tries to avoid losing the queen and appears to be able to do so using a lookahead tree of depth 6, but a little deeper it becomes obvious that the queen is going to be lost.

Piece	Value
Pawn	1
Knight	3
Bishop	3
Rook	5
Queen	9



The Horizon Effect

The **positive** horizon effect:

MAX may not realize that something good is going to be achievable.

For example, MAX would like to take MIN's queen and that can happen
- but the restricted horizon prevents MAX from making the right choices to realize this possibility