

Artificial Intelligence

CS4365 --- Fall 2022

Constraint Satisfaction Problem

Instructor: Yunhui Guo

Standard Search Problem

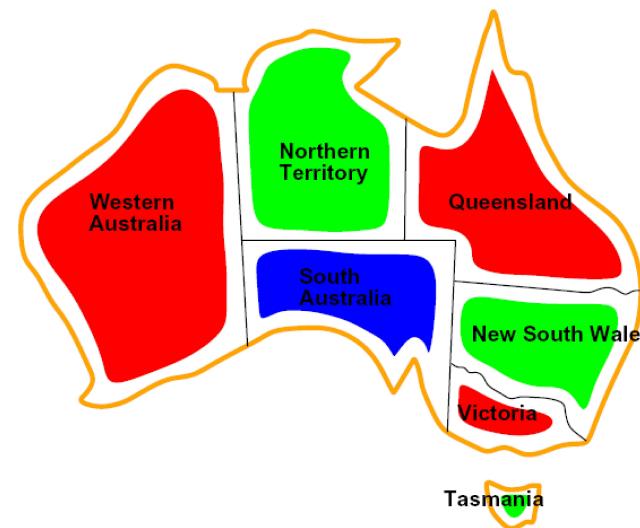
- Standard Search Problems:
 - State is a “black box”: arbitrary data structure
 - Goal test can be any function over states
 - Successor function can also be anything

Constraint Satisfaction Problems (CSP)

- Constraint satisfaction problems (CSPs):
 - A **special subset** of search problems
 - **State** is defined by variable X_i with values from a domain D (sometimes D depends on i)
 - Goal test is **a set of constraints** specifying allowable combinations of values for subsets of variables
- Why CSPs?
 - Allows useful **general-purpose algorithms** with more power than standard search algorithms

Example: Map-Coloring Problem

- Color each region either **red**, **green** or **blue** in such a way that no neighboring regions have the same color



- Variables: different regions
- Domains: colors

Constraint Satisfaction Problems (CSP)

- A powerful representation for (discrete) search problems.
- A **Constraint Satisfaction Problem** (CSP) is defined by:
 - **X** is a set of n variables X_1, X_2, \dots, X_n , each defined by a finite domain D_1, D_2, \dots, D_n of possible values
 - **C** is a set of constraints C_1, C_2, \dots, C_m . Each C_i involves some subset of the variables; specifies the allowable combinations of values for that subset.

A **solution** is an assignment of values to the variables that satisfies all constraints.

Cryptarithmetic as a CSP

- **Variables:**

FOUR TW C₁ C₂ C₃ TWO

- Domains:

{0, 1, 2, 3, 4, 5, 6, 7, 8, 9}

- Constraints:

O + O = R + 10 * C₁ FOUR

$$C_1 + W + W = U + 10 * C_2$$

$$C_2 + T + T = O + 10^* C_3$$

$$C_3 = F$$

alldiff(F, O, U, R, T, W)

Cryptarithmetic Constraint Graph

- Variables:

F O U R T W C₁ C₂ C₃

- Domains:

{0, 1, 2, 3, 4, 5, 6, 7, 8, 9}

- Constraints:

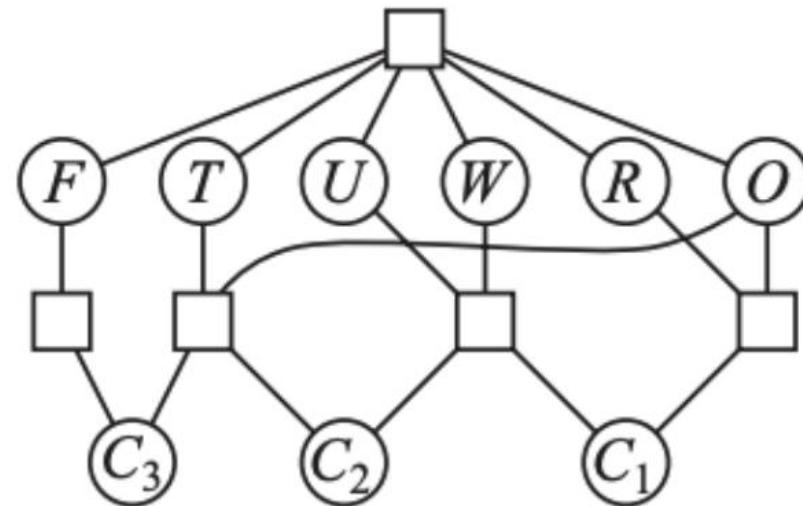
$$O + O = R + 10 * C_1$$

$$C_1 + W + W = U + 10 * C_2$$

$$C_2 + T + T = O + 10 * C_3$$

$$C_3 = F$$

$$\text{alldiff}(F, O, U, R, T, W)$$



Map-Coloring Problem

- **Variables:**

- WA, NT, Q, NSW, V, SA, T

- **Domains:**

- $\{\text{red}, \text{green}, \text{blue}\}$

- **Constraints:**

- Adjacent regions must have different colors

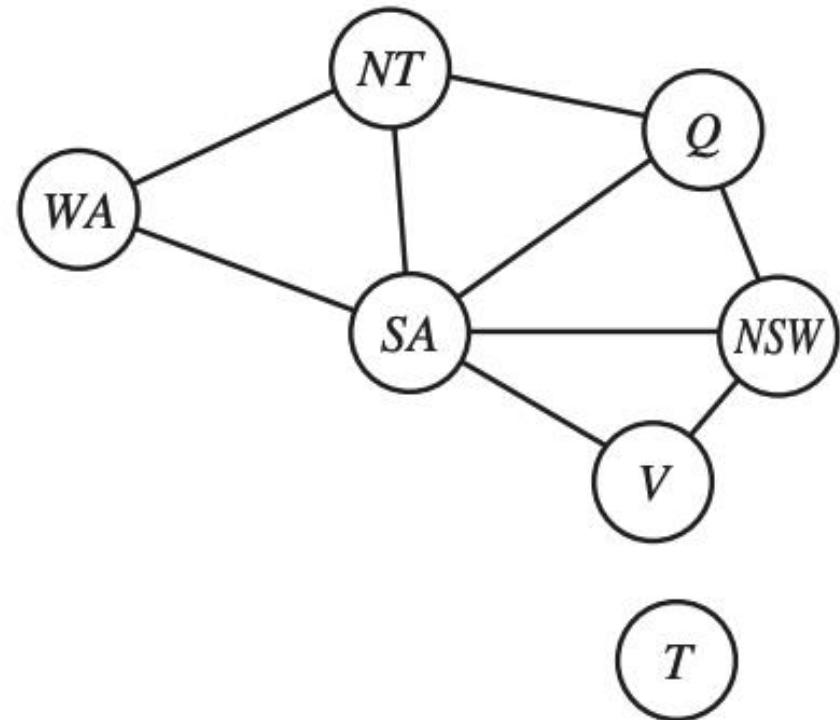
- Implicit: $WA \neq NT$

- Explicit: $(WA, NT) \in \{(\text{red}, \text{green}), (\text{red}, \text{blue}), \dots\}$



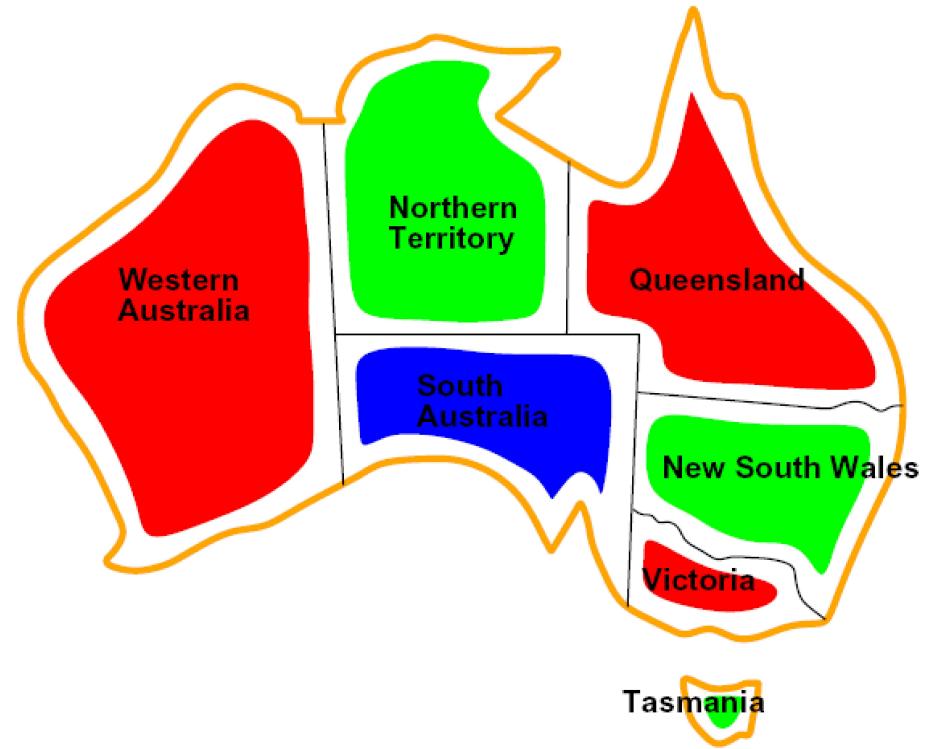
Constraint Graph

- Variables:
WA, NT, Q, NSW, V, SA, T
- Domains:
 $\{\text{red, green, blue}\}$
- Constraints:
Adjacent regions must have different colors
Implicit: WA \neq NT
Explicit: $(WA, NT) \in \{(\text{red, green}), \{\text{red, blue}\}, \dots\}$



Map-Coloring Problem

- Variables:
WA, NT, Q, NSW, V, SA, T
- Domains:
{red, green, blue}
- Constraints:
Adjacent regions must have different colors
Implicit: WA \neq NT
Explicit: (WA, NT) \in {(red, green), {red, blue}, ...}



Varieties of CSPs

- Discrete variables
 - Finite domains: each variable can be only assigned to finite number of values.
 - Infinite domains: each variable can be assigned to infinite number of values, e.g., set of integers
- Continuous variables
 - Linear programming

$$\begin{aligned} & \text{Minimize} && z = 0.6x_1 + 0.35x_2 \\ & \text{subject to:} && \\ & && 5x_1 + 7x_2 \geq 8 \\ & && 4x_1 + 2x_2 \geq 15 \\ & && 2x_1 + x_2 \geq 3 \\ & && x_1 \geq 0, x_2 \geq 0. \end{aligned}$$

Varieties of Constraints

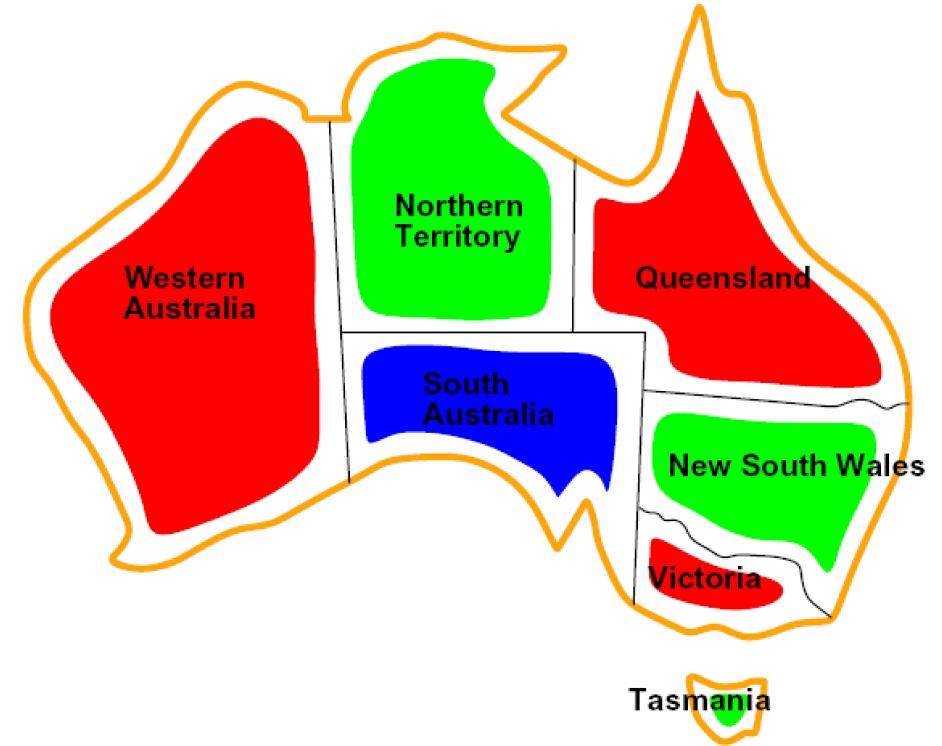
- Unary constraints involve a single variable (equivalent to reducing domains), e.g.:
 - SA ≠ Green
- Binary constraints involve pairs of variables, e.g.:
 - SA ≠ WA
- Higher-order constraints involve 3 or more variables
- Preferences (soft constraints):
 - E.g., red is better than green

Real-World CSPs

- Assignment problems
 - e.g., who teaches what class
- Timetabling problems
 - e.g., which class is offered when and where?
- Transportation scheduling
- Factory scheduling

Constraint Satisfaction Problems (CSP)

- For a given CSP the problem is one of the following:
 - Find all solutions
 - Find one solution
 - The optimal solution given an objective function
 - Determine if a solution exists



How to View a CSP as a Search Problem?

- Initial State -- state in which all the variables are unassigned.
- Successor function -- assign a value to a variable from a set
- Goal test -- check if all the variables are assigned and **all the constraints are satisfied**. We call the assignment is complete.
- Path cost -- assumes constant cost for each step

Backtracking Search

- Depth-first search for CPSs with single-variable assignments is called backtracking search
 - Backtracks when a variable has no legal values left to assign

Backtracking search is the basic uninformed algorithm for CSPs.

The Backtracking Search Algorithm

Backtrack(assignment, csp):

 If assignment is complete then return assignment

 var \leftarrow select one unassigned variable

 for each value in the domain of var do

 if value is consistent with assignment then

 add {var = value} to assignment

 result \leftarrow Backtrack(assignment, csp)

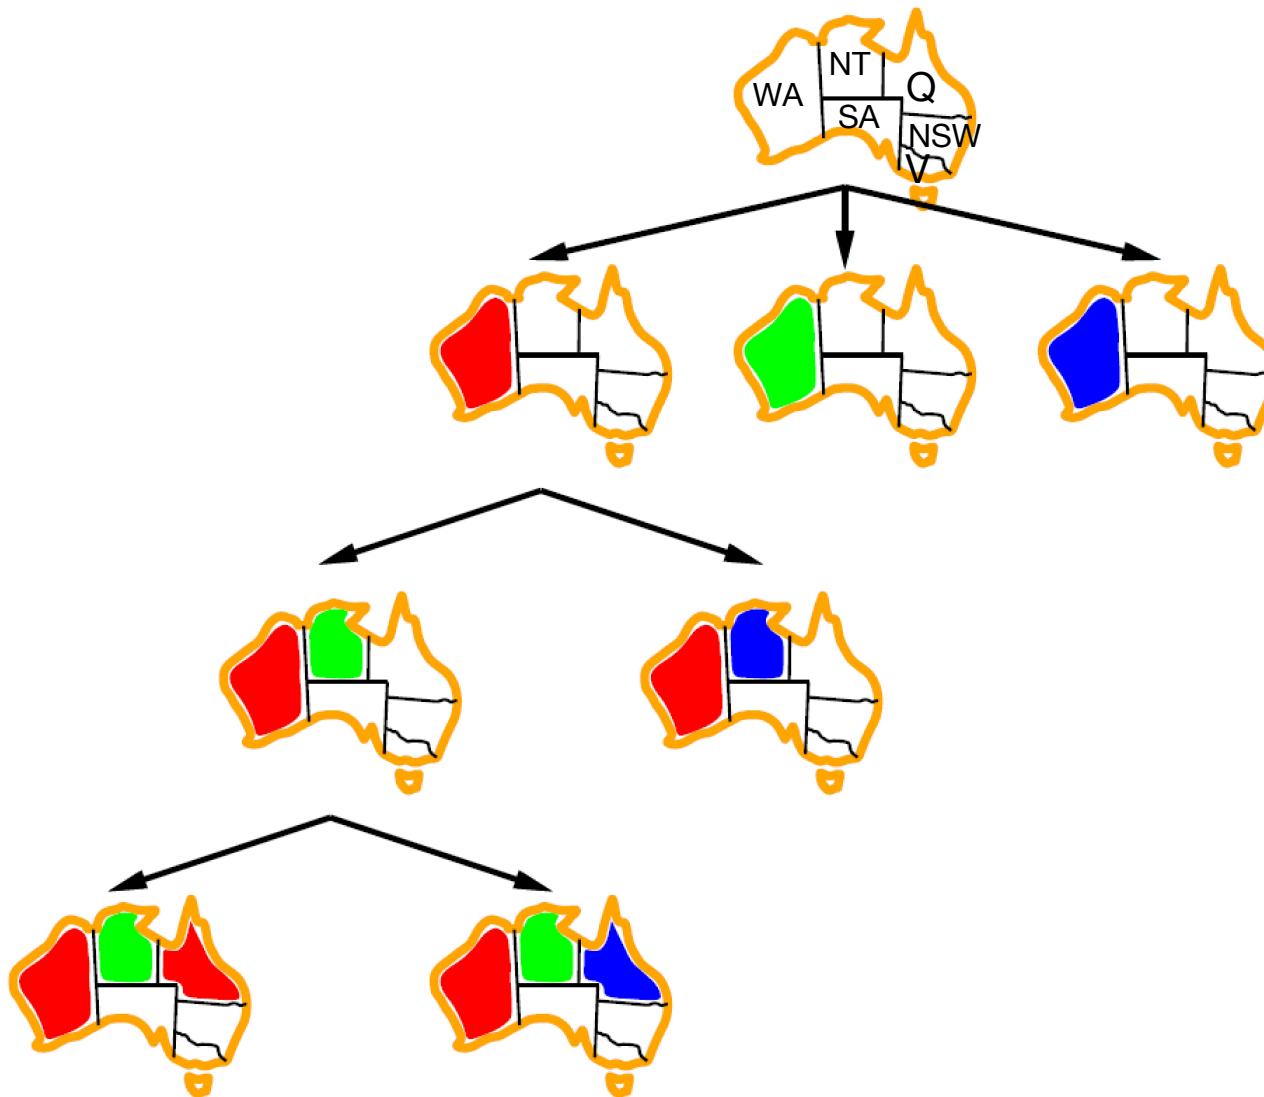
 if result \neq failure **then**

return result

 remove {var = value} from assignment

return failure

Backtracking Example



The Backtracking Search Algorithm

Backtrack(assignment, csp):

 If assignment is complete then return assignment

 var \leftarrow select one unassigned variable \leftarrow which one to select?

 for each value in the domain of var do \leftarrow which one to use?

 if value is consistent with assignment then

 add {var = value} to assignment \leftarrow detect failure early

 result \leftarrow Backtrack(assignment, csp)

 if result \neq failure **then**

return result

 remove {var = value} from assignment

return failure

Improving Backtracking Efficiency

General-purpose methods can give huge gains in speed:

1. Which variable should be assigned next?

- Choose the most constrained variable, and break ties by choosing the most constraining variable

2. In what order should its values be tried?

- Choose the least constraining value

3. Can we detect inevitable failure early?

- Forward checking, constraint propagation

Improving Backtracking Efficiency

General-purpose methods can give huge gains in speed:

1. Which variable should be assigned next?

- Choose the most constrained variable, and break ties by choosing the most constraining variable

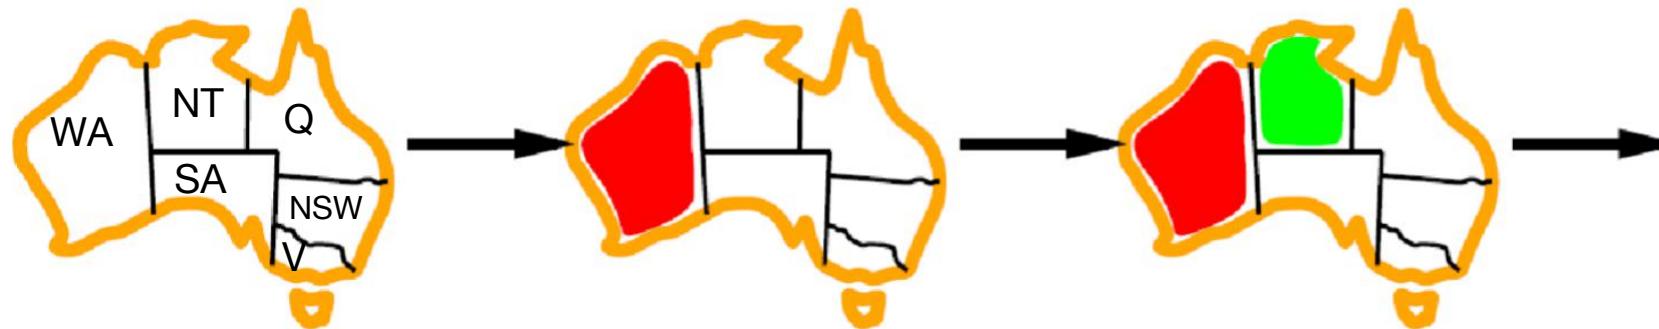
2. In what order should its values be tried?

- Choose the least constraining value

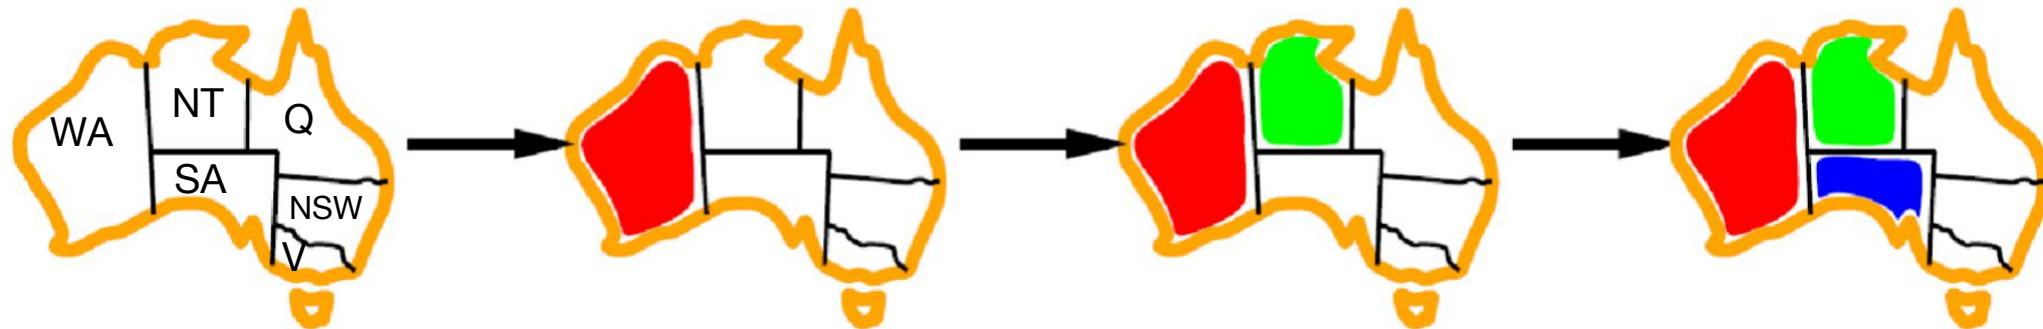
3. Can we detect inevitable failure early?

- Forward checking, constraint propagation

Variable Selection Heuristic 1: Most Constrained Variable



Variable Selection Heuristic 1: Most Constrained Variable



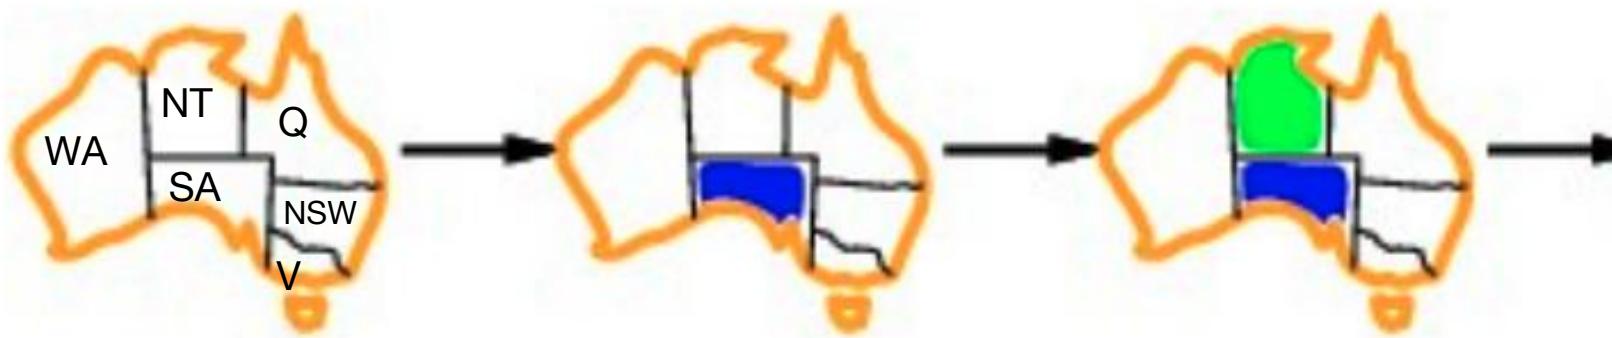
Most constrained variable:

Choose the variable with the fewest legal values

a.k.a. **minimum remaining values (MRV)** heuristic

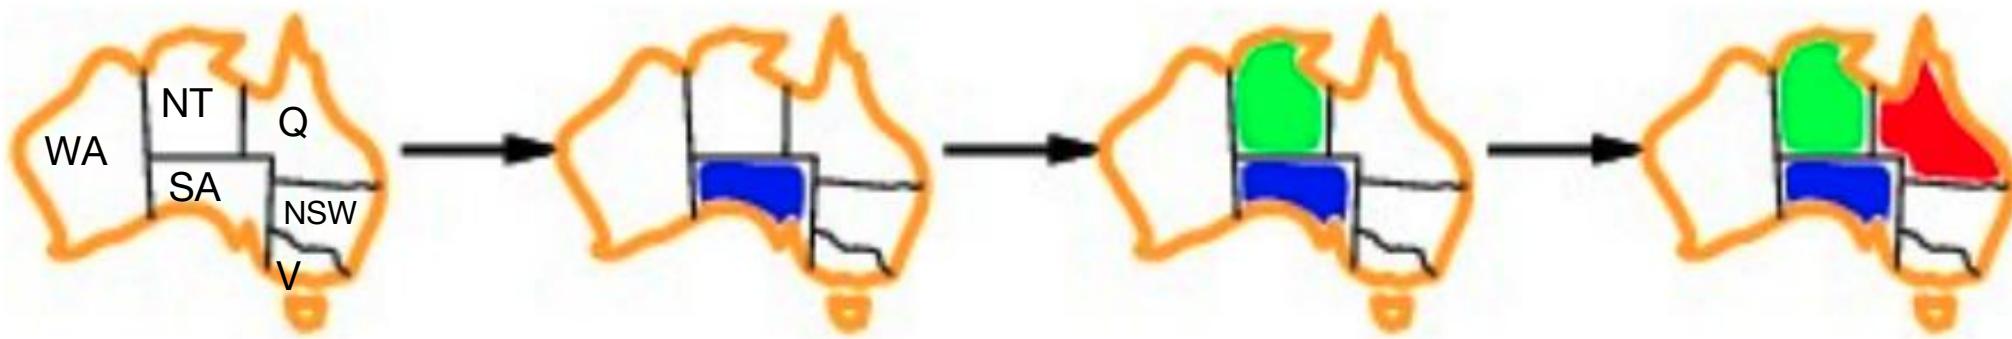
- Why min rather than max?
- “Fail-fast” ordering
- $1000 \times$ faster or more

Variable Selection Heuristic 2: Most Constraining Variable



- Tie-breaker among most constrained variables

Variable Selection Heuristic 2: Most Constraining Variable



- Tie-breaker among most constrained variables

Most constraining variable:

choose the variable with the most constraints on remaining variables

Improving Backtracking Efficiency

General-purpose methods can give huge gains in speed:

1. Which variable should be assigned next?

- Choose the most constrained variable, and break ties by choosing the most constraining variable

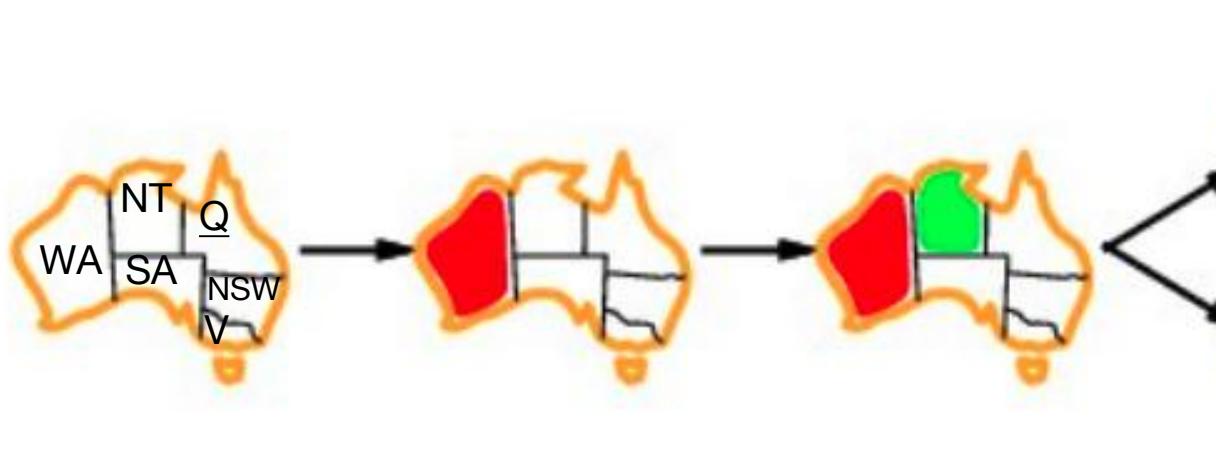
2. In what order should its values be tried?

- Choose the least constraining value

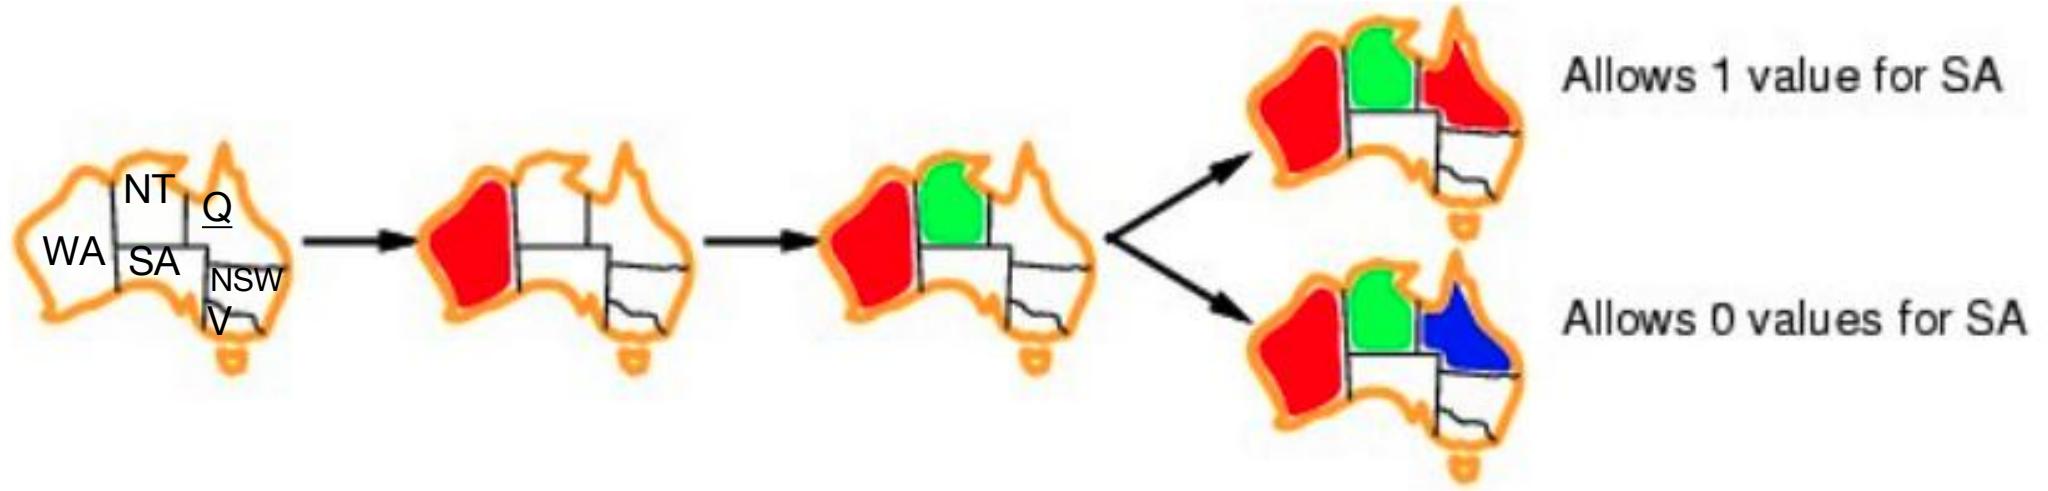
3. Can we detect inevitable failure early?

- Forward checking, constraint propagation

Value Select Heuristic: Least Constraining Value



Value Select Heuristic: Least Constraining Value



- Given a variable, choose the least constraining value:
the one that rules out the fewest values in the remaining variables
- Why least rather than most?

Variable and Value Ordering Heuristics

- The heuristics should be applied whenever a variable/value needs be chosen during the search process
- They should not be applied once and for all at the begining of the search process

Improving Backtracking Efficiency

General-purpose methods can give huge gains in speed:

1. Which variable should be assigned next?

- Choose the most constrained variable, and break ties by choosing the most constraining variable

2. In what order should its values be tried?

- Choose the least constraining value

3. Can we detect inevitable failure early?

- Forward checking, constraint propagation

Consistency-Enforcing Procedure 1: Forward Checking

- Idea:
 - Keep track of **remaining legal values** for **unsigned variables**
 - Terminate search when any variable has no legal values



WA

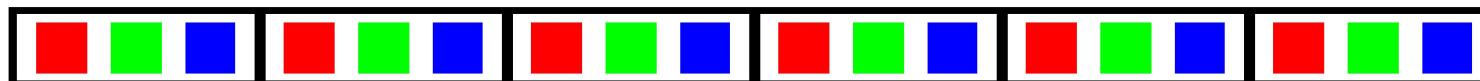
NT

Q

NSW

V

SA



Consistency-Enforcing Procedure 2: Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:



WA	NT	Q	NSW	V	SA
Red	Green	Blue	Red	Green	Blue
Red		Green	Blue	Red	Green
Red		Blue	Green	Red	Blue

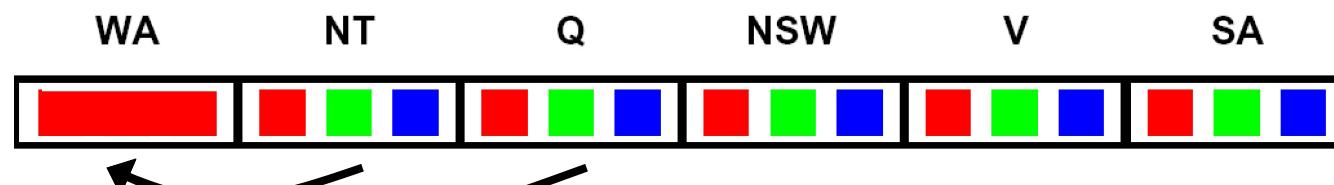
- NT and SA cannot both be blue!
- Constraint propagation repeatedly enforces constraints locally.

Consistency-Enforcing Procedure 2: Constraint Propagation

- Using the **constraints** to reduce the number of **legal values** for a variable
- Constraint propagation may be intertwined with search, or it may be done as a preprocessing step, before search starts
- Sometimes this preprocessing can solve the whole problem, so no search is required at all!
- Local consistency

Arc Consistency

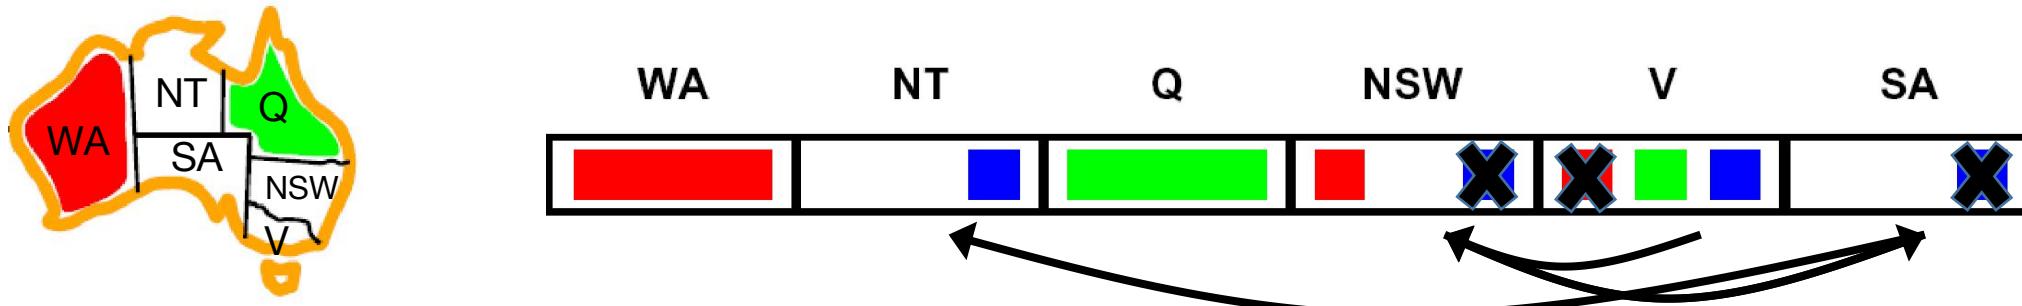
- An arc $X \rightarrow Y$ is **consistent** iff for every x in the tail there is some y in the head which could be assigned without violating a constraint



- Forward checking: Enforcing consistency of arcs pointing to each new assignment

Arc Consistency

- An arc $X \rightarrow Y$ is **consistent** iff for every x in the tail there is some y in the head which could be assigned without violating a constraint



- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency detects failure earlier than forward checking
- Can be run as a preprocessor or after each assignment

Arc Consistency

function AC-3(*csp*) **returns** the CSP, possibly with reduced domains

inputs: *csp*, a binary CSP with variables $\{X_1, X_2, \dots, X_n\}$

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REMOVE-INCONSISTENT-VALUES(X_i, X_j) **then**

for each X_k **in** NEIGHBORS[X_i] **do**

 add (X_k, X_i) to *queue*

function REMOVE-INCONSISTENT-VALUES(X_i, X_j) **returns** true iff succeeds

removed \leftarrow false

for each x **in** DOMAIN[X_i] **do**

if no value y in DOMAIN[X_j] allows (x, y) to satisfy the constraint $X_i \leftrightarrow X_j$

then delete x from DOMAIN[X_i]; *removed* \leftarrow true

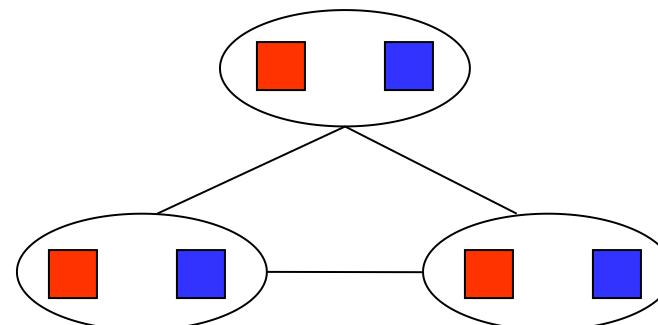
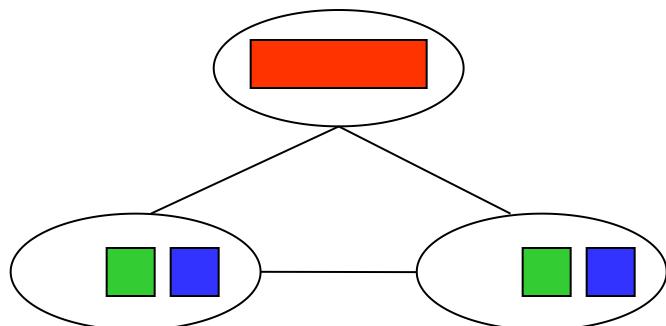
return *removed*

Runtime: $O(n^2d^3)$

n : # of variables d : domain size

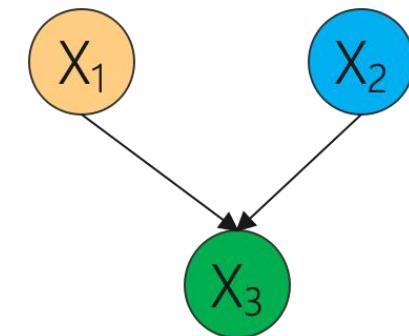
Limitations of Arc Consistency

- After enforcing arc consistency:
 - Can have one solution left
 - Can have multiple solutions left
 - Can have no solutions left (and not know it)



K-Consistency

- Increasing degrees of consistency
 - 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints
 - 2-Consistency (Arc Consistency): For each pair of nodes, any consistent assignment to one can be extended to the other
 - K-Consistency: For each k nodes, any consistent assignment to k-1 can be extended to the kth node.



Consistency-Enforcing Procedures

- Forward checking
 - Constraint propagation
 - Variable/value ordering heuristics
-
- They can used in combination to improve search process.

General Purpose CSP Algorithm

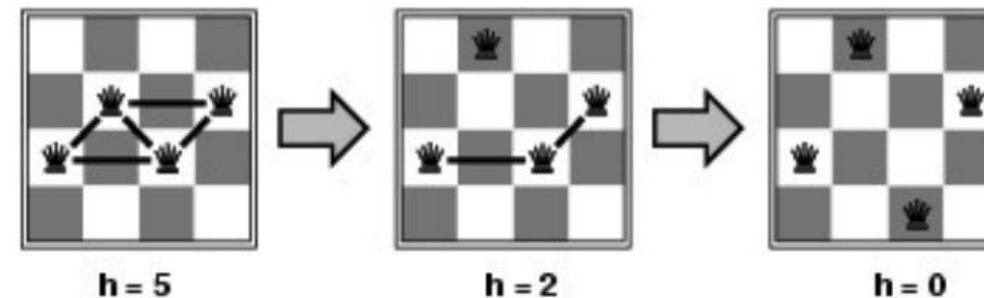
1. If all values have been successfully assigned then stop, else go on to 2.
2. Apply the consistency enforcing procedure (e.g. forward checking if feeling computationally mean, or constraint propagation if extravagant.)
3. If a deadend is detected then backtrack (e.g., DFS-type backtrack.). Else go on to step 4.
4. Select the next variable to be assigned (using your variable ordering heuristic)
5. Select a promising value (using your value ordering heuristic)
6. Create a new search state. Cache the info you need for backtracking. And go back to 1.

Local Search for CSPs

- Hill-climbing, simulated annealing typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - allow states with unsatisfied constraints operators **reassign** variable values
 - Variable selection: randomly select any conflicted variable
 - Value selection by **min-conflicts** heuristic:
 - Choose value that violates the fewest constraints
 - i.e., hill-climb with $h(n)$ = total number of violated constraints

Example: 4-Queens

- States: 4 queens in 4 columns (256 states)
- Actions: move queen in column
- Goal test: no attacks
- Evaluation: $h(n) = \text{number of attacks}$



- Give random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)

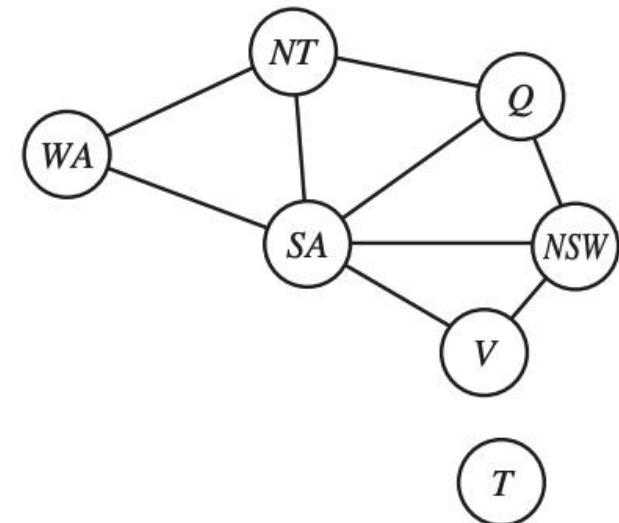
Problem Encoding

- How can the problem structure help to find a solution quickly?
- Decomposition of the original problem:

$$CSP = \bigcup_i CSP_i$$

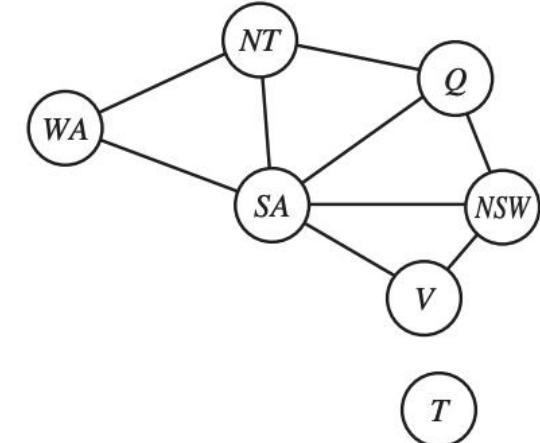
- Solutions can be combined to solve the original problem

$$S = \bigcup_i S_i$$

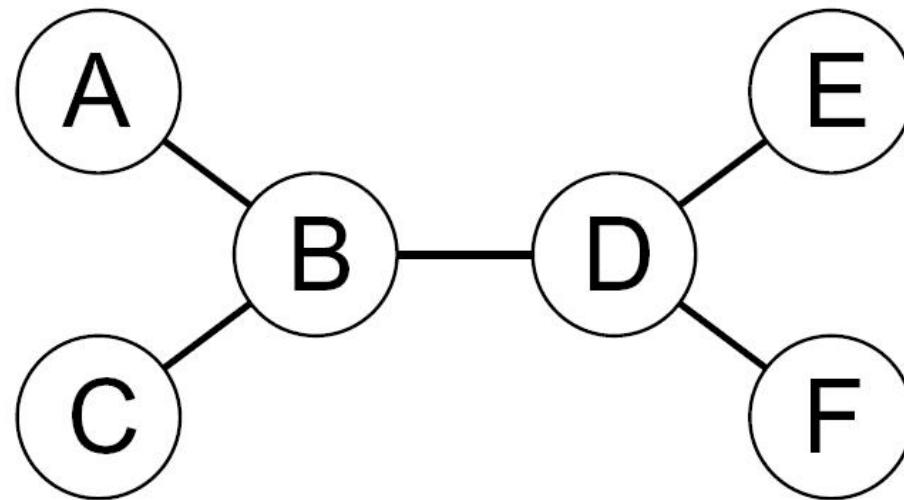


Problem Encoding

- Extreme case: **independent** subproblems
 - Example: Tasmania and mainland do not interact
- **Independent subproblems** are identifiable as connected components of constraint graph
- Suppose a graph of n variables can be broken into subproblems of only c variables:
 - Worst-case solution cost is $O((n/c)(d^c))$, linear in n
 - E.g., $n = 80$, $d = 2$, $c = 20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec



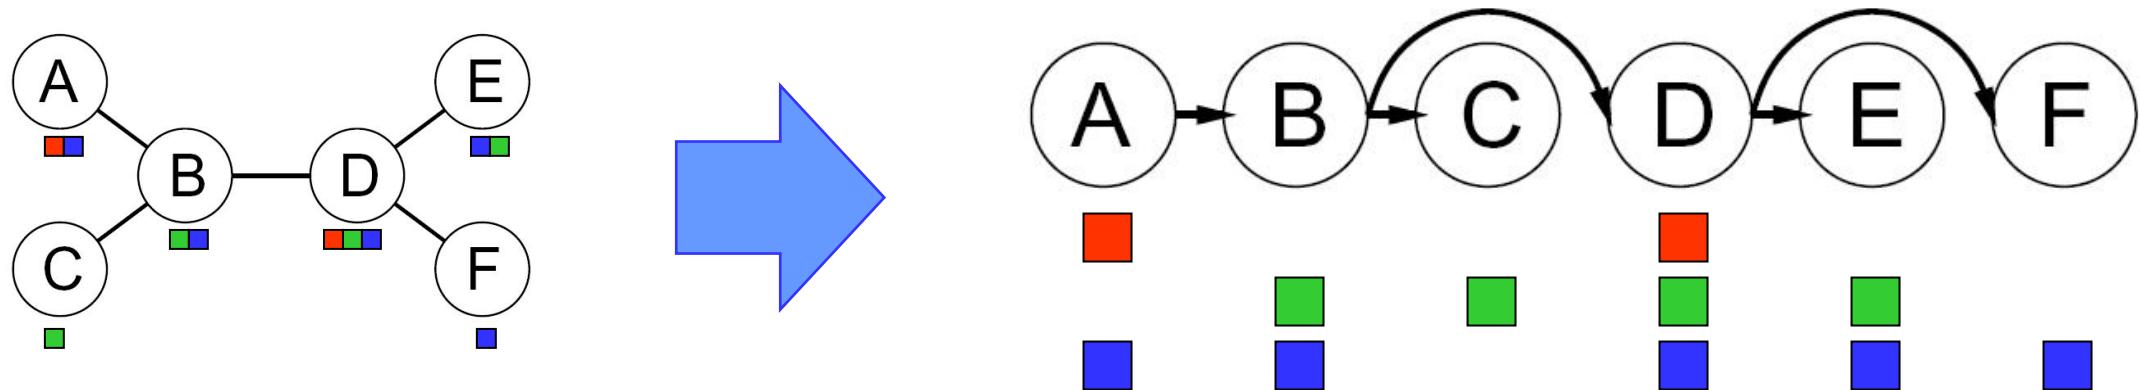
Tree-Structured CSPs



- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(nd^2)$ time
- Compare to general CSPs, where worst-case time is $O(d^n)$

Tree-Structured CSPs

- Algorithm for tree-structured CSPs:
 - Order: Choose a root variable, order variables so that parents precede children



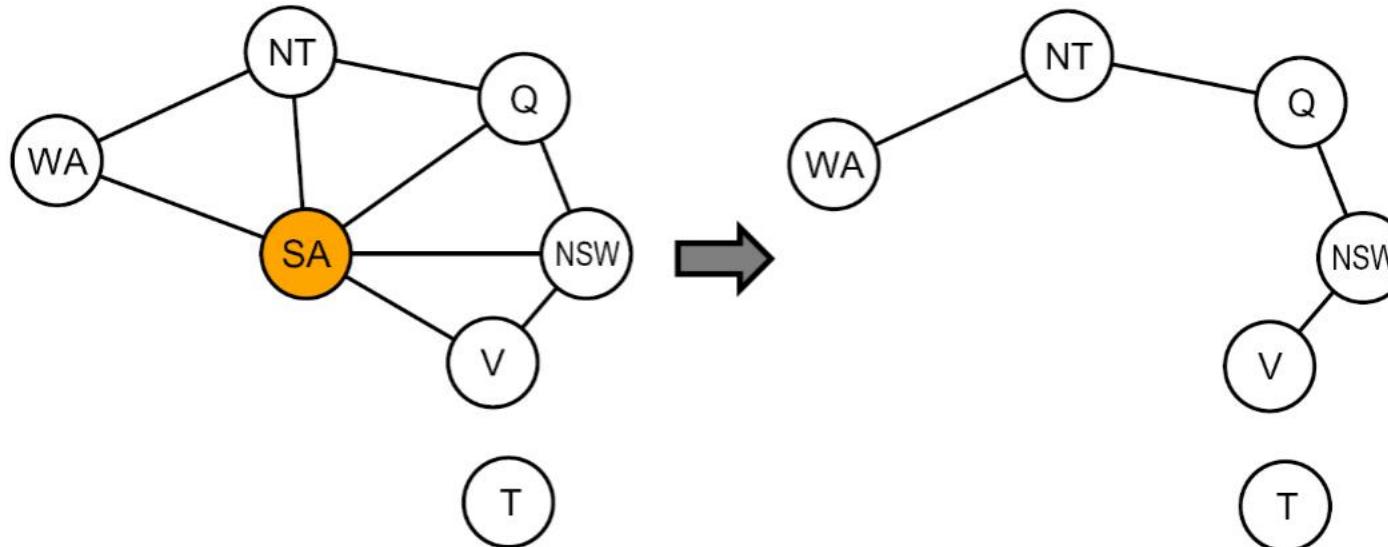
- Remove backward: For $i = n : 2$, apply **RemoveInconsistent**($\text{Parent}(X_i)$, X_i)
- Assign forward: For $i = 1 : n$, assign X_i consistently with $\text{Parent}(X_i)$

General CSPs

- Removing nodes:
 - cutset conditioning: remaining variables form a tree
- Collapsing nodes
 - Tree decomposition: construct a set of connected subproblems.

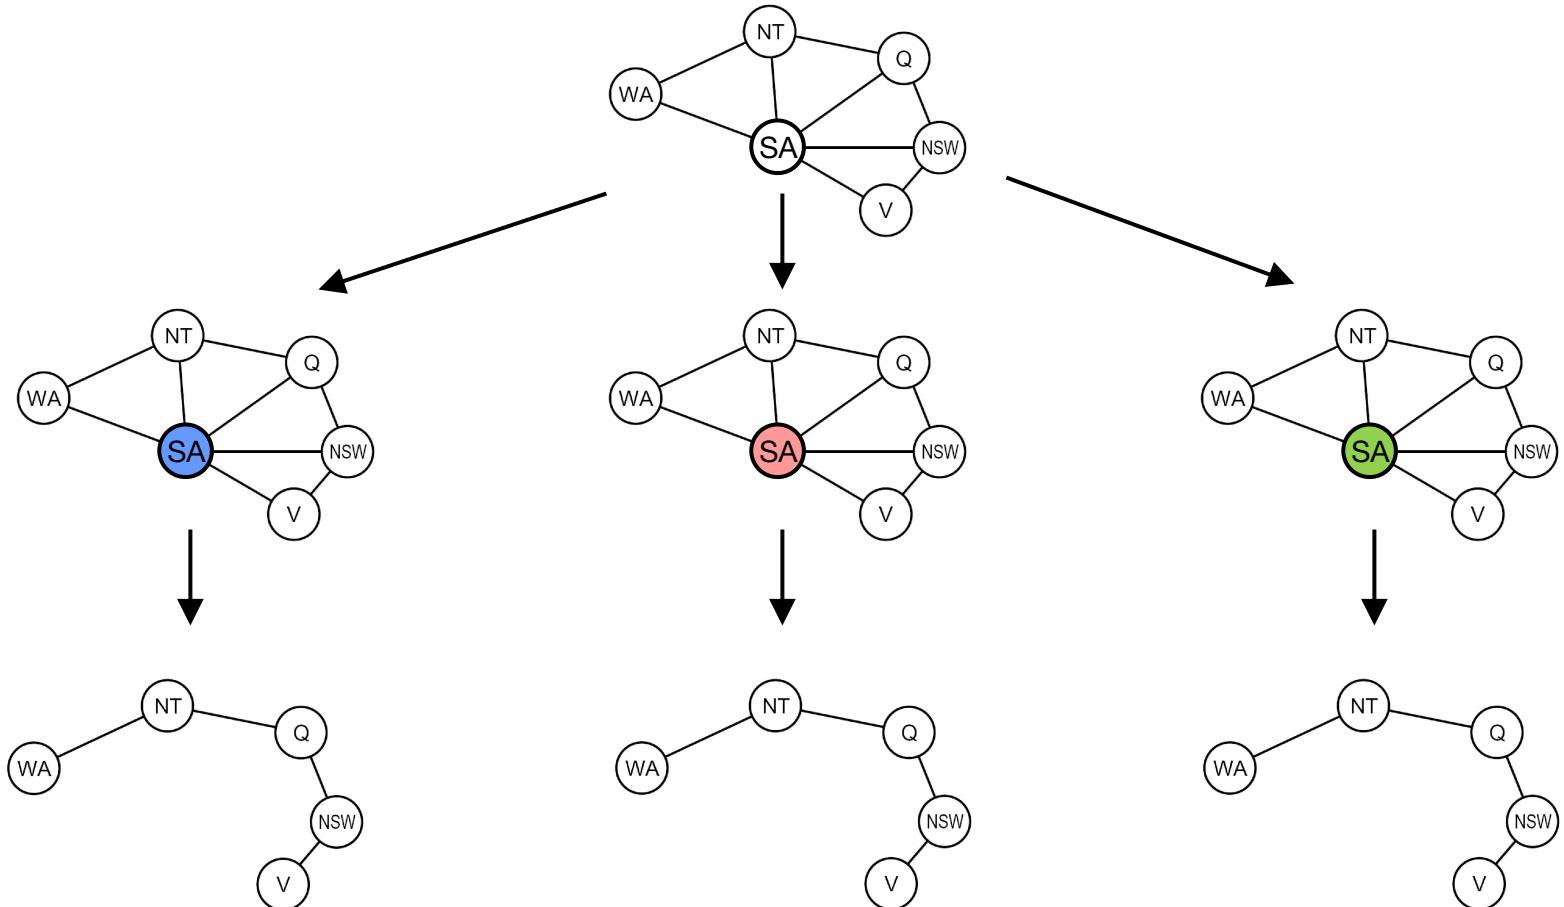
Nearly Tree-Structured CSPs

- **Conditioning:** instantiate a variable, prune its neighbors' domains



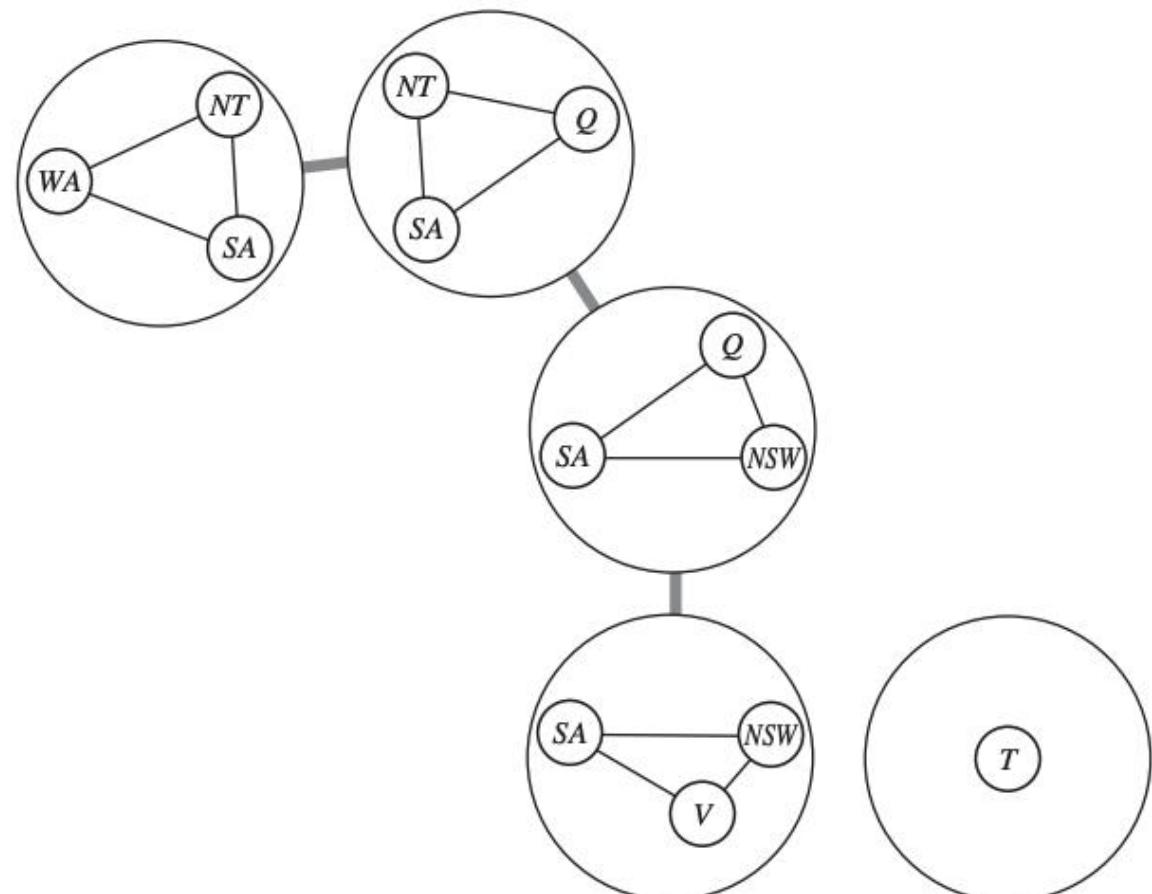
Nearly Tree-Structured CSPs

- **Cutset conditioning:**
 - instantiate (in all ways) a set of variables S such that the remaining constraint graph is a tree
 - Solve the remaining constraint graph and return it with the assignment of S .
 - If the cutset has size c , then the run time is $O(d^c \cdot (n-c)d^2)$



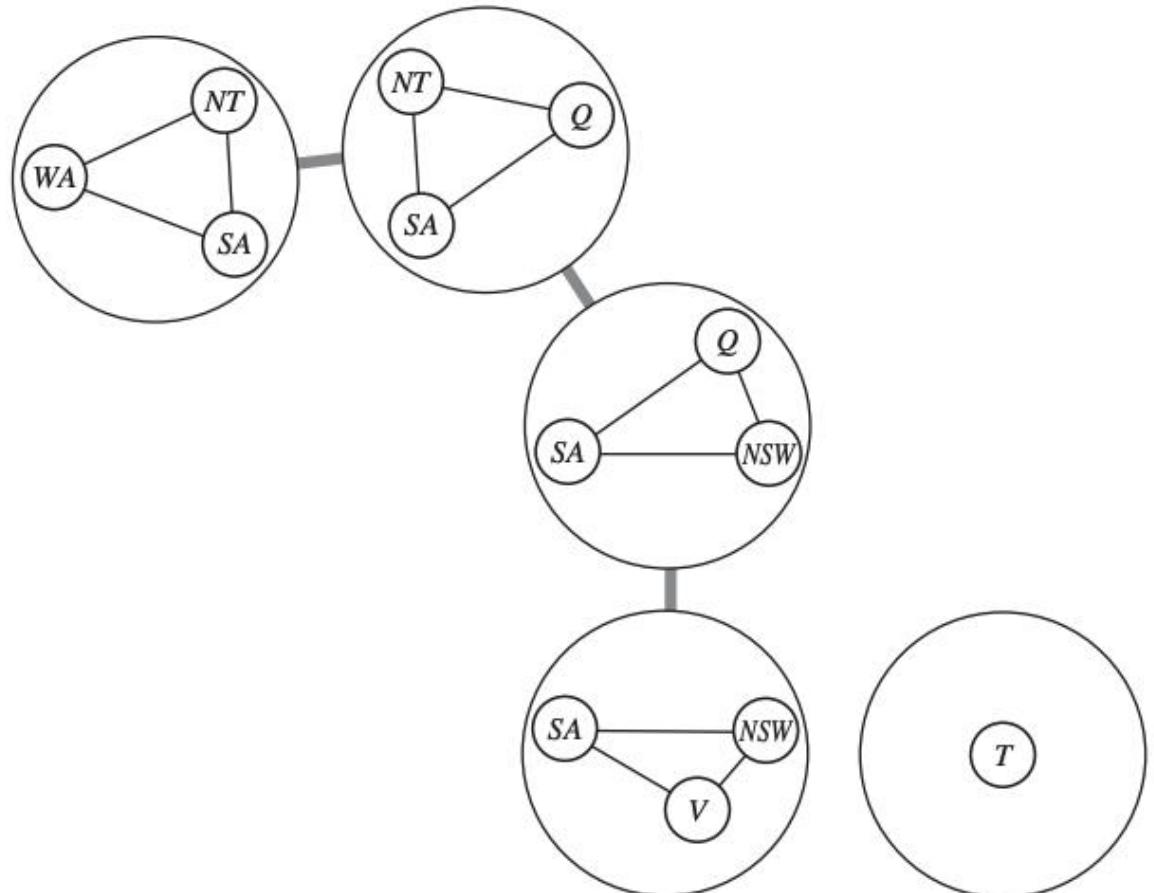
Tree Decompositon

- Every variable must appear in **at least one subproblem**
- If two variables are connected by a constraint in the original problem, they must appear together (along with the constraint) in at least one of the subproblems.
- If a variable occurs in two subproblems in the tree, it must appear in every subproblem on the path that connects the two



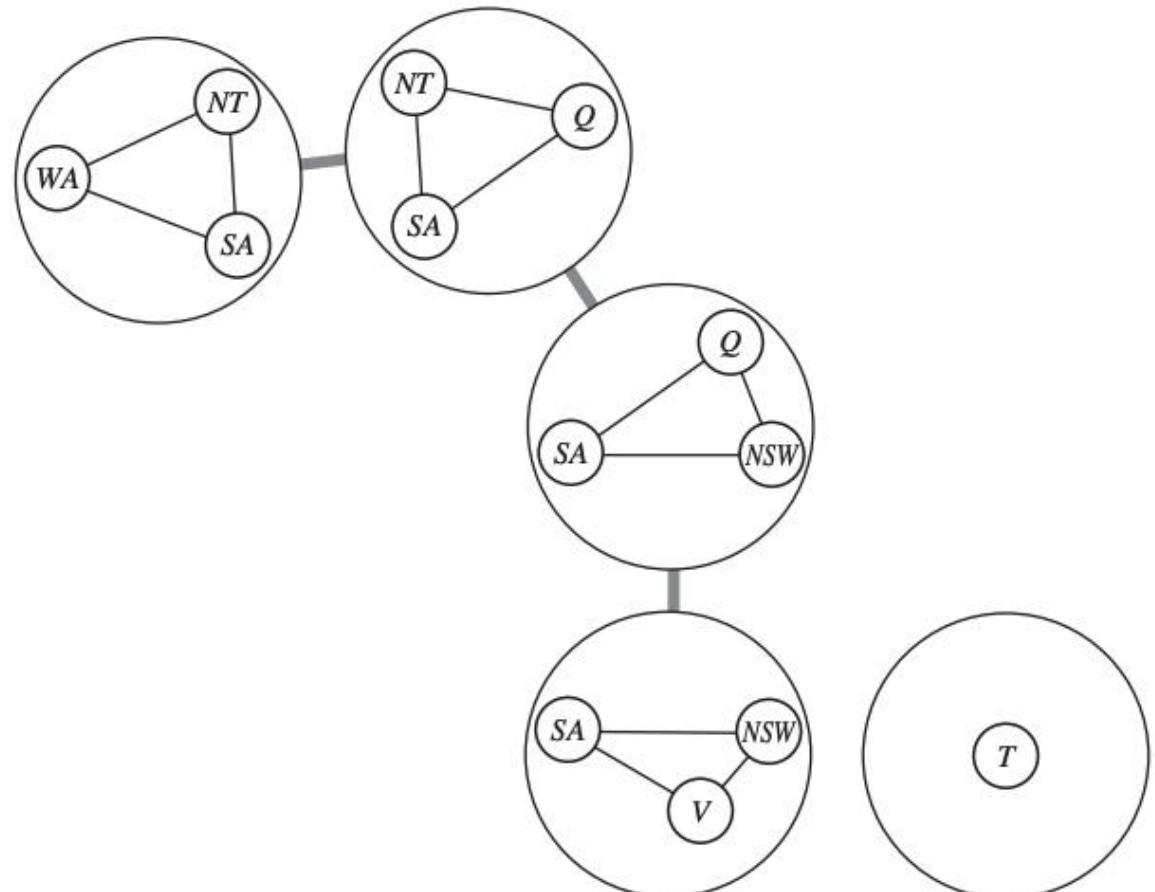
Tree Width

- The subproblem should be as small as possible..
- Tree width w of a tree decomposition is the size of **largest** sub-problem minus 1
- Tree width of a graph is minimal tree width over all possible tree decompositions (hard to find)



Tree Decompositon

- Solve for all solutions of each subproblem.
- Use tree-structured algorithm, treating the subproblem solutions as variables for those subproblems.
- $O(nd^{w+1})$ where w is the tree width of the graph



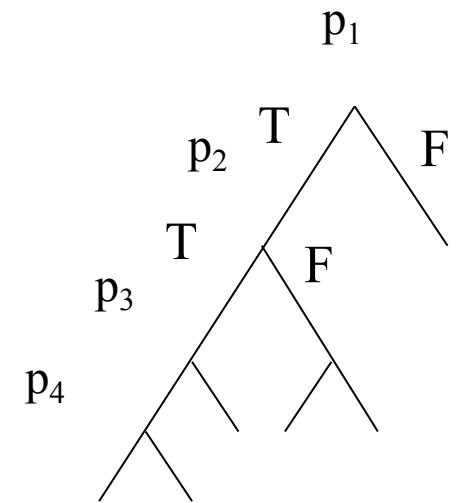
Application: Satisfiability

- 3SAT:
 - find assignments s.t. $(p_1 \vee p_3 \vee p_4) \wedge (\neg p_1 \vee p_2 \vee \neg p_3)$
 $\wedge \dots$ is satisfied
- DFS, GSAT, random walk SAT, ...

Davis-Putnam-Logemann-Loveland (DPLL) tree search algorithm

$$(p_1 \vee p_3 \vee p_4) \wedge (\neg p_1 \vee p_2 \vee \neg p_3) \wedge \dots$$

- Backtrack when some clause becomes empty



- **Unit propagation** (for variable & value ordering):

- if some clause only has one literal left, assign that variable the value that satisfies the clause (never need to check the other branch)

Variants of CSPs

- Dynamic CSPs:
 - The constraints are evolved.
 - Can be considered as a sequence of static CSPs.
- Flexible CSPs:
 - Not all constraints need to be satisfied.
- Distributed CSPs:
 - The variables are distributed among multiple agents.

Summary

- CSPs are a special kind of problem:
 - **states** defined by values of a fixed set of variables
 - **goal test** defined by constraints on variable values
- Backtracking = DFS with one variable assigned per node
- Variable ordering and value selection heuristics help significantly
- Forward checking prevents assignments that guarantee later failure
- Constraint propagation (e.g., arc consistency) does additional work to constrain values and detect inconsistencies
- Iterative min-conflicts is usually effective in practice