

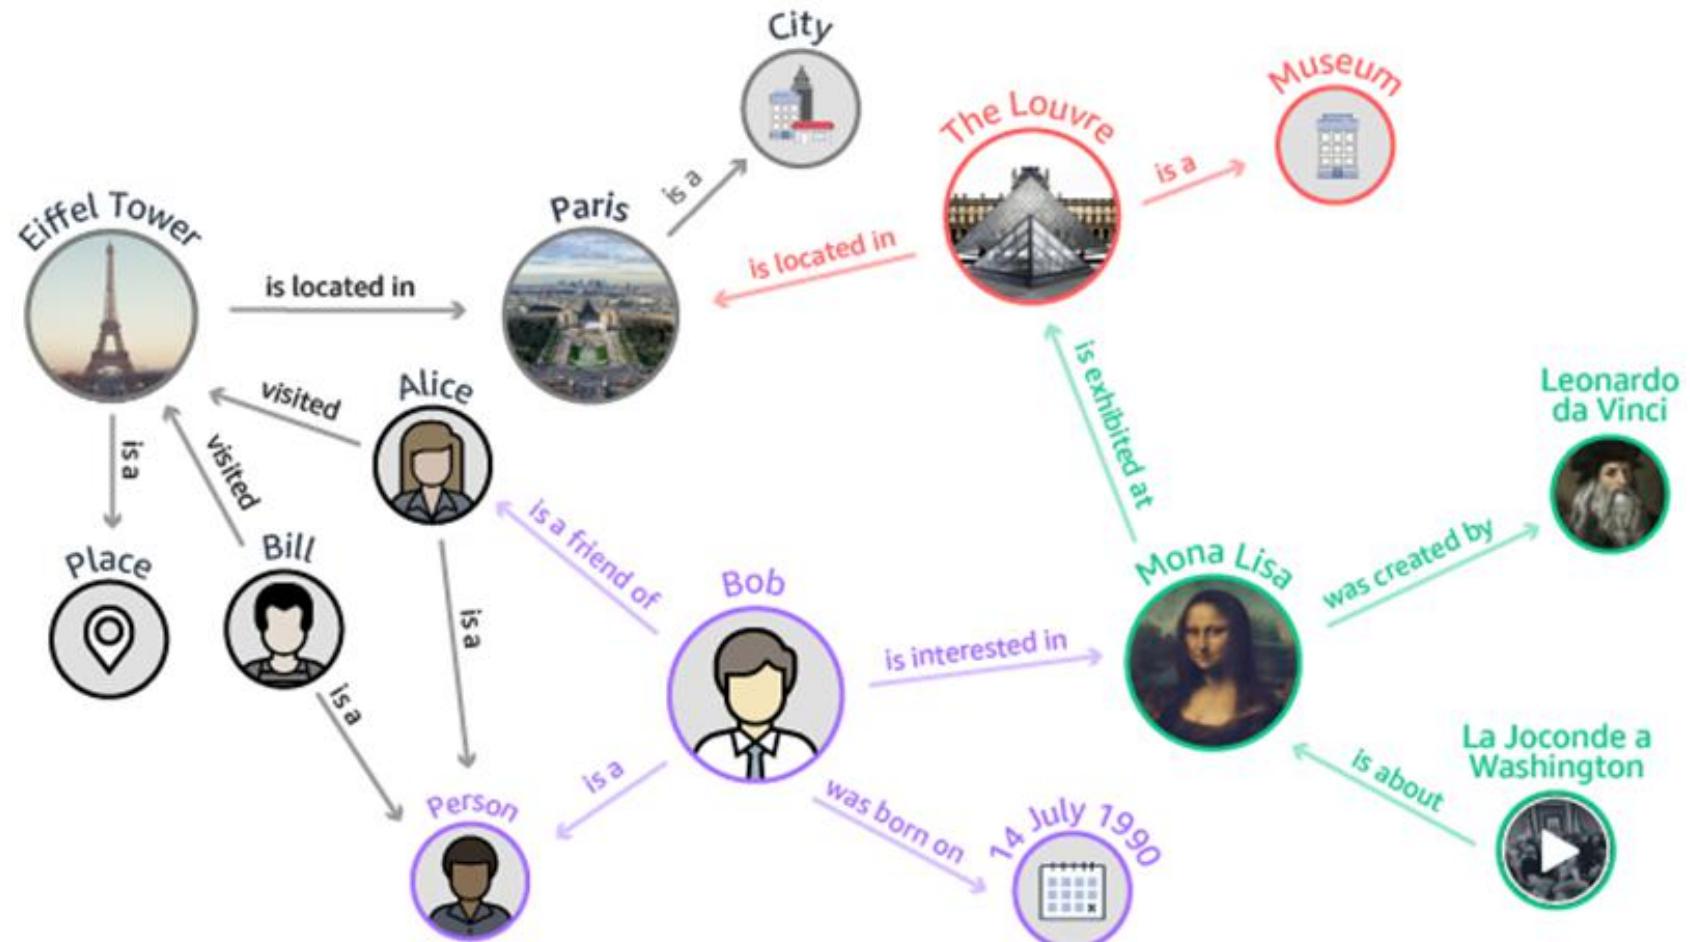
Artificial Intelligence

CS4365 --- Fall 2022

Knowledge Representation and Reasoning

Instructor: Yunhui Guo

Knowledge and Reasoning



Knowledge and Reasoning

- **Knowledge:**
 - the fact or condition of knowing something with familiarity gained through experience or association
- **Reasoning:**
 - the drawing of inferences or conclusions through the use of reason

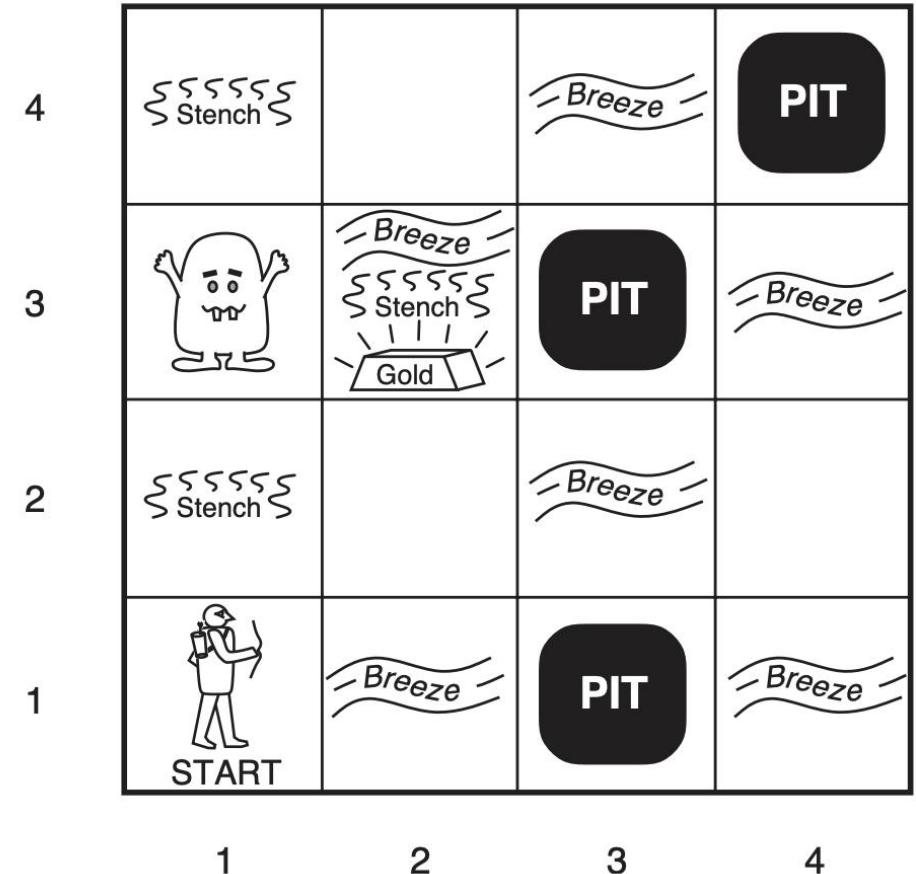
Knowledge + Reasoning → New Knowledge

The Wumpus World

- A decision-maker needs to represent **knowledge** of the world and **reason** with it in order to safely explore this world.

E.g. In the squares directly adjacent to a pit, the agent will perceive a Breeze

→ There a pit in [2, 2] or [3, 1] or both



Knowledge Representation

- Human intelligence relies on a lot of **background knowledge** (the more you know, the easier many tasks become / “knowledge is power”)
 - E.g. SEND + MORE = MONEY puzzle.
- Natural language understanding
 - Time flies like an arrow.
 - Fruit flies like bananas.
 - The spirit is willing but the flesh is weak. (English)
 - The vodka is good but the meat is rotten. (Russian)
- Or: Plan a trip to L.A.

Knowledge Representation

Q. How did we encode (domain) knowledge so far?
For search problems?

Fine for limited amounts of knowledge / well-defined domains.

Otherwise: **knowledge-based systems approach**

Knowledge-Based Systems / Agents

Key components

- knowledge base: a set of **sentences** expressed in some knowledge representation language
- Inference / reasoning mechanisms to query what is known and to derive new information or make decisions

Knowledge-Based Systems / Agents

- Natural candidate: **logical language** (propositional / first-order) combined with a logical inference mechanism
- How close to human thought?
- In any case, appears reasonable strategy for machines

Logic

- Logic:
 - defines a **formal language** for logical reasoning
- It gives us a tool that helps us to understand how to construct a valid argument
- Logic defines:
 - the **meaning** of statements
 - the rules of **logical inference**

Logic as a Knowledge Representation

Three components:

- syntax: specifies which sentences can be constructed in a given formal logic
 - E.g. $x + y = 4$
- semantics: specifies what a sentence means
 - $x + y = 4$ is True if $x = 2$ and $y = 2$
- proof theory: a set of **general purpose rules** that allow efficient derivation of new information from the sentences in the **knowledge base**

Logic as a Knowledge Representation

Model: a **truth assignment** to every propositional symbol

Logic entailment:

A sentence follows logically from another sentence:

$$\alpha \models \beta$$

In **every** model in which α is true, β is also true.

Logic as a Knowledge Representation

- Logic inference:
 - Given a knowledge base KB and a sentence α
 - Does a KB semantically entail α ? $KB \models \alpha$

One possible approach:

Model Checking: enumerate all the possible models to check if α is true in all models in which KB is true

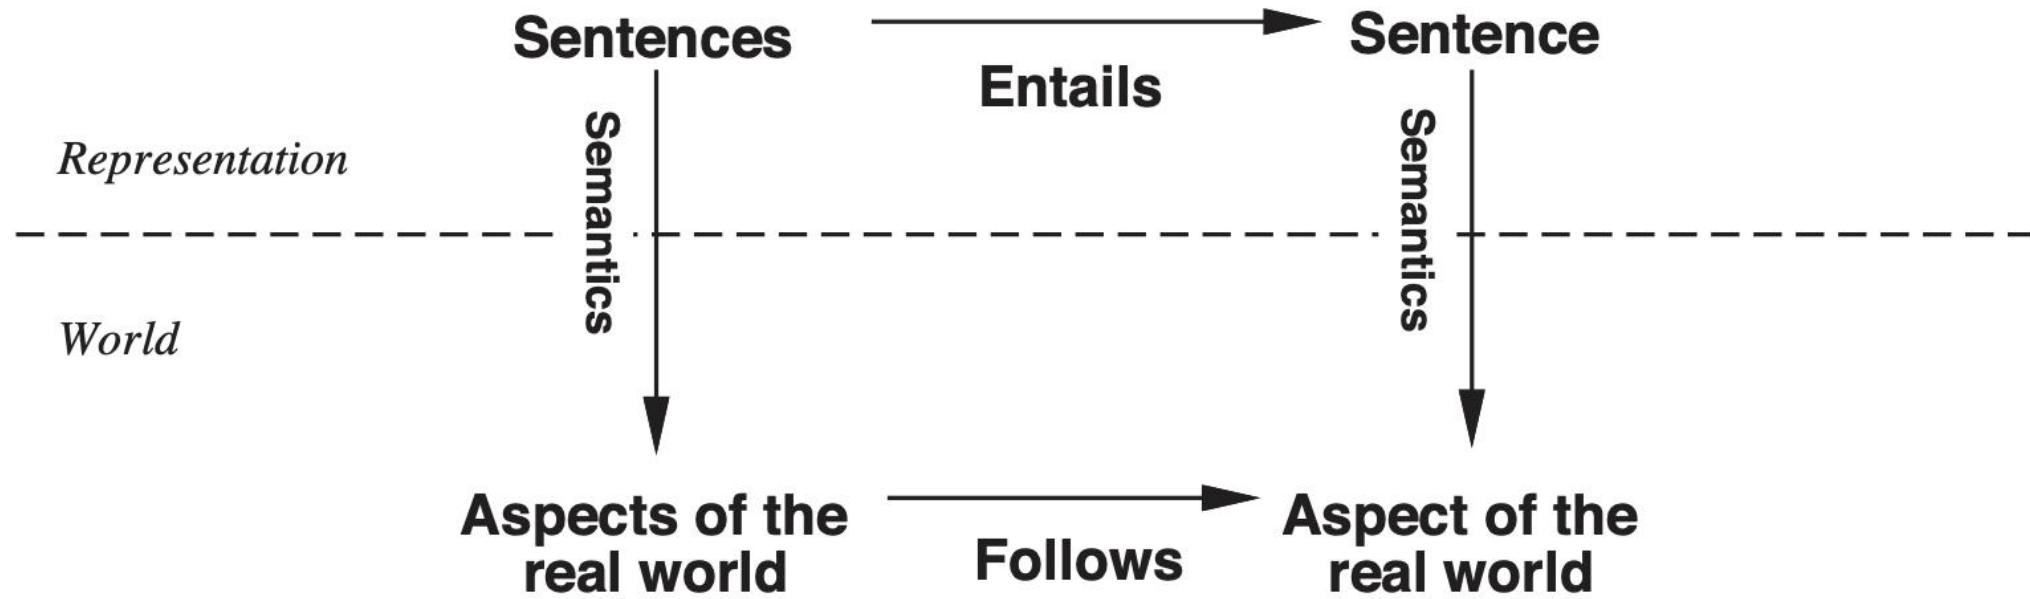
Logic as a Knowledge Representation

Proof theory:

Sound: An inference algorithm that derives only **entailed sentences**

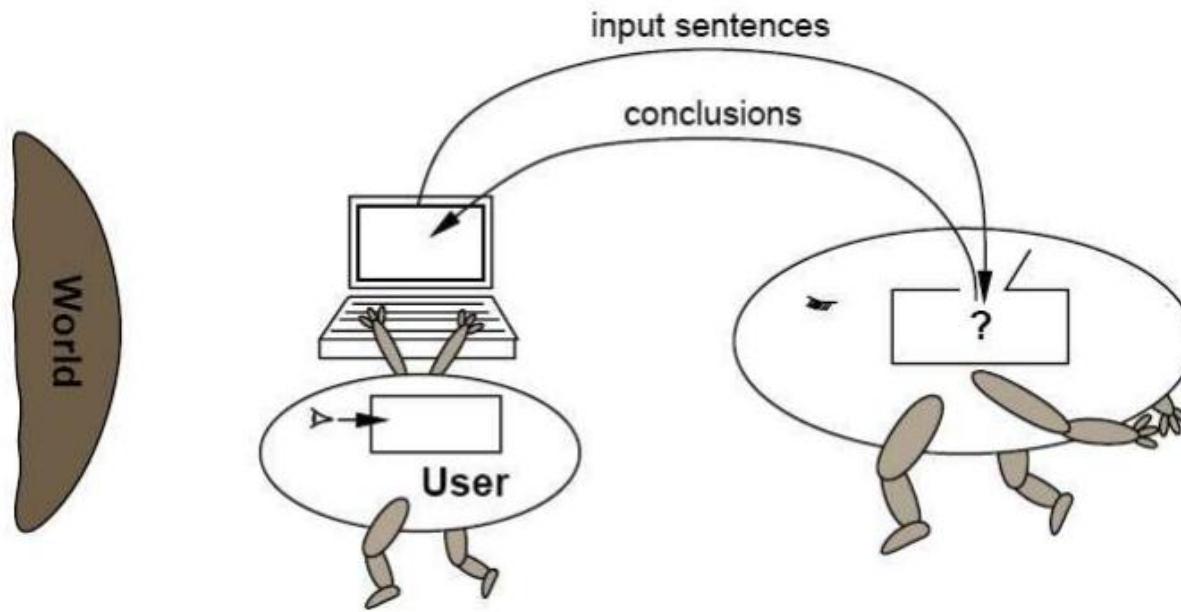
Complete: an inference algorithm is complete if it can derive any sentence that is entailed

Connecting Sentences to the Real World



- Logical reasoning should ensure that the new configurations represent aspects of the world that actually follow from the aspects that the old configurations represent.

Tenuous Link to Real World



- All computers have **sentences** (hopefully about the world).

KR Language: Propositional Logic

- Go back to 3rd century B.C. studied by Stoic school of philosophy
- Real development began in the mid-19th century and was initiated by the English mathematician G. Boole
- The classical propositional calculus was first formulated as a formal axiomatic system by the eminent German logician G. Frege in 1879.

KR Language: Propositional Logic

- The simplest logic
- Definition
 - A **proposition** is a statement that is either **true** or **false**.
- Example:
 - $5 + 2 = 8$ (F)
 - It is raining today
 - (either T or F)

KR Language: Propositional Logic

- Literal: an atomic formula or its negation
 - Positive literal: P, Q
 - Negative literal: $\neg P$, $\neg Q$
- Syntax: build sentences from atomic propositions, using connectives:
 - \wedge : and
 - \vee : or
 - \neg : not
 - \Rightarrow : implies
 - \Leftrightarrow : equivalence (biconditional)

KR Language: Propositional Logic

Syntax: build **sentences** from atomic propositions, using connectives \wedge , \vee , \neg , \Rightarrow , \Leftrightarrow
(and / or / not / implies / equivalence (biconditional))

E.g.: $\neg P$
 $Q \wedge R$

$(\neg P \vee (Q \wedge R)) \Rightarrow S$

KR Language: Propositional Logic

- Clause: a disjunction of literals
E.g.: $Q \vee R$
- Conjunctive normal form (CNF): a conjunction of clauses
E.g.: $(Q \vee R) \wedge (P \vee R)$
- Every formula can be equivalently written as a formula in conjunctive normal form
$$(Q \wedge R) \vee P \rightarrow (Q \vee P) \wedge (R \vee P)$$

Semantics

Semantics specifies what something means.

In propositional logic, the semantics (i.e., meaning) of a sentence is the set of interpretations (i.e., **truth assignments**) in which the sentence evaluates to True.

Example:

The semantics of the sentence $P \vee Q \Rightarrow R$ is

- P is True, Q is True, R is True
- P is True , Q is False, R is True
- P is False , Q is True , R is True
- P is False , Q is False , R is True
- P is False , Q is False , R is False

Interpretations: The Key to Semantics

An interpretation is a logician's word for “truth assignment”

- Given 3 propositional symbols P, Q, R, there are 8 interpretations.
- Given n propositional symbols P_1, P_2, \dots, P_n , there are 2^n interpretations

In propositional logic:

- an interpretation is a mapping from **propositional symbols** to **truth values**.
- the **meaning** of a sentence is the set of interpretations in which the sentence evaluates to True

How to evaluate a sentence under a given interpretation?

Evaluating a sentence under interpretation I

We can evaluate a sentence using a **truth table**

P	Q
<i>false</i>	<i>false</i>
<i>false</i>	<i>true</i>
<i>true</i>	<i>false</i>
<i>true</i>	<i>true</i>

Evaluating a sentence under interpretation I

We can evaluate a sentence using a **truth table**

P	Q	$\neg P$
<i>false</i>	<i>false</i>	<i>true</i>
<i>false</i>	<i>true</i>	<i>true</i>
<i>true</i>	<i>false</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>

Evaluating a sentence under interpretation I

We can evaluate a sentence using a **truth table**

P	Q	$\neg P$	$P \wedge Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>

Evaluating a sentence under interpretation I

We can evaluate a sentence using a **truth table**

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>

Evaluating a sentence under interpretation I

We can evaluate a sentence using a **truth table**

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>

Evaluating a sentence under interpretation I

We can evaluate a sentence using a **truth table**

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>

Evaluating a sentence under interpretation I

We can evaluate a sentence using a **truth table**

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>

Note: \Rightarrow is somewhat counterintuitive

What's the true value of “5 is even implies Sam is smart”

If P is True, then I claim Q is True

Three Important Concepts

- Logic Equivalence
- Validity
- Satisfiability

Logic Equivalence

- Two sentences are **equivalent** if they are true in the same set of models.
- We write this as $\alpha \equiv \beta$. $\alpha \equiv \beta$ if and only if $\alpha \vDash \beta$ and $\beta \vDash \alpha$

For example:

- I. If Lisa is in Denmark, then she is in Europe
- II. If Lisa is not in Europe, then she is not in Denmark

Logic Equivalence

- $(\alpha \wedge \beta) \equiv (\beta \wedge \alpha)$ commutativity of \wedge
- $(\alpha \vee \beta) \equiv (\beta \vee \alpha)$ commutativity of \vee
- $((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma))$ associativity of \wedge
- $((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma))$ associativity of \vee
- $\neg(\neg \alpha) = \alpha$ double-negation
- $(\alpha \Rightarrow \beta) \equiv (\neg \beta \Rightarrow \neg \alpha)$ contraposition
- $(\alpha \Rightarrow \beta) \equiv (\neg \alpha \vee \beta)$ implication elimination
- $(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha))$ biconditional elimination

Logic Equivalence

- $\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta)$ De Morgan
- $\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta)$ De Morgan
- $(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma))$ distributivity of \wedge over \vee
- $(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma))$ distributivity of \vee over \wedge

These equivalences play much the same role in logic as arithmetic identities do in ordinary mathematics.

Validity

- Some sentences are very true! For example

1) True

2) $P \Rightarrow P$

3) $(P \wedge Q) \Rightarrow Q$

A valid sentence is one whose meaning includes **every** possible interpretation.

$$((P \vee H) \wedge (\neg H)) \Rightarrow P$$

P	H	$P \vee H$	$(P \vee H) \wedge \neg H$	$((P \vee H) \wedge \neg H) \Rightarrow P$
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>

The truth table shows that $((P \vee H) \wedge (\neg H)) \Rightarrow P$ is valid

We write $\models ((P \vee H) \wedge (\neg H)) \Rightarrow P$

Satisfiability

- An unsatisfiable sentence is one whose meaning has **no interpretation** (e.g., $P \wedge \neg P$)
- A satisfiable sentence is one whose meaning has **at least** one interpretation.
- A sentence must be either **satisfiable** or **unsatisfiable** but it can't be both.
- If a sentence is valid then it's satisfiable.
- If a sentence is satisfiable then it may or may not be valid.

Satisfiability

- The SAT problem is to determine the **satisfiability** of sentences
- Connection to **validity**:
 - α is valid iff $\neg\alpha$ is unsatisfiable
 - α is satisfiable iff $\neg\alpha$ is not valid
- Proving by checking the unsatisfiability:
 $\alpha \vDash \beta$ if and only if the sentence $(\alpha \wedge \neg\beta)$ is unsatisfiable

Knowledge Base and Models

- Knowledge base: **a set of sentences**. Each sentence represents some assertion about the world.
- A model of a set of sentences (KB) is a truth assignment in which each of the KB sentences evaluates to True.
- With more and more sentences, the models of KB start looking more and more like the “real-world”.

Models

If a sentence α holds (is True) in **all models** of a KB, we say that α is entailed by the KB.

α is of interest, because whenever KB is true in a world α will also be True.

We write $KB \models \alpha$

Entailment Examples

KB

R1: CS4365Lectures \Rightarrow (TodayIsTuesday \vee TodayIsThursday)

R2: \neg TodayIsThursday

R3: TodayIsSaturday \Rightarrow SleepLate

R4: Rainy \Rightarrow GrassIsWet

R5: CS4365Lectures \vee TodayIsSaturday

R6: \neg SleepLate

Entailment Examples

KB

R1: CS4365Lectures \Rightarrow (TodayIsTuesday \vee TodayIsThursday)

R2: \neg TodayIsThursday

R3: TodayIsSaturday \Rightarrow SleepLate

R4: Rainy \Rightarrow GrassIsWet

R5: CS4365Lectures \vee TodayIsSaturday

R6: \neg SleepLate

Then which of these are correct entailments?

$KB \models \neg$ SleepLate

$KB \models \neg$ SleepLate \vee GrassIsWet

$KB \models$ GrassIsWet

$KB \models$ TodayIsTuesday

Entailment Examples

- KB
 - Propositional symbols:
 - CS4365Lectures, TodayIsTuesday, TodayIsThursday, TodayIsSaturday, SleepLate, Rainy, GrassIsWet
 - Model checking:
 - **Enumerate all the possible models** to check if α is true in all models in which KB is true

Entailment Examples

KB

R1: CS4365Lectures \Rightarrow (TodayIsThursday V TodayIsThusday)

R2: \neg TodayIsTuesday

R3: TodayIsSaturday \Rightarrow SleepLate

R4: Rainy \Rightarrow GrassIsWet

R5: CS4365Lectures V TodayIsSaturday

R6: \neg SleepLate

CS4365Lectures: T TodayIsThursday: T TodayIsTuesday: F

TodayIsSaturday: F SleepLate: F Rainy: F/T GrassIsWet: T/F

Entailment Examples

- KB is **True** when

- CS4365Lectures: T
- TodayIsThursday: T
- TodayIsTuesday: F
- TodayIsSaturday: F
- SleepLate: F
- Rainy: F/T
- GrassIsWet: T/F

$KB \models \neg \text{SleepLate}$ T

$KB \models \neg \text{SleepLate} \vee \text{GrassIsWet}$ T

$KB \models \text{GrassIsWet}$ F

$KB \models \text{TodayIsTuesday}$ F

- Complexity: $O(2^N)$

Logical Inference

- Problem definition:
 - The computer has a **knowledge base KB**.
 - The user inputs a **sentence**.
 - The computer tells the user whether the sentence is entailed by the knowledge base.

Humans who are doing proofs almost never use this brute-force approach. Then how to do logical inference **efficiently**?

Proof Theory

- A set of purely **syntactic** rules for efficiently determining entailment
- We write: $\text{KB} \vdash \alpha$, i.e., α can be **deduced** from KB or α is **provable** from KB .

Key property:

Both in **propositional** and in **first-order logic** we have a proof theory (“calculus”) such that:
 \models and \vdash are equivalent

Proof Theory (cont.)

If $\text{KB} \vdash \alpha$ implies $\text{KB} \vDash \alpha$, we say the proof theory is **sound**

If $\text{KB} \vDash \alpha$ implies $\text{KB} \vdash \alpha$, we say the proof theory is **complete**.

Why so important?

Allow computer to ignore semantics and “just push symbols”!

Example Proof Theory

One rule of inference: **Modus Ponens**

From α and $\alpha \Rightarrow \beta$ it follows that β .

Semantic soundness can easily be verified (using truth table).

Another rule of inference: **And-Elimination**

From $\alpha \wedge \beta$, it follows that α and β .

Example Proof Theory

Axiom schemas:

$$(\text{Ax. I}) \alpha \Rightarrow (\beta \Rightarrow \alpha)$$

$$(\text{Ax. II}) ((\alpha \Rightarrow (\beta \Rightarrow \gamma)) \Rightarrow ((\alpha \Rightarrow \beta) \Rightarrow ((\alpha \Rightarrow \gamma)))$$

$$(\text{Ax. III}) (\neg\alpha \Rightarrow \beta) \Rightarrow ((\neg\alpha \Rightarrow \neg\beta) \Rightarrow \alpha)$$

Note: α , β , γ stand for arbitrary sentences. So, we have an infinite collection of axioms.

Example Proof

- Now, α can be **deduced** from a set of sentences φ iff there exists **a sequence of applications** of modus ponens that leads from φ to α (possibly using axioms).
- One can prove that:
 - Modus ponens with the above axioms will generate exactly all (and only those) statements logically entailed by φ .

So, we have a way of generating entailed statements in a purely syntactic manner!

(Sequence is called a proof. Finding it can be hard ...)

Example Proof

Lemma. 1) For any α , we have $\vdash (\alpha \Rightarrow \alpha)$.

Proof.

$(\alpha \Rightarrow ((\alpha \Rightarrow \alpha) \Rightarrow \alpha)) \Rightarrow ((\alpha \Rightarrow (\alpha \Rightarrow \alpha)) \Rightarrow (\alpha \Rightarrow \alpha))$, Ax. II

$\alpha \Rightarrow ((\alpha \Rightarrow \alpha) \Rightarrow \alpha)$, Ax. I

$(\alpha \Rightarrow (\alpha \Rightarrow \alpha)) \Rightarrow (\alpha \Rightarrow \alpha)$; Modus Ponens

$\alpha \Rightarrow (\alpha \Rightarrow \alpha)$, Ax. I

$\alpha \Rightarrow \alpha$, Modus Ponens

Another Example Proof

Lemma. 2) For any α and β , we have $\beta, \neg\beta \vdash \alpha$

Proof.

$(\neg\alpha \Rightarrow \beta) \Rightarrow ((\neg\alpha \Rightarrow \neg\beta) \Rightarrow \alpha)$, (Ax. III)

β , (hyp.)

$\beta \Rightarrow (\neg\alpha \Rightarrow \beta)$, (Ax. I)

$\neg\alpha \Rightarrow \beta$, (Modus Ponens)

$(\neg\alpha \Rightarrow \neg\beta) \Rightarrow \alpha$, (Modus Ponens)

$\neg\beta$, (hyp.)

$\neg\beta \Rightarrow (\neg\alpha \Rightarrow \neg\beta)$, (Ax. I)

$\neg\alpha \Rightarrow \neg\beta$, (Modus Ponens)

α , (Modus Ponens)

Another Example Proof

Why are lemma 1 and lemma 2 true **semantically**?

I.e., $\models \alpha \Rightarrow \alpha$ and $\beta, \neg\beta \models \alpha$

- Logic entailment: $\varphi \models \gamma$
 - In **every** model in which φ is **true**, γ is also **true**.
 - No model ω in which φ is **true** and γ is **false**.
- False entails anything (Principle of explosion)

Another Example Proof

Why are lemma 1 and lemma 2 true **semantically**?

I.e., $\models \alpha \Rightarrow \alpha$ and $\beta, \neg\beta \models \alpha$

Note: **proofs** are purely **syntactic** --- machines does not need to know anything about the meaning of the sentences!

Whatever is **syntactically** derived will be **semantically** true, and we can derive everything syntactically that is semantically true.

How hard is it to find proofs?

Sound rules of inference

Rule	Premise	Conclusion
Modus Ponens	$A, A \rightarrow B$	B
And Introduction	A, B	$A \wedge B$
And Elimination	$A \wedge B$	A
Double Negation	$\neg\neg A$	A

Monotonicity

- The set of **entailed sentences** can only **increase** as information is added to the knowledge base.
- For any sentence α and β
if $KB \models \alpha$ then $KB \wedge \beta \models \alpha$
- Propositional logic is monotonic

Key Properties

We have the following properties (also for first-order logic):

For a **sound** and **complete** proof theory, the following three conditions are equivalent:

- (I) $\varphi \vDash \alpha$
- (II) $\varphi \vdash \alpha$
- (III) $\varphi, \neg\alpha$ is inconsistent (i.e., can be refuted)

(I) is semantic; (II) syntactic; (III) at high-level semantic but we have a nice syntactic automatic procedure: **resolution**.

What common proof technique does III represent?

Resolution

First need **canonical form**: CNF

Conjunction of disjunctions / CNF (conjunctive normal form)

Example: $\neg(P \Rightarrow Q) \vee (R \Rightarrow P)$

$$\neg(\neg P \vee Q) \vee (\neg R \vee P)$$

$$(P \wedge \neg Q) \vee (\neg R \vee P) \text{ (de Morgan's law)}$$

$$(P \vee \neg R \vee P) \wedge (\neg Q \vee \neg R \vee P) \text{ (assoc. and distri. laws)}$$

$$(P \vee \neg R) \wedge (\neg Q \vee \neg R \vee P)$$

$$\{(P \vee \neg R), (\neg Q \vee \neg R \vee P)\}$$

$$\{(P, \neg R), (\neg Q, \neg R, P)\} \text{ (just notation)}$$

Resolution

Given a CNF, a single inference rule (and no axioms) will allow us to determine **inconsistency**.

So, using property III (above) and **resolution**, we have a **sound** and **complete** proof procedure for propositional logic (which can be extended to first-order logic).

The Resolution Rule (clausal form)

From $\alpha \vee p$ and $\neg p \vee \beta$, we can derive:

$\alpha \vee \beta$ (α and β are **disjunctions of literals**, where a literal is a propositional variable or its negation).

- Reason:
 - If p is true, β must be true
 - If p is false, α must be true
 - So β or α holds

The Resolution Rule (clausal form)

From $\alpha \vee p$ and $\neg p \vee \beta$, we can derive:

$\alpha \vee \beta$ (**resolvent**)

(α and β are **disjunctions of literals**, where a literal is a propositional variable or its negation).

Note: $\neg\alpha \Rightarrow p$ and $p \Rightarrow \beta$ gives $\neg\alpha \Rightarrow \beta$

It is a “chaining rule”.

General Resolution Rule

If $(L_1 \vee L_2 \vee \dots \vee L_k)$ is True,

and $(\neg L_k \vee L_{k+1} \vee \dots \vee L_m)$ True,

then we can conclude that

$(L_1 \vee L_2 \vee \dots \vee L_{k-1} \vee L_{k+1} \vee \dots \vee L_m)$ is True.

A **resolution-based theorem prover** can, for any sentences α and β in propositional logic, decide whether $\alpha \models \beta$

General Resolution Rule

- We can derive the **empty clause** via resolution iff the set of clauses is **inconsistent**. ($\neg p \wedge p$)
- Method relies on property III. It's **refutation complete**. Note that method does not generate theorems from scratch

E.g., we have $p \wedge r \vDash p \vee r$, but we can't get $p \vee r$ from $\{\{p\}, \{r\}\}$.
(Why not??)

But, given $\{\{p\}, \{r\}\}$ and the negation of $p \vee r$, we get the set $\{\{p\}, \{r\}, \{\neg p\}, \{\neg r\}\}$. Resolving on this set gives the **empty clause**. Thus contradiction. Thus proof.

Algorithm: Resolution Proof

- Negate the theorem to be proved, and add the result to the list of sentences in the KB. We show that $(KB \wedge \neg\alpha)$ is unsatisfiable
- Put the list of sentences into conjunctive normal form.
- Until there is no resolvable pair of clauses,
 - Find resolvable clauses and resolve them.
 - Add the results of resolution to the list of clauses.
 - If NIL (empty clause) is produced, stop and report that the (original) theorem is true. (empty clause represents contradiction)
- Report that the (original) theorem is false

Algorithm: Resolution Proof

- May seem cumbersome but note that can be easily **automated**. Just “smash” clauses till empty clause or no more new clauses
- Guaranteed sound and (refutation) complete.
- Q. Why is method with axioms more difficult to implement?

Example

- 1) ARM-OK
- 2) \neg MOVES
- 3) ARM-OK \wedge LIFTABLE \Rightarrow MOVES
- 4) \neg ARM-OK \vee \neg LIFTABLE \vee MOVES

Prove: \neg LIFTABLE

- 5) LIFTABLE (assert)
- 6) \neg ARM-OK \vee MOVES (resolving 5 and 4)
- 7) \neg ARM-OK (from 6 and 2)
- 8) Nil (empty clause / contradiction, from 7 and 1).

Another Example

$$KB = (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1}$$

$$\alpha: \neg P_{1,2}$$

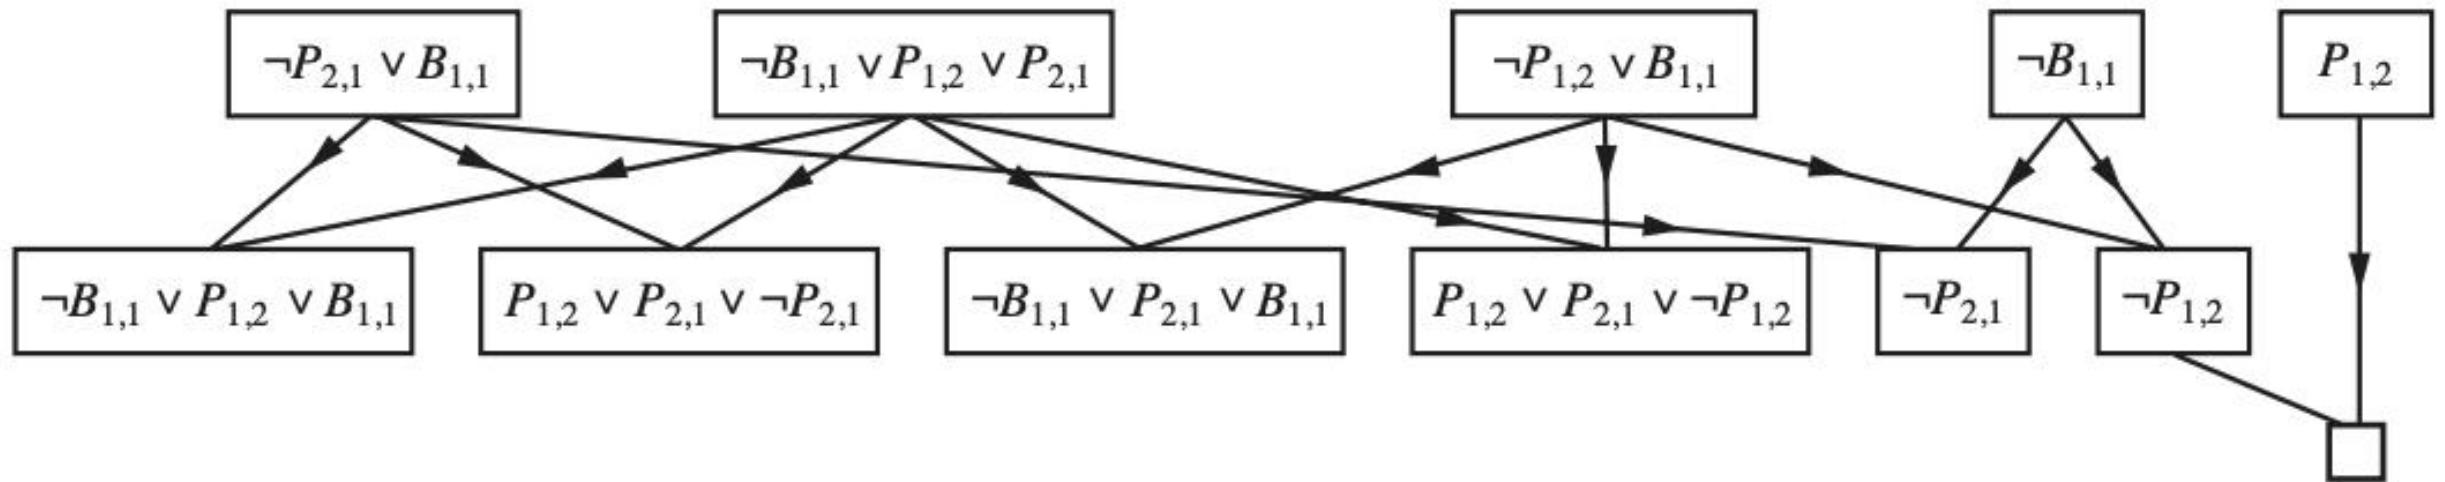
- First convert $(KB \wedge \neg\alpha)$ into CNF

Another Example

- First convert $(KB \wedge \neg\alpha)$ into CNF
- $(B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1} \wedge P_{1,2}$
- $((B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})) \wedge ((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1})) \wedge \neg B_{1,1} \wedge P_{1,2}$ (biconditional)
- $((\neg B_{1,1} \vee (P_{1,2} \vee P_{2,1})) \wedge (\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1})) \wedge \neg B_{1,1} \wedge P_{1,2}$ (Implication elimination)
- $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge ((\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1}) \wedge \neg B_{1,1} \wedge P_{1,2}$ (De Morgan)
- $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1}) \wedge \neg B_{1,1} \wedge P_{1,2}$ (Distri.)

Another Example

- Resolution proof



Completeness of Resolution

- Resolution closure (RC): the set of all clauses derivable by repeated application of the resolution rule to clauses in S or their derivatives
- Ground resolution theorem: If a set of clauses is unsatisfiable, then the resolution closure of those clauses contains the empty clause
- Proof by contrapositive:
 - if the closure $RC(S)$ does not contain the empty clause, then S is satisfiable

Completeness of Resolution

- Construct a **model** for S with **suitable truth values** for P_1, \dots, P_k

For i from 1 to k ,

- If a clause in $RC(S)$ contains the literal $\neg P_i$ and all its other literals are **false** under the assignment chosen for P_1, \dots, P_{i-1} , then assign **false** to P_i .
- Otherwise, assign **true** to P_i .

This assignment to P_1, \dots, P_k is a **model** of S .

Completeness of resolution

- Assume at some stage i in the construction, assigning P_i causes some C to be false.
- C must look like $(\text{false} \vee \text{false} \dots \vee P_i)$ or $(\text{false} \vee \text{false} \dots \vee \neg P_i)$
- C can only be falsified if both of these clauses are in $RC(S)$
- $RC(S)$ will contain resolvent of these two clauses (already falsified)
- Contradiction

Horn clauses and Definite Clauses

- **Definite clauses**: a disjunction of literals of which exactly one is positive
 - $\neg L_{1,1} \vee \neg P_{1,2} \vee B_{1,1}$,
- **Horn clauses**: a disjunction of literals of which at most one is positive
- Horn clauses are closed under resolution

Horn clauses and Definite Clauses

- Every **definite clause** can be written as an implication whose premise is a conjunction of positive literals and whose conclusion is a single positive literal.
 - $(L_{1,1} \wedge P_{1,2}) \Rightarrow B_{1,1}$,
- Inference with Horn clauses can be done through the **forward-chaining** and **backward-chaining** algorithms
- Deciding entailment with **Horn clauses** can be done in time that is **linear** in the size of the knowledge base

Forward Chaining

- From facts to conclusions
- It begins from known facts (positive literals) in the knowledge base
- If all the premises of an implication are known, then its conclusion is added to the set of known facts
- Given $(L_{1,1} \wedge P_{1,2}) \Rightarrow B_{1,1}$, If $L_{1,1}$ and $P_{1,2}$, then $B_{1,1}$ can be added to the knowledge base

Forward Chaining

- It begins from **known facts** (positive literals) in the knowledge base
- If all the **premises** of an **implication** are known, then its **conclusion** is added to the set of known facts

$$P \Rightarrow Q$$

$$L \wedge M \Rightarrow P$$

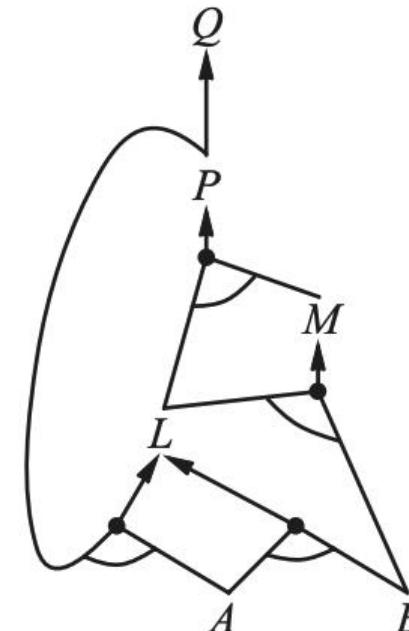
$$B \wedge L \Rightarrow M$$

$$A \wedge P \Rightarrow L$$

$$A \wedge B \Rightarrow L$$

$$A$$

$$B$$



Forward Chaining

- **Sound:** every inference is essentially an application of Modus Ponens
- **Complete:** every entailed atomic sentence will be derived.

Backward Chaining

- From **goal** to **data**
- Assume the **query q is true** and works **backward**
- If all the premises of one of those implications can be proved true (by backward chaining), then q is true

Propositional logic is a weak language

- Hard to identify “**individuals**” (e.g., Mary, 3)
- Can’t directly talk about **properties** of individuals or **relations** between individuals (e.g., “Bill is tall”)
- **Generalizations, patterns, regularities** can’t easily be represented (e.g., “all triangles have 3 sides”)

Length of Resolution Proof

Consider **Pigeon-Hole (PH) problem**: Formula encodes that you cannot place $n+1$ pigeons in n holes (one per hole).

PH takes **exponentially** many steps (no matter in what order)!

PH hidden in many practical problems. Makes theorem proving expensive.

Pigeon-Hole Principle

$P_{i,j}$ for Pigeon i in hole j.

$$P_{1,1} \vee P_{1,2} \vee P_{1,3} \dots P_{1,n}$$

$$P_{2,1} \vee P_{2,2} \vee P_{2,3} \dots P_{2,n}$$

....

$$P_{(n+1),1} \vee P_{(n+1),2} \vee P_{(n+1),3} \dots P_{(n+1),n}$$

and ??

Pigeon-Hole Principle

$(\neg P_{1,1} \vee \neg P_{1,2}), (\neg P_{1,1} \vee \neg P_{1,3}), (\neg P_{1,1} \vee \neg P_{1,4}), \dots$

$(\neg P_{1,(n-1)} \vee \neg P_{1,n}),$

$(\neg P_{2,1} \vee \neg P_{2,2}) \dots (\neg P_{2,(n-1)} \vee \neg P_{2,n})$

etc.

$(\neg P_{1,1} \vee \neg P_{2,1}), (\neg P_{1,1} \vee \neg P_{3,1}), \dots$

$(\neg P_{1,2} \vee \neg P_{2,2}), (\neg P_{1,2} \vee \neg P_{3,2}),$ **etc.**

Pigeon-Hole Principle

Resolution proof of **inconsistency** requires at least an **exponential** number of clauses, no matter in what order you resolve things!

A More Concise Formulation

$$\forall x \exists y (x \in \text{Pigeons}) (y \in \text{Holes}) \text{IN}(x, y)$$
$$\forall x \forall x' \forall y (\text{IN}(x, y) \wedge \text{IN}(x', y) \dots ??)$$
$$\forall x \forall y \forall y' (\text{IN}(x, y) \wedge \text{IN}(x, y') \dots ??)$$

Pigeons = $\{p_1, p_2, \dots, p_{n+1}\}$,

Holes = $\{h_1, h_2, \dots, h_n\}$

We have **first-order logic** with some set-theory notation.

Midterm review

- Uninformed search
 - Breadth first search
 - Depth first search
 - Uniform cost search
 - Iterative deepening search
 - bidirectional search

Midterm review

Generic Tree-Search Algorithm

Add initial state to the **frontier**

Loop

node = **remove-frontier()** -- and save in order to return as part of path to **goal**

 if **goal-test(node)** = **true** **return** path to **goal**

S = **successors(node)**

 Add **S** to **frontier**

until **frontier** is empty

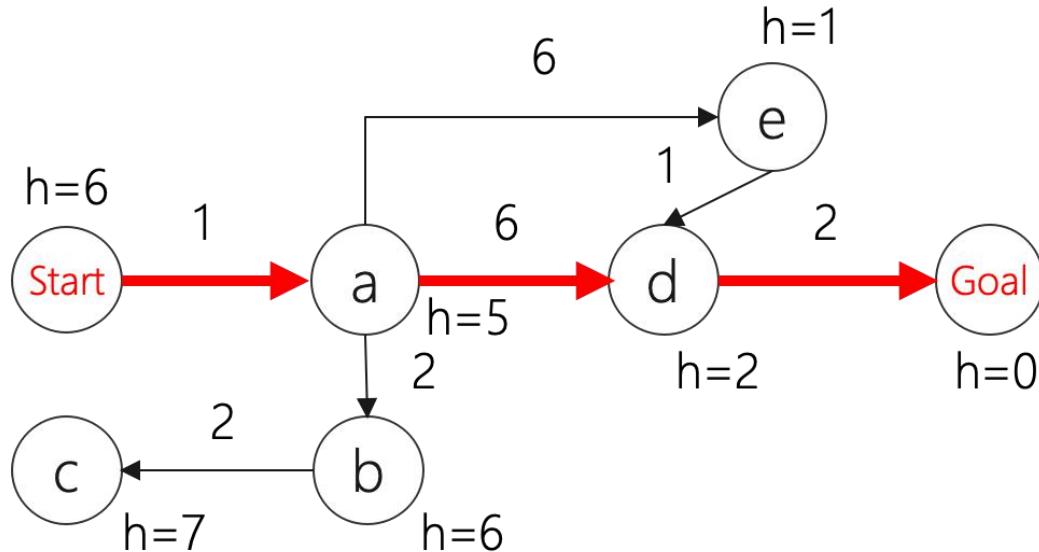
return **failure**

Midterm review

- Informed search
 - Greedy Best-First Search
 - Expands the node that is “closest” to the goal as measured by $h(n)$
 - A* search
 - Combining Uniform-Cost Search and Greedy Best-First Search
 - $f(n) = g(n) + h(n)$
 - $g(n)$: the path cost from the start node to node n
 - $h(n)$: the estimated cost of the cheapest path from node n to the goal node

Midterm review

A* Search



- Uniform-cost orders by path cost $g(n)$
- Greedy best first search orders by estimated goal proximity $h(n)$
- A* search combines $g(n)$ and $h(n)$

Midterm review

- If h is **admissible**, then the A^* tree search is optimal
- If h is **consistent**, then the A^* graph search is optimal

Midterm review

Hill Climbing Search

- Move in the direction of increasing value
- Terminate when it reaches a “peak” where no neighbor has a higher value.

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

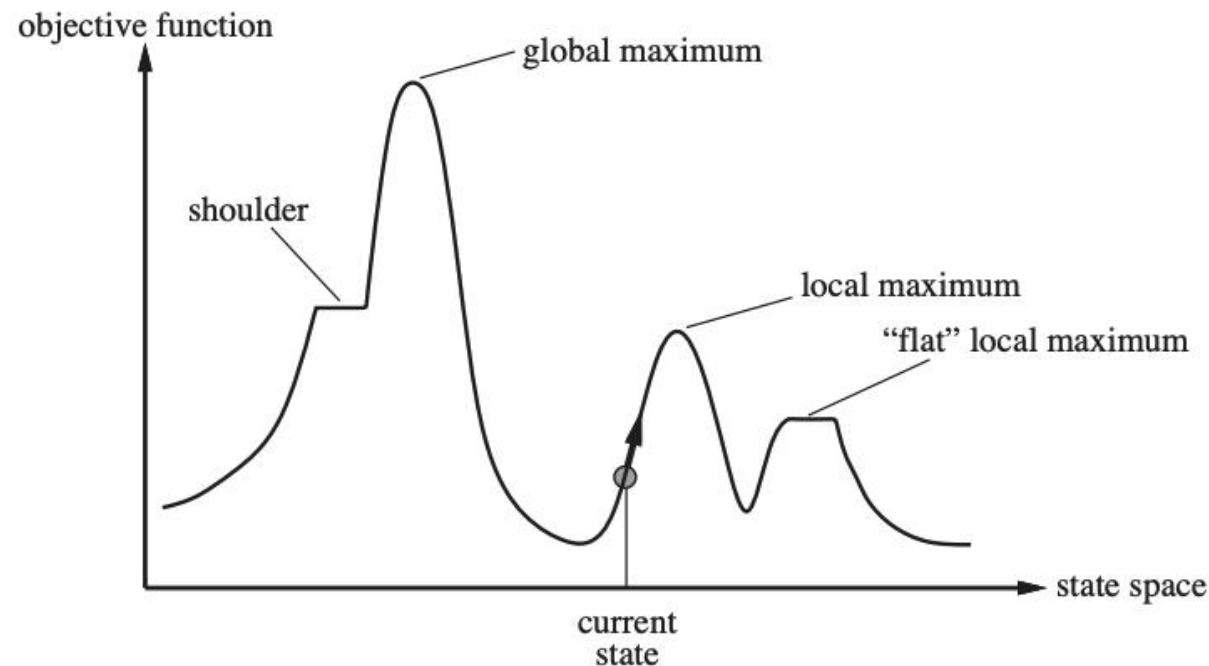
if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

Midterm review

State-Space Landscape

- Local Search Methods rely on an objective function or a cost function to compute the state-space landscape



- Local search algorithms explore this landscape

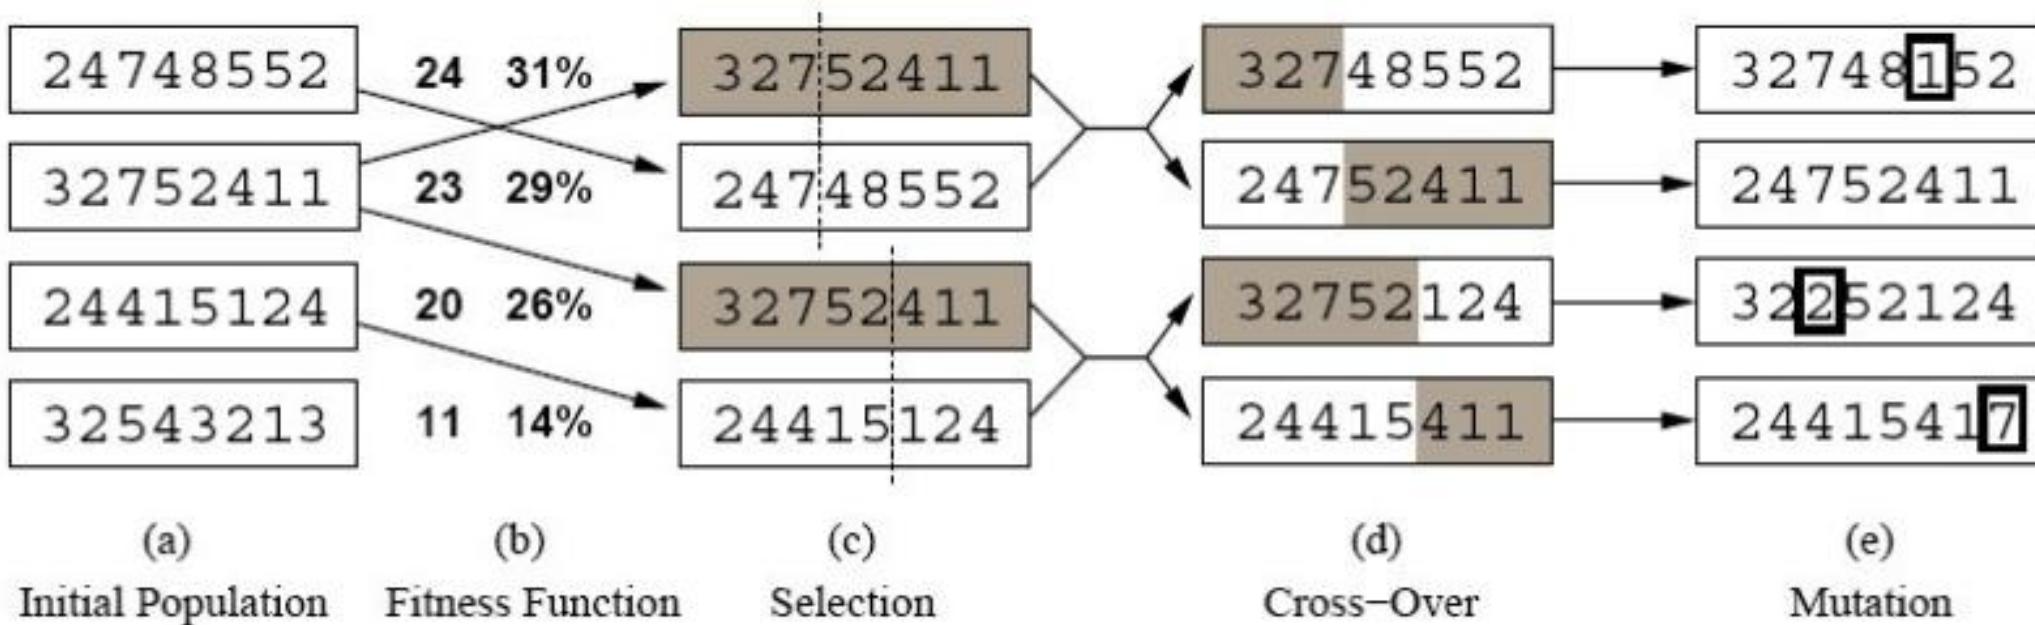
Midterm review

Improvements to Basic Local Search

- Issue: How to move more quickly to successively higher plateaus and avoid getting “stuck” / local minima
- Idea: Introduce uphill moves (“noises”) to escape from long plateaus (or true local minima).
- Strategies:
 - Simulated Annealing
 - Random-restart hill-climbing
 - Tabu search
 - Local beam search
 - Genetic Algorithms

Midterm review

Genetic Algorithm



Midterm review

Crossover Operators

- Single-point crossover:

Parent A: 1 0 0 1 0 1 1 1 0 1

Parent B: 0 1 0 1 1 1 0 1 1 0

Child AB: 1 0 0 1 0 1 0 1 1 0

Child BA: 0 1 0 1 1 1 1 1 0 1

Midterm review

Mutation

- Mutation: randomly flip one bit

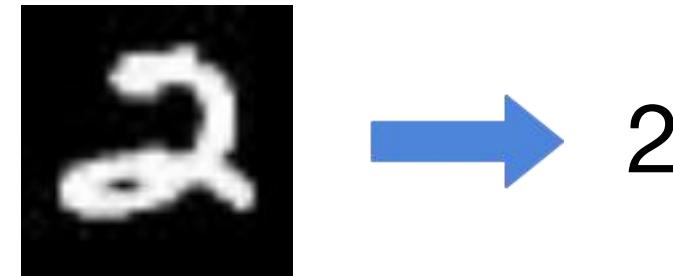
Individual A: 1 0 0 1 0 1 1 1 0 1



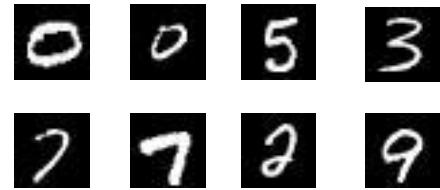
Individual A': 1 1 0 0 0 1 1 1 0 1

Machine Learning

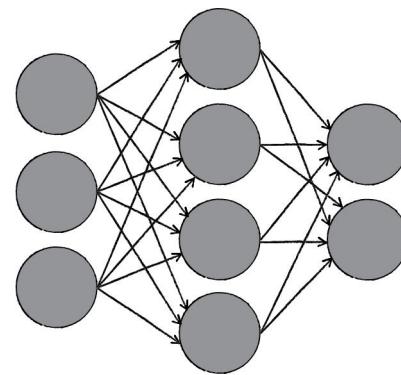
Task: classify a given image



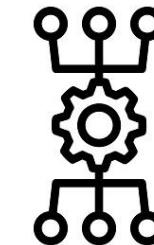
Data



Model

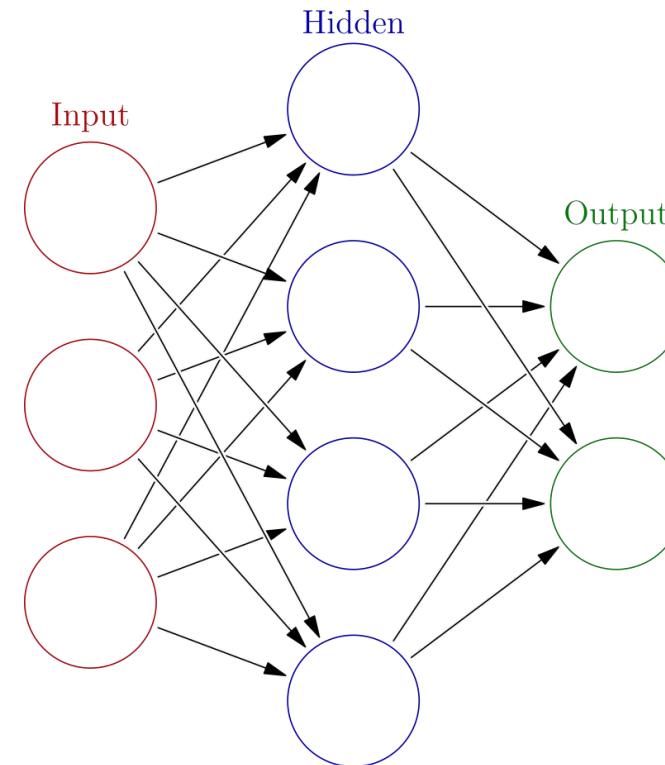
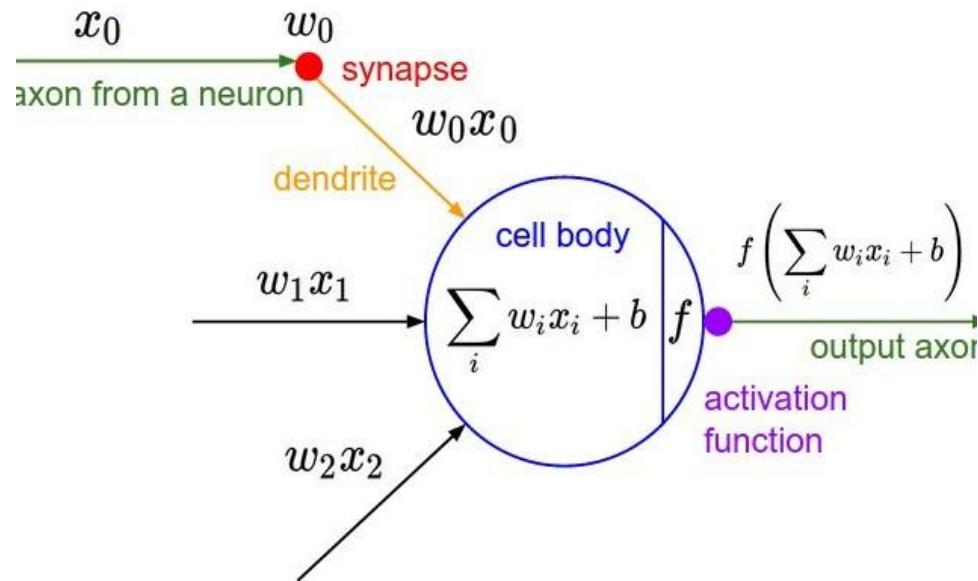


Algorithm

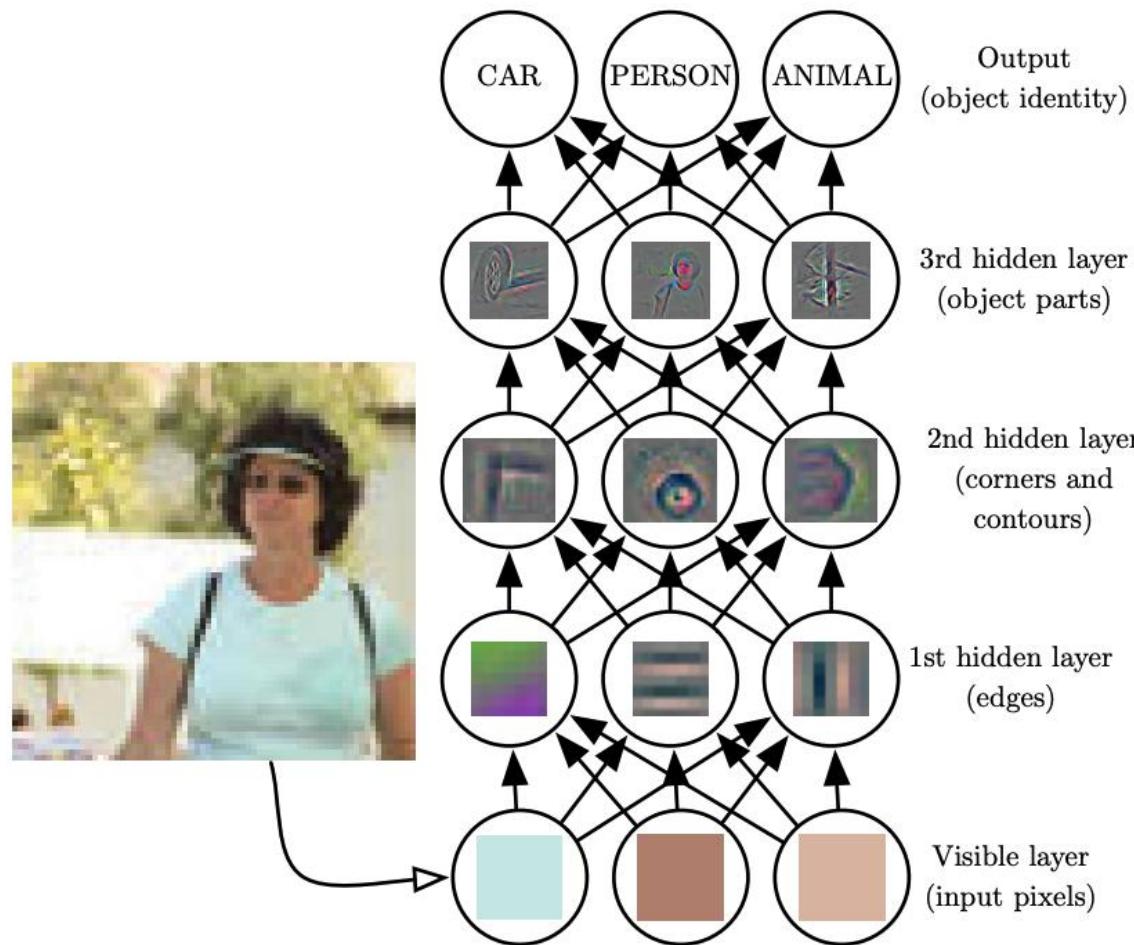


Neural Networks

- Neural networks consist of a collection of nodes which model biological neurons in human brain.
- Neural networks have one input layer, one or multiple hidden layers and one output layer



What are learnt by neural networks

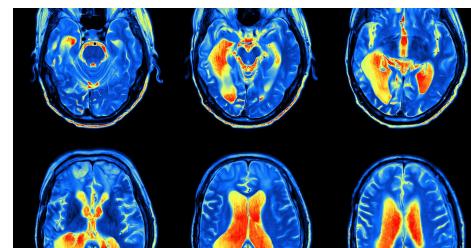
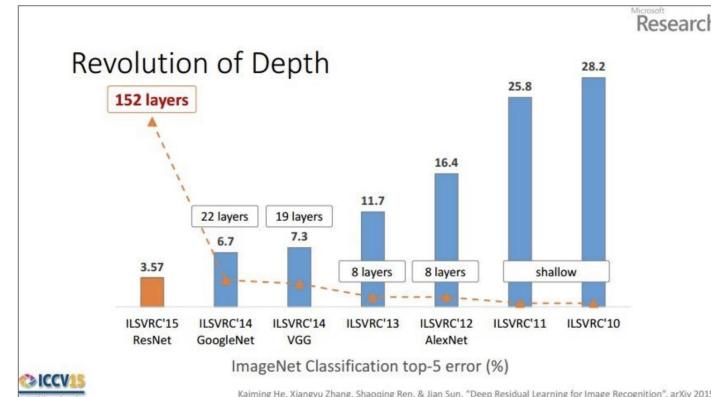


Deep Neural Networks: The Success

Large-Scale Dataset



Top-5 Error on ImageNet

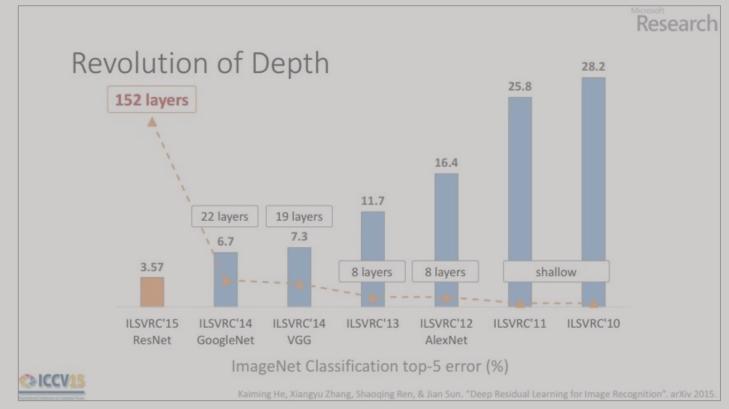


Practical Deployment

Large-Scale Dataset



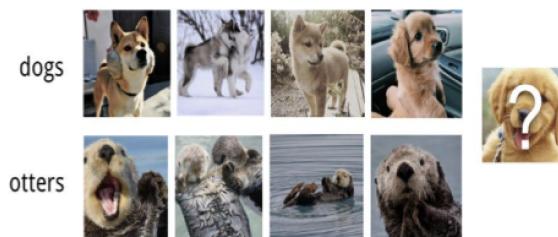
Top-5 Error on ImageNet



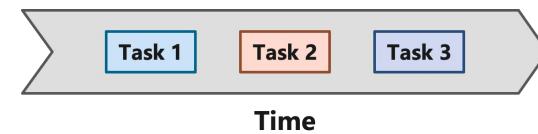
Limited Resources



Limited Examples

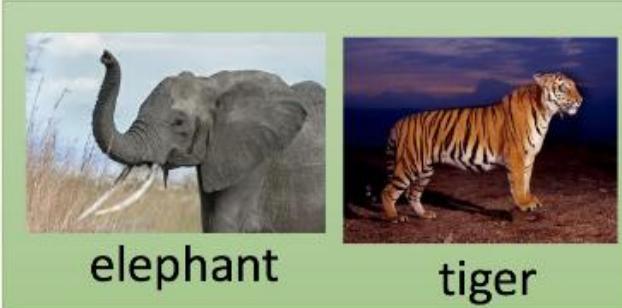


A Sequence of Tasks



Transfer Learning

Dataset 1



elephant

tiger

Knowledge extracted
from dataset 1 can be
reused for learning on
dataset 2



Dataset 2



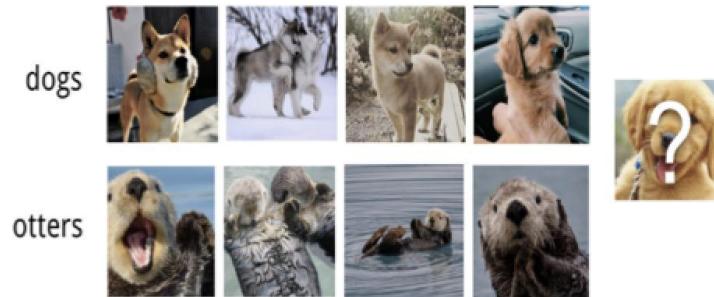
dog

cat

Few-shot Learning

N-way K-shot few-shot learning task:

- Each task has a total of N classes.
- Each class has K examples.



2-way 4-shot

Sequential Tasks

Rainy



Cloudy



Sunny



Dog or Cat



Tiger or Lion



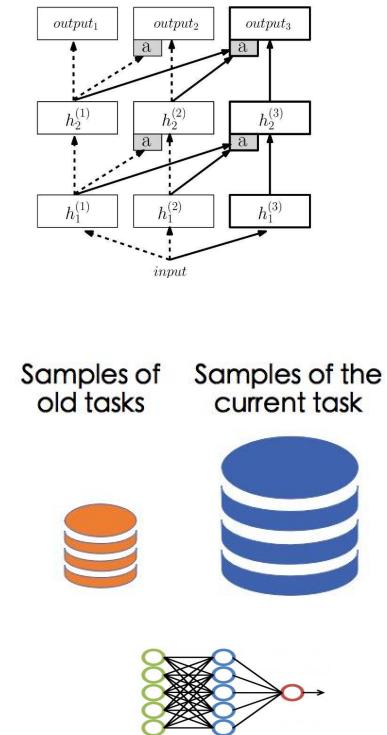
Dolphin or Shark



Lifelong Learning

- Regularization based approach
 - Prevent **important weights** for old tasks from changing drastically.
- Knowledge transfer based method
 - Build **connections** between the models for the new task and old tasks.
- Episodic memory based approach
 - Store **a small subset of examples** from old tasks

$$\mathcal{L}(\theta) = \mathcal{L}_B(\theta) + \sum_i \frac{\lambda}{2} F_i (\theta_i - \theta_{A,i}^*)^2$$



Semantic Segmentation

