

Artificial Intelligence

CS4365 --- Fall 2022

Uninformed Search

Instructor: Yunhui Guo

Problem Solving as Search

- **Search** is needed to solve many real-world problems

8-puzzle problem

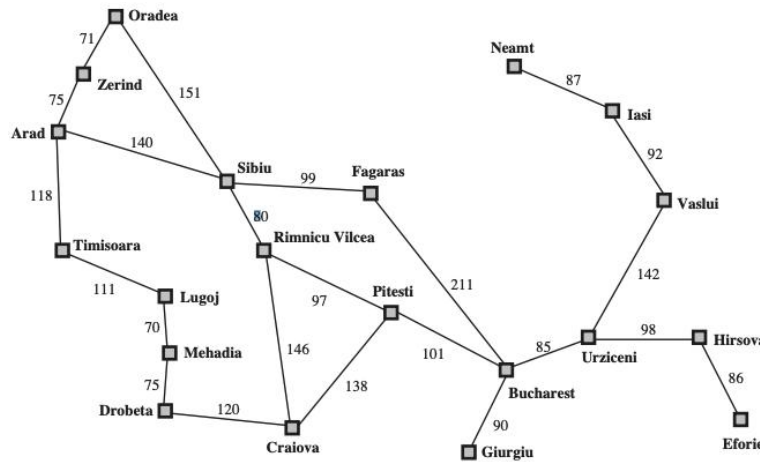
5	4	
6	1	8
7	3	2

Start State


1	2	3
8		4
7	6	5

Goal State

Planning




Go




Problem Solving as Search

- Search is a central topic in AI
 - Originated with Newell and Simon's work on problem solving.
Famous book:
“Human Problem Solving” (1972)
 - Automated reasoning is a natural search task
 - More recently: Given that almost all AI formalisms (planning, learning, etc.) are NP-complete or worse, some form of search is generally unavoidable (no “smarter” algorithm available).

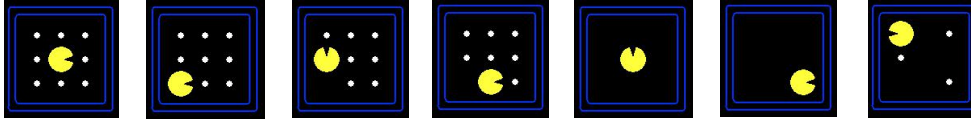


Pac-Man



Define a Search Problem

- A **search problem** consists of:

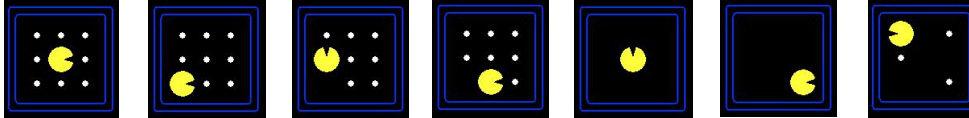
- **State space:** 

- A **successor function:**
(action + cost)


- A **start state** and a **goal test**

Define a Search Problem

- A **search problem** consists of:

- **State space:** 

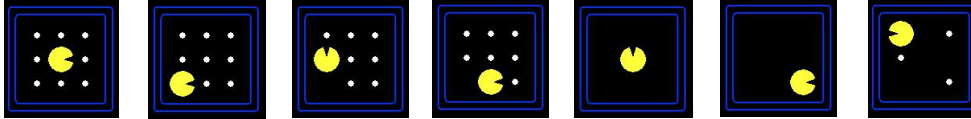
- A **successor function:**
(action + cost)




- A **start state** and a **goal test**

Define a Search Problem


- A **search problem** consists of:

- **State space:** 

- A **successor function:**
(action + cost)




- A **start state** and a **goal test**



Define a Search Problem

- A **path** is any sequence of states connected by a sequence of actions.
- **Path cost** – function that assigns a cost to a path; relevant if more than one path leads to the goal, and we want the shortest path.
- A **solution** is a sequence of actions (a plan) which transforms the **start state** to a **goal state**



Example: The 8-Puzzle

- State space: ?
- Initial state: ?
- Goal test: ?
- Successor function: ?
- Path cost: ?

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Example: The 8-Puzzle

- **State space:**
 - The location of each tile
- **Initial state:**
 - Any state can be the initial state
- **Goal test:**
 - Whether the state matches the goal state
- **Successor function:**
 - The movement of the blank space (Left, Right, Up or Down)
- **Path cost:**
 - Each step costs 1

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Example: Cryptarithmic

$$\begin{array}{r} \text{SEND} \\ + \text{MORE} \\ \hline \text{MONEY} \end{array}$$

- Find **substitution of digits for letters** such that the resulting sum is arithmetically correct.
- Each letter must stand for a different digit

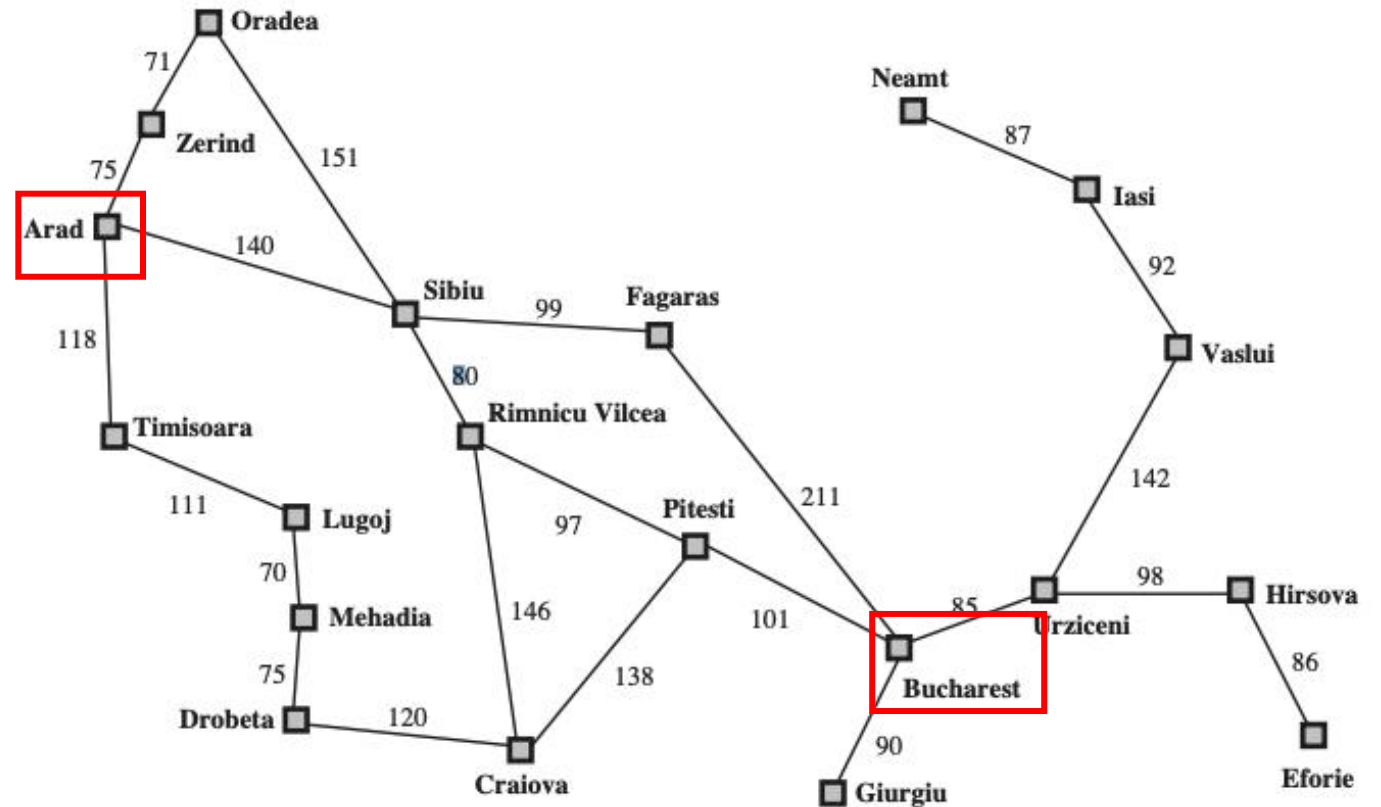
Cryptarithmic, cont.

- **State space**: an 8-tuple indicating a (partial) assignment of digits to letters.
- **Goal test**: all letters have been assigned digits and sum is correct
- **Successor function**: represents the act of assigning digits to letters
- **Path cost**: all solutions are equally valid; step cost = 0

```
      SEND
    + MORE
    -----
      MONEY
```

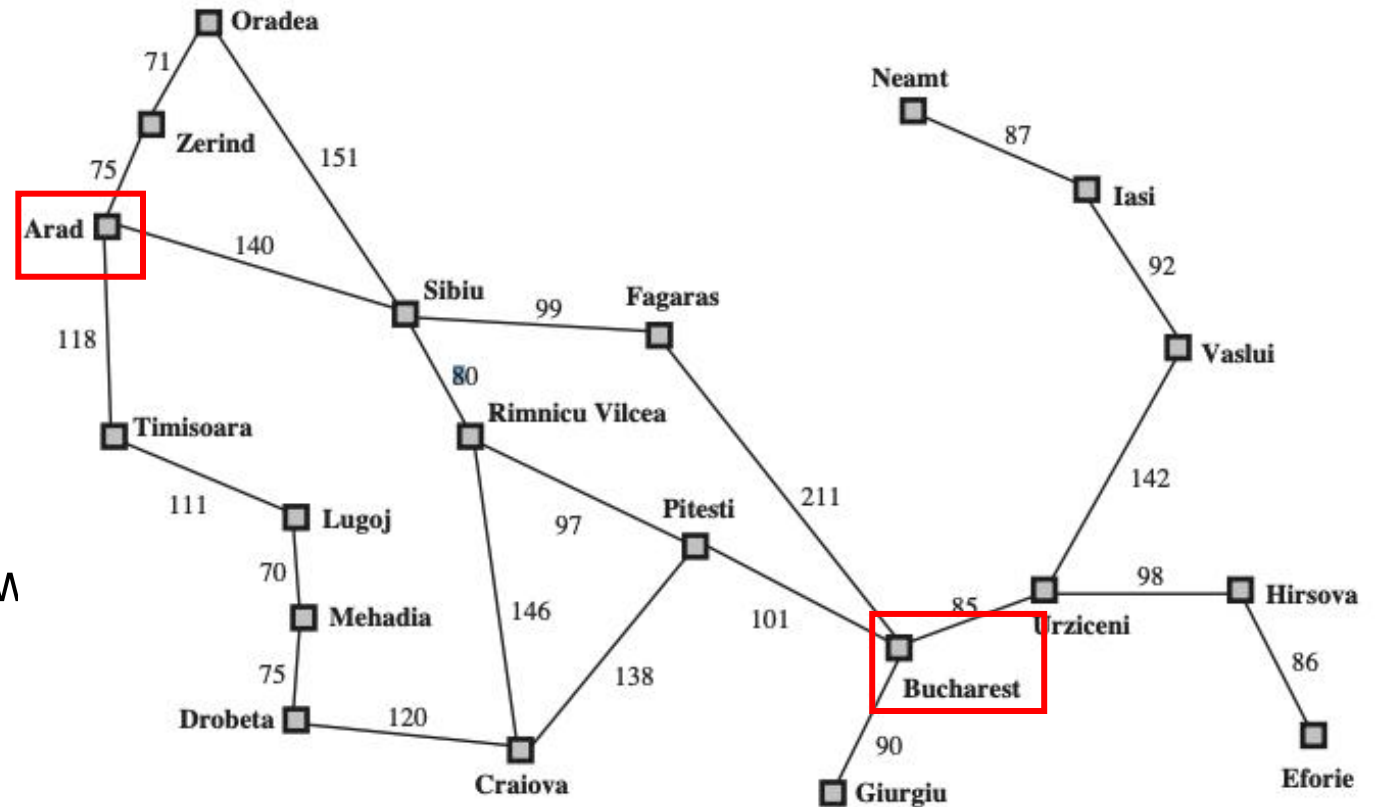
Example: Traveling in Romania

- State space: ?
- Initial state: ?
- Goal test: ?
- Successor function: ?
- Path cost: ?



Example: Traveling in Romania

- **State space:**
 - Cities
- **Initial state:**
 - Arad
- **Goal test:**
 - Is state == Bucharest
- **Successor function:**
 - Roads: go to adjacent cities w cost = distance
- **Path cost:**
 - The cost of a path



Solving a Search Problem: State Space Search

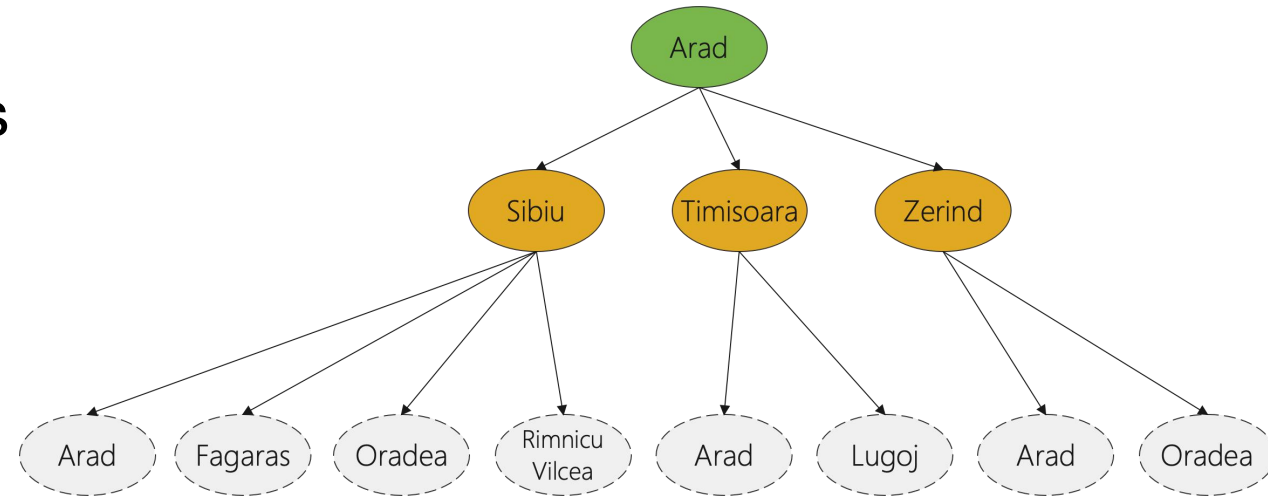
- Input:
 - Initial state
 - Goal test
 - Successor function
 - Path cost function
- Output: **path** from **initial state** to **goal**. Solution quality is measured by the **path cost function**
- Expanding: apply each legal action to the current state
- The leaf nodes available for expansion is called the frontier

Search procedure defines a search tree

root node — initial state

Children of a node — successor states

Leaves of tree (frontier) — states not yet expanded



Search strategy — algorithm for deciding which leaf node to expand next

Search procedure defines a search tree

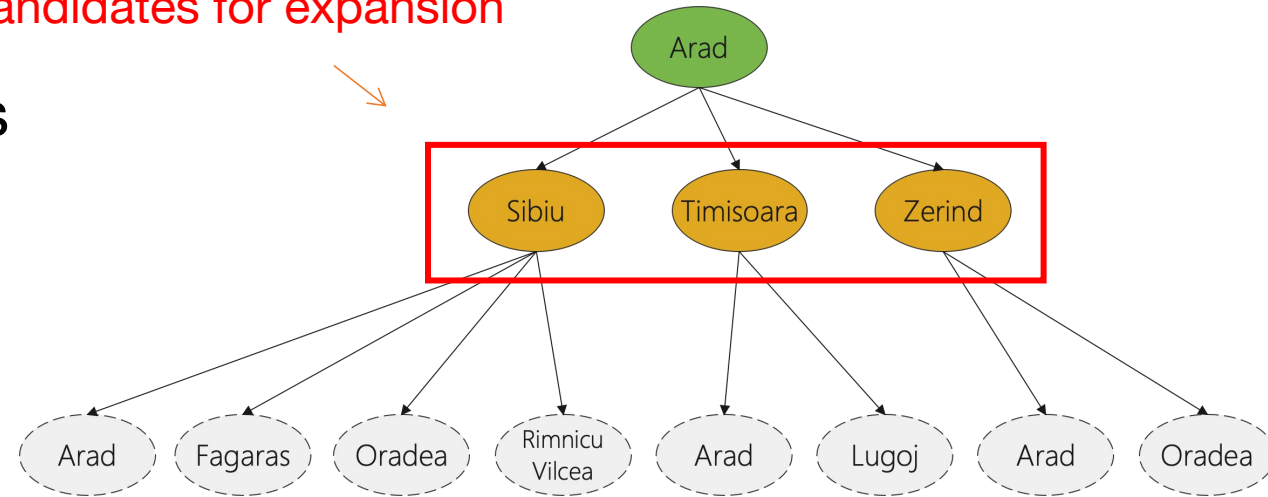
root node — initial state

Children of a node — successor states

Leaves of tree (frontier) — states not yet expanded


Search strategy — algorithm for deciding which leaf node to expand next

Candidates for expansion




Node Data Structure

- Node data structure is used to **keep track of** the search tree



Node Data Structure

- Node data structure is used to **keep track of** the search tree



- Nodes vs. States
 - A node is a **bookkeeping data structure** used to represent the search tree
 - A state corresponds to a **configuration** of the world.

Evaluating a Search Strategy

- Completeness: is the strategy guaranteed to find a solution when there is one?
- Time complexity: how long does it take to find a solution?
- Space complexity: how much memory does it need?
- Optimality: does the strategy find the highest-quality solution when there are several different solutions?

Generic Tree-Search Algorithm

Add **initial state** to the **frontier**

Loop

node = **remove-frontier**() -- and save in order to return as part of
path to **goal**

 if **goal-test**(**node**) = **true** return path to **goal**

S = **successors**(**node**)

 Add **S** to **frontier**

until **frontier** is empty

return **failure**

Generic Tree-Search Algorithm

Add **initial state** to the **frontier**

Loop

node = **remove-frontier**() -- and save in order to return as part of path to **goal**

if **goal-test**(**node**) = **true** return path to **goal**

S = **successors**(**node**)


Add **S** to **frontier**

until **frontier** is empty

return **failure**

Tree Search vs. Graph Search


- Tree search allows a state to be expanded more than once



- Failure to detect repeated states can cause more work
- Advantage:

Tree Search vs. Graph Search

- Tree search allows a state to be expanded more than once



- Failure to detect repeated states can cause more work
- Advantage: memory-efficient

Graph Search


- Idea: never expand a state twice

Graph Search

- Idea: never expand a state twice
- How to implement:
 - Tree search + set of expanded states (“closed set”)
 - Expand the search tree node-by-node, but...
 - Before expanding a node, check to make sure its state has never been expanded before
 - If not new, skip it, if new add to closed set

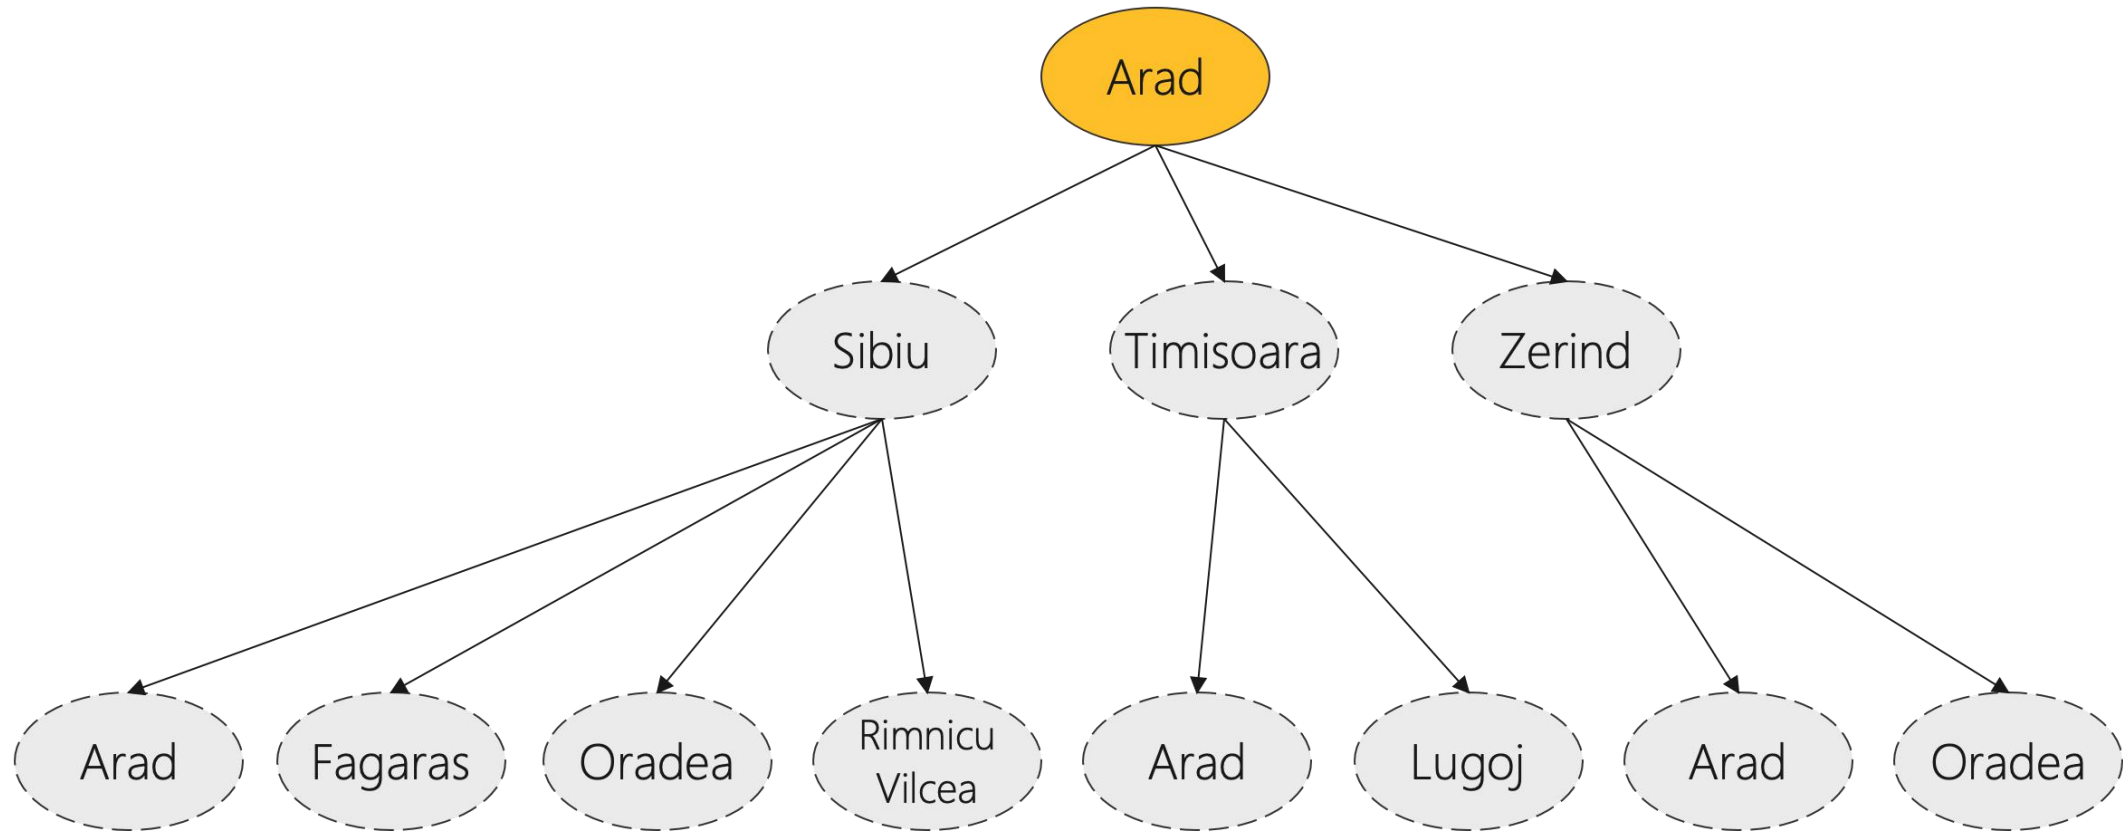
Uninformed Search: BFS

- Use the first-in, first out or FIFO queue to store the frontier

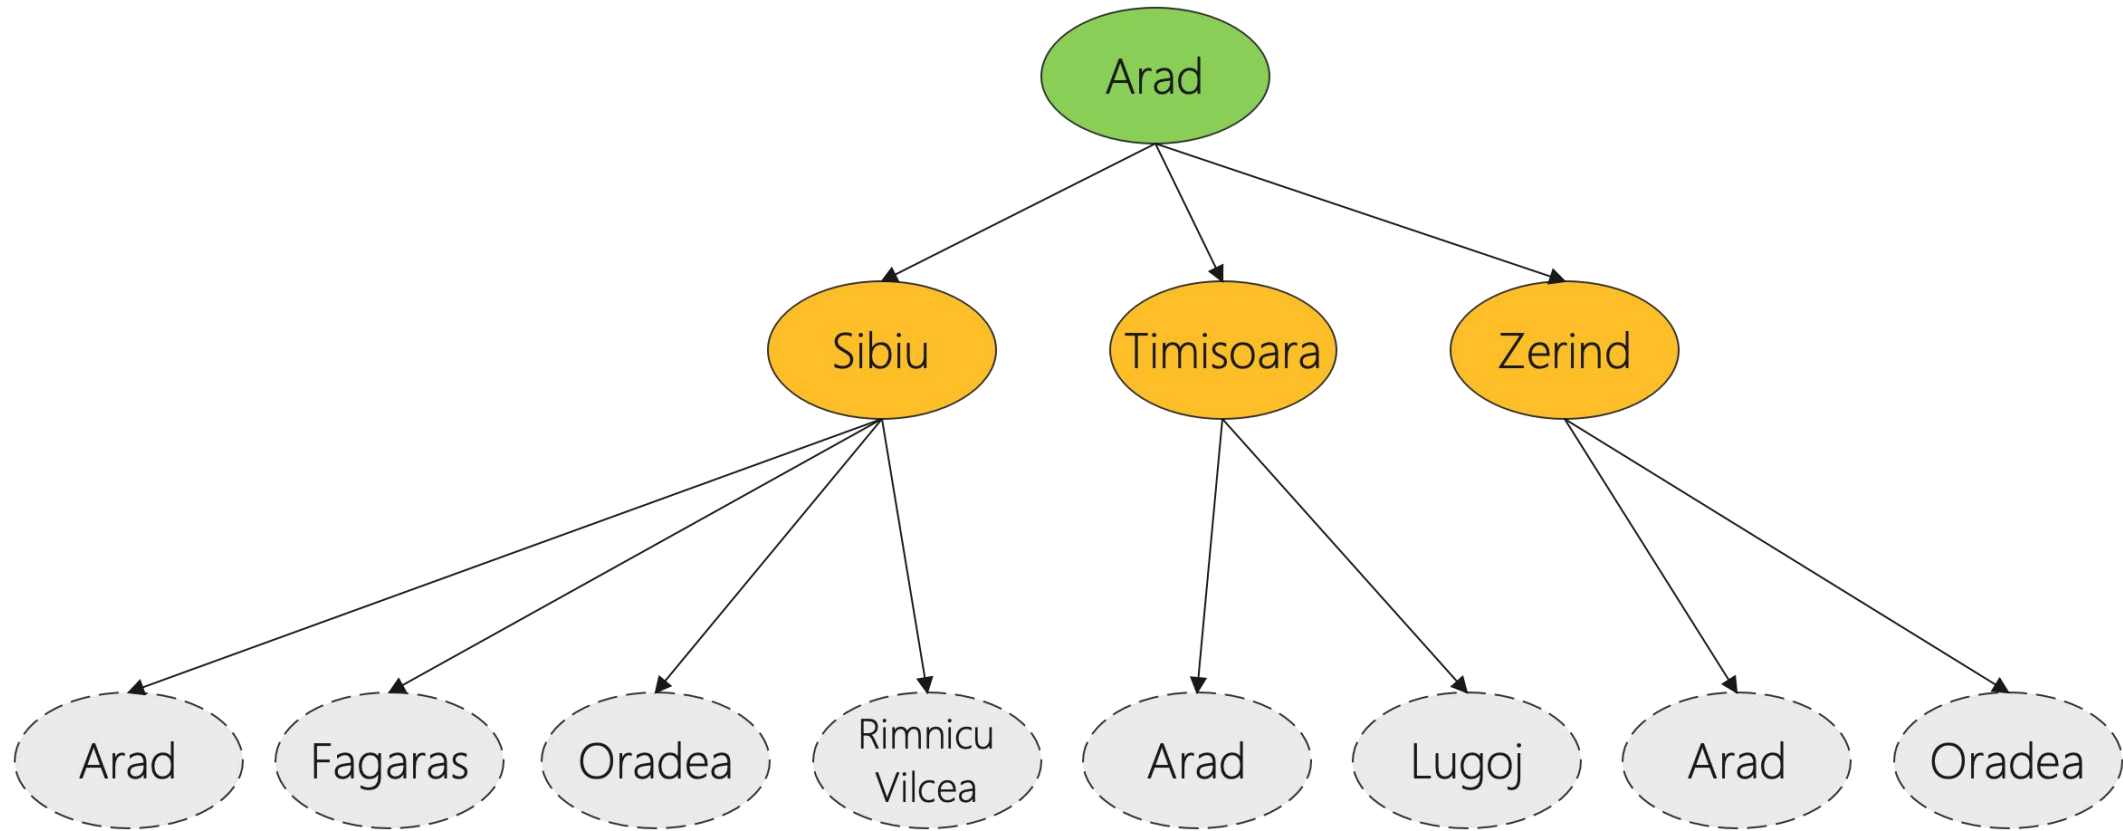


- Consider paths of length 1, then of length 2, then of length 3, then of length 4,

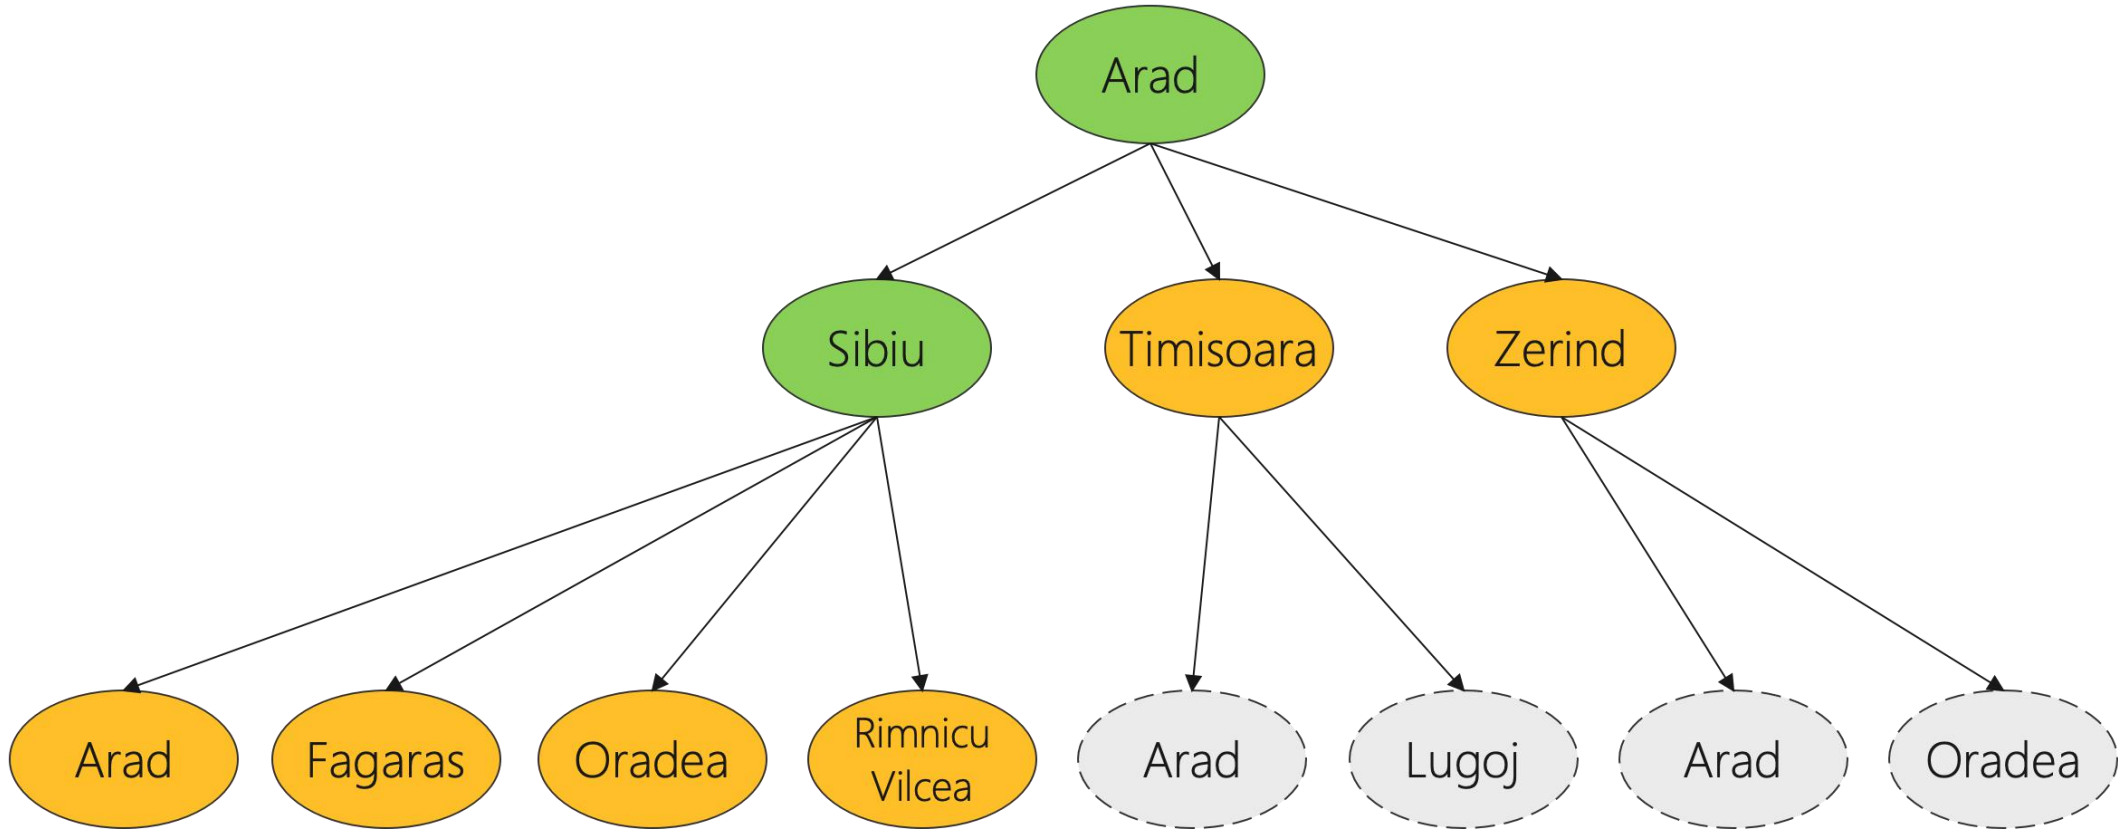
BFS: Example



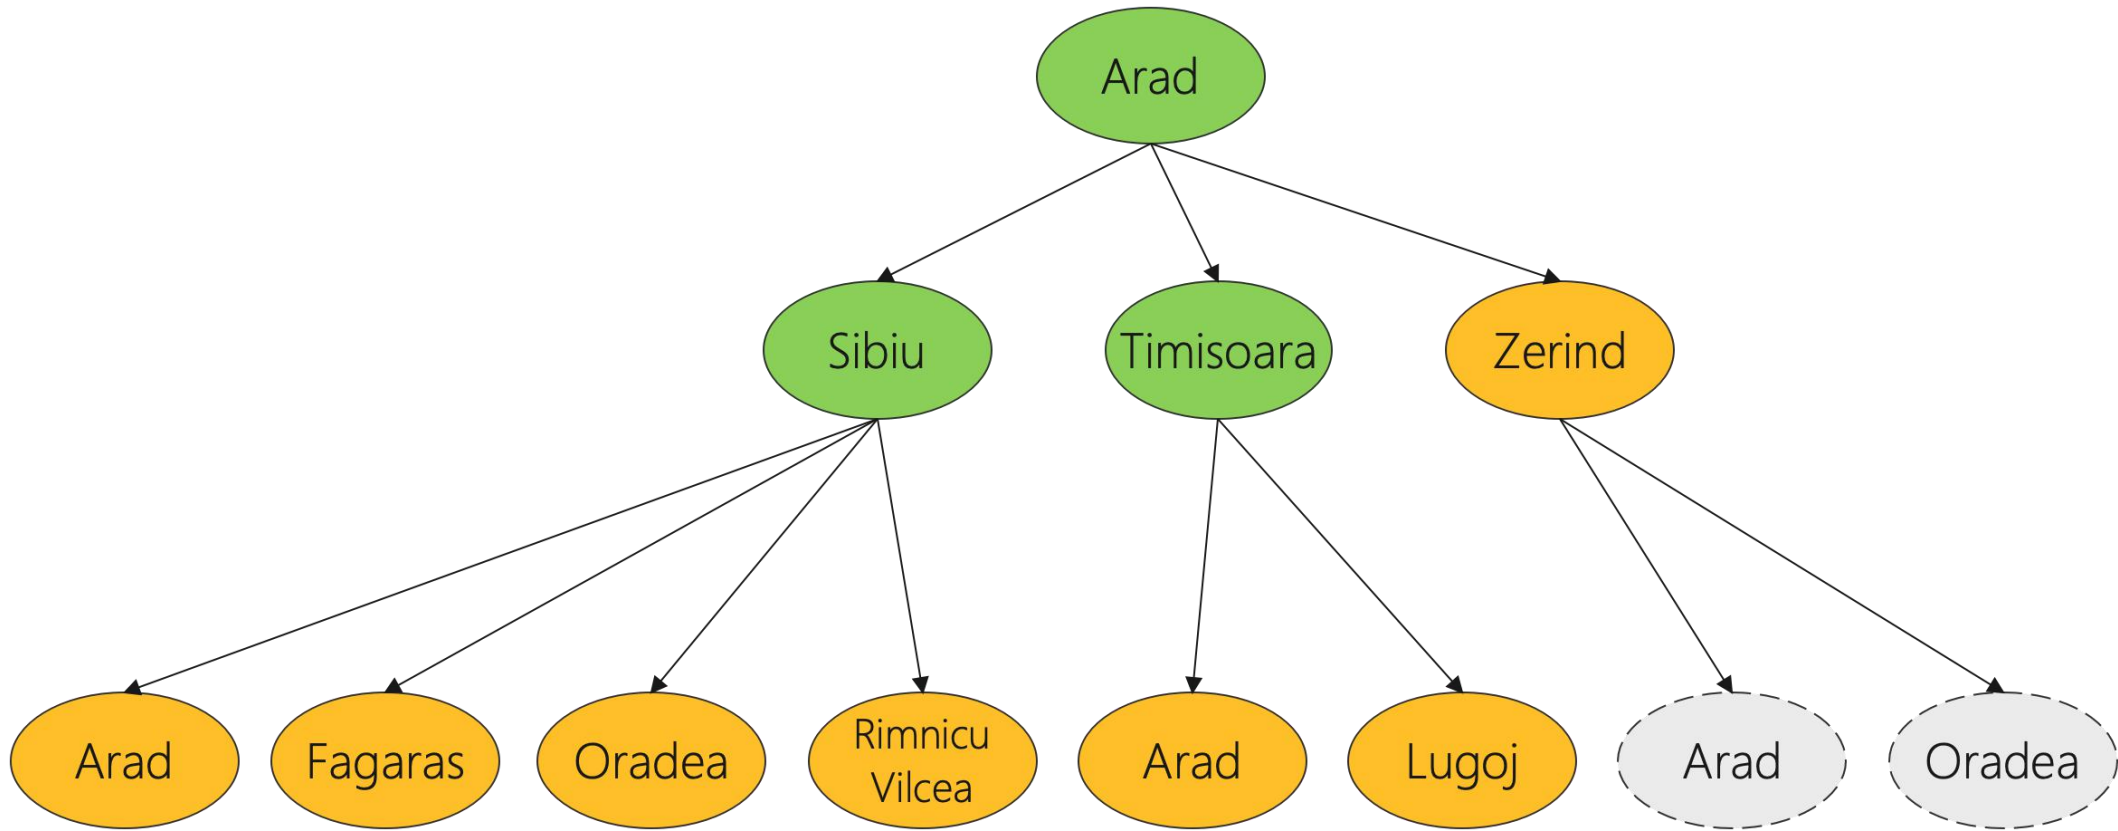
BFS: Example




BFS: Example



BFS: Example




BFS: Example



Time and Memory Requirements for BFS

- Let b = **branching factor** -> **maximum number of successors of any node**
- d = **solution depth** -> **the shallowest goal node**
- Then the maximum number of nodes generated is:
$$b + b^2 + \dots + b^d = O(b^d)$$
- For graph search,
 $O(b^{d-1})$ in the closed set and $O(b^d)$ in the frontier



Time and Memory Requirements for BFS

- **branching factor** = 10
- 1 million nodes / second
- each node requires 1000 bytes of storage

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

Time and Memory Requirements for BFS


- branching factor = 10
- 1 million nodes / second
- each node requires 1000 bytes of storage

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

125000 * 8GB

Uniform-Cost Search


- Use BFS, but always expand the **lowest-cost node** on the frontier as measured by path cost $g(n)$



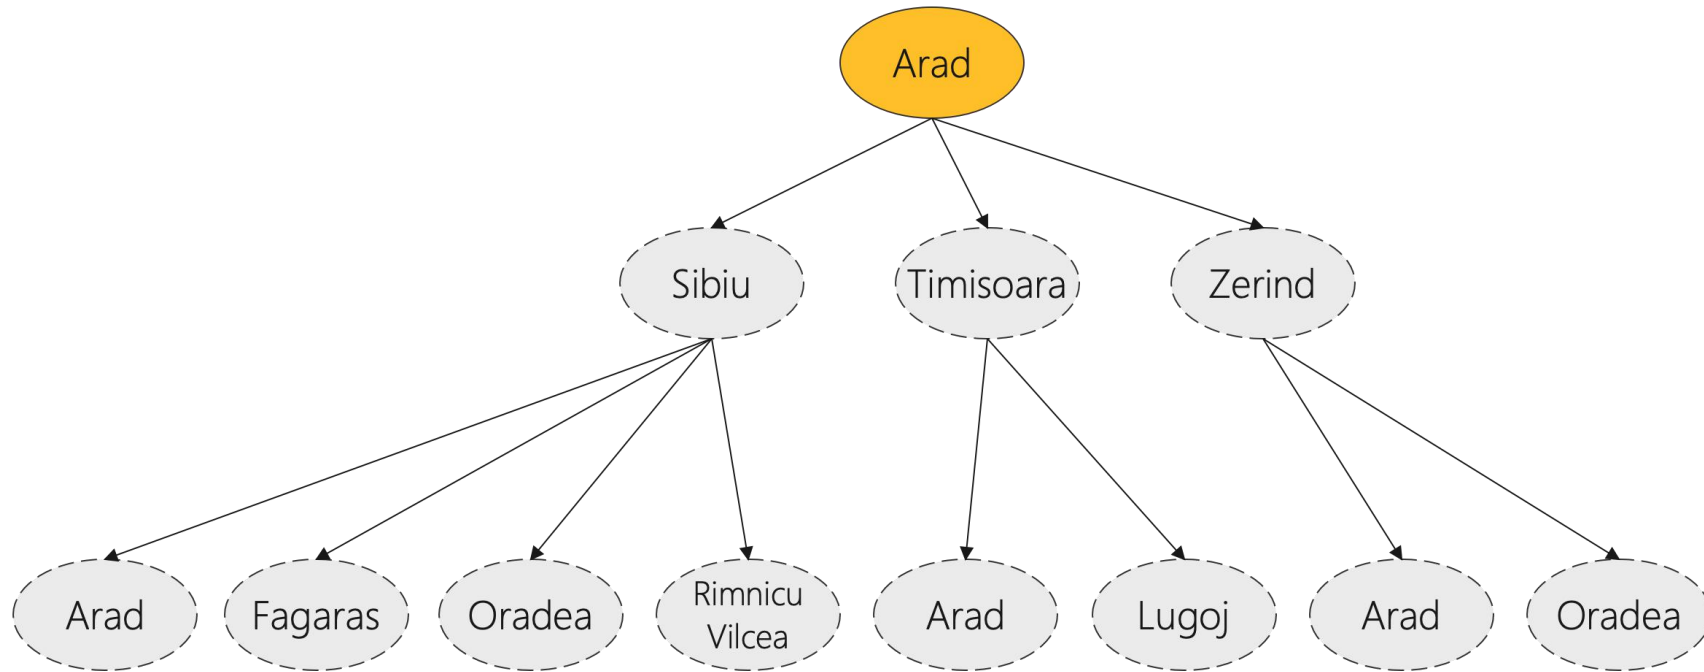
- $g(\text{Successor}(n)) > g(n)$ is a necessary condition for completeness and a sufficient condition for optimality

Uninformed search: DFS

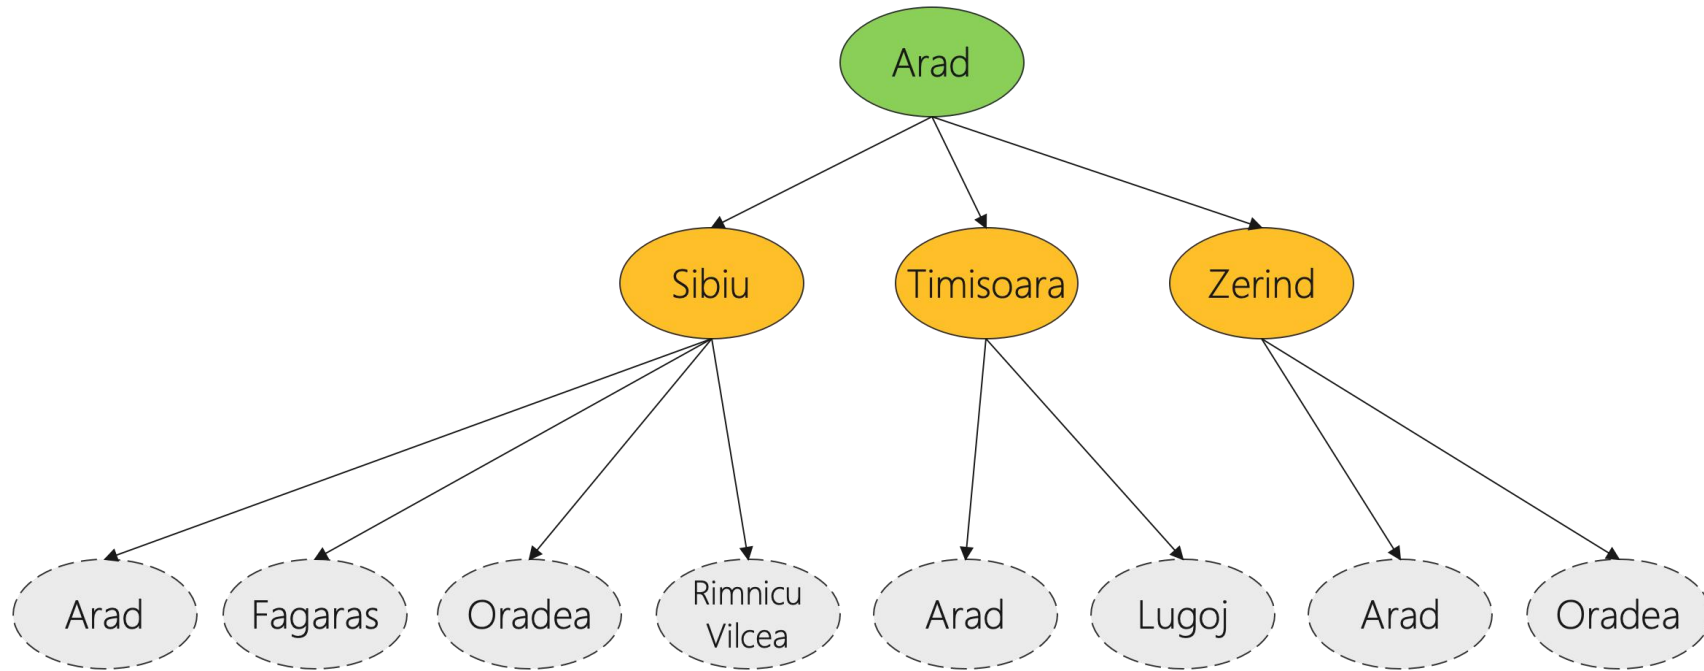
- Use the last-in, first out or LIFO queue to store the frontier



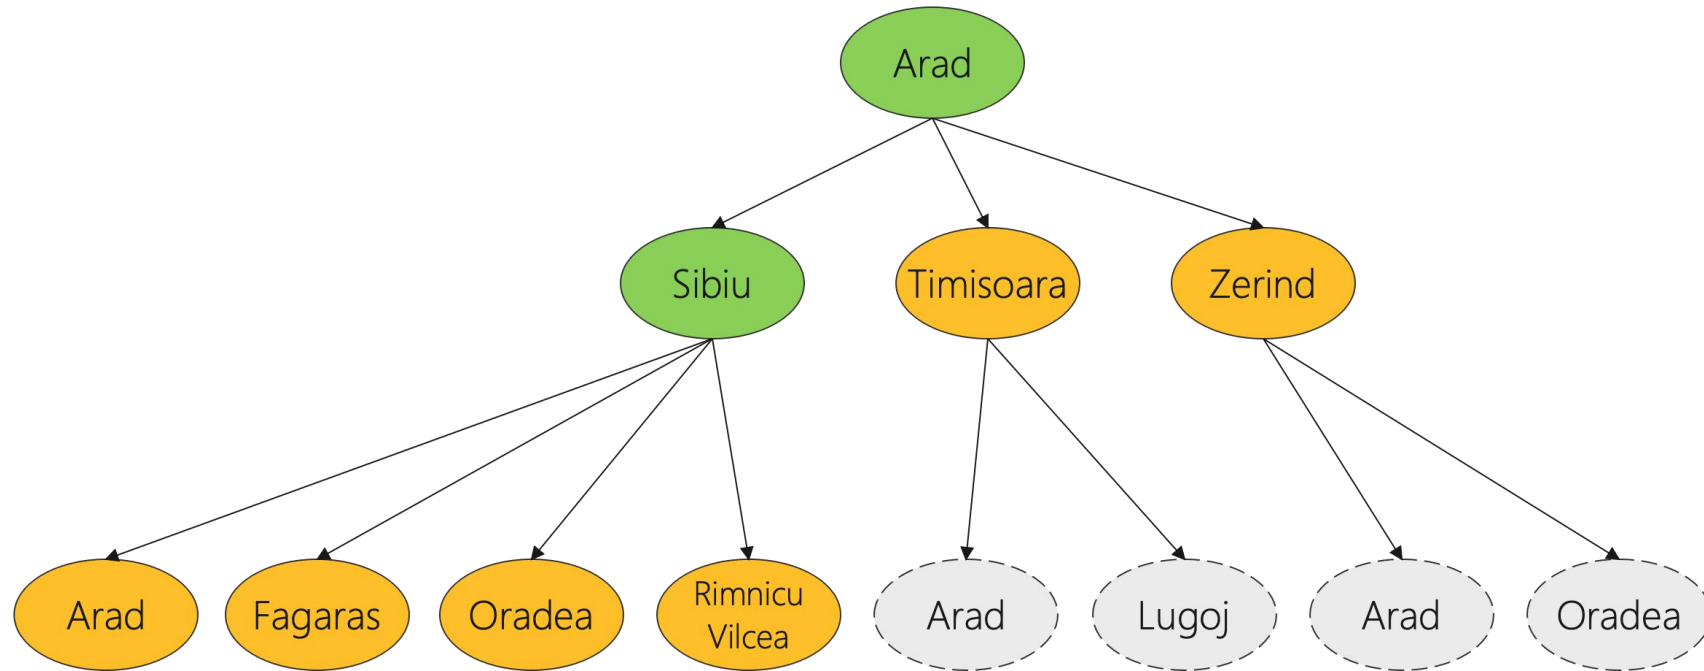
DFS: Example



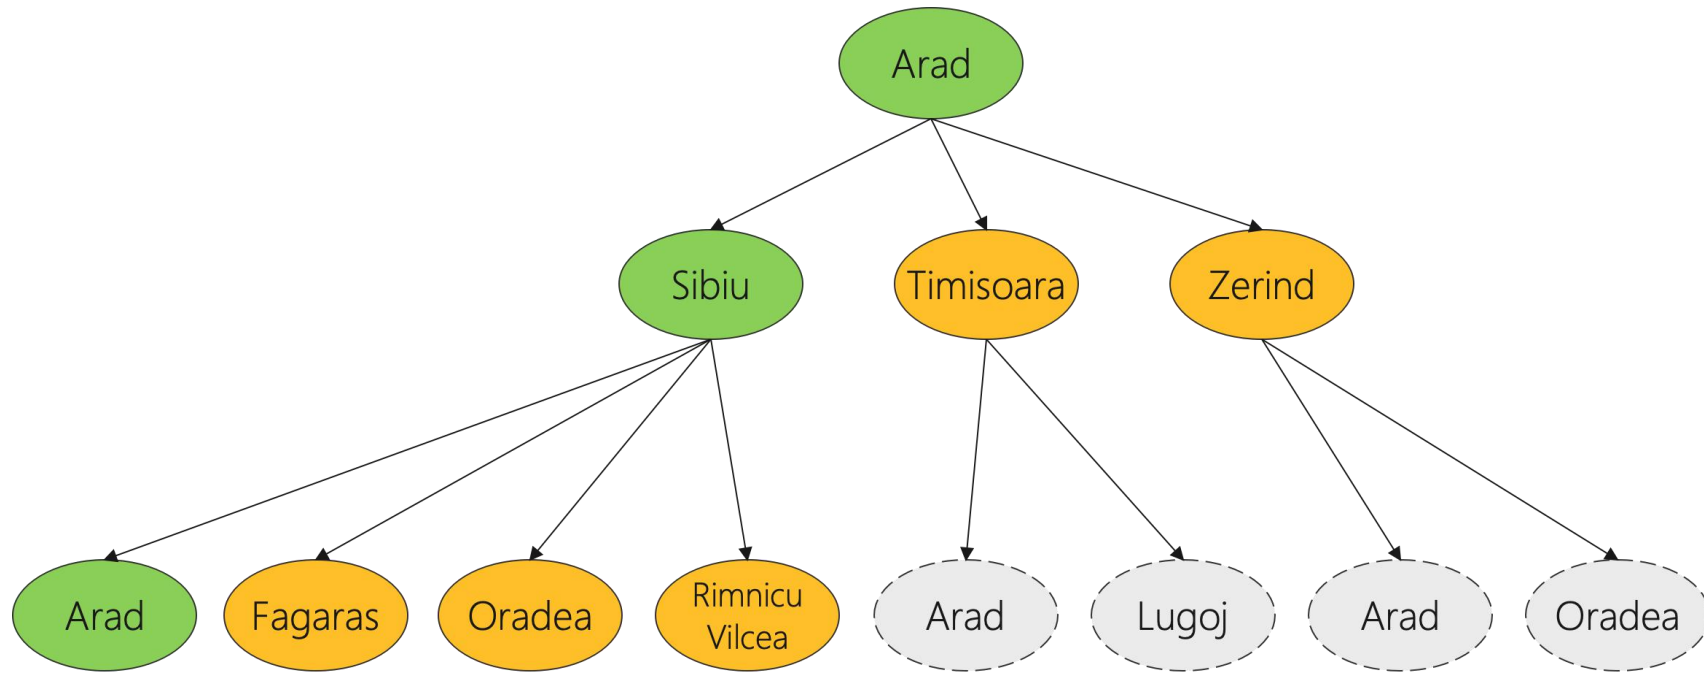
DFS: Example




DFS: Example



DFS: Example



DFS: Example



Time and Memory Requirements for DFS


- Let b = **branching factor** \rightarrow **maximum number of successors of any node**
- m = **maximum depth of any node**

Time and Memory Requirements for DFS

- Let b = **branching factor** -> **maximum number of successors of any node**
- m = **maximum depth of any node**
- Time: for tree search, then the maximum number of nodes generated is: $O(b^m)$
- Space: for tree search, only need to store **?** nodes

Time and Memory Requirements for DFS

- Let b = **branching factor** -> **maximum number of successors of any node**
- m = **maximum depth of any node**
- Time: for tree search, then the maximum number of nodes generated is: $O(b^m)$
- Space: for tree search, only need to store **$O(bm)$** nodes
 - at depth $l < d$ we have $b-1$ nodes
 - at depth d we have b nodes
 - total = $(m-1)*(b-1) + b = O(bm)$



DFS vs. BFS

	Complete?	Optimal?	Time	Space
BFS	YES	YES	$O(b^d)$	$O(b^d)$
DFS	finte depth	NO	$O(b^m)$	$O(bm)$

d: depth of the shallowest solution m: maximum depth

- Takeaways:
 - If the solution is not far from the root: BFS might be faster
 - If the search tree is very deep: DFS may never find solution
 - If the search tree is wide: BFS might take too much memory
 - If d is known: DFS is preferred
 - If optimal solution is needed: BFS is preferred

Iterative Deepening [Korf 1985]

- Problem of DFS:

Iterative Deepening [Korf 1985]

- Problem of DFS: cannot avoid infinite loops

Iterative Deepening [Korf 1985]

- Problem of DFS: cannot avoid infinite loops

- Idea:

Use an artificial depth cutoff, c .

If search to depth c succeeds, we are done. if not, increase c by 1 and start over.

Each iteration searches using DFS.

Iterative Deepening [Korf 1985]

- Problem of DFS: cannot avoid infinite loops

- Idea:

Use an artificial depth cutoff, c .

If search to depth c succeeds, we are done. if not, increase c by 1 and start over.

Each iteration searches using DFS. Why?

Iterative Deepening [Korf 1985]

- Problem of DFS: cannot avoid infinite loops

- Idea:

Use an artificial depth cutoff, c .

If search to depth c succeeds, we are done. if not, increase c by 1 and start over.

Each iteration searches using DFS.

Combine the benefit of the DFS and BFS

Iterative Deepening

Limit = 0 ●

Iterative Deepening

Limit = 0



Limit = 1



Iterative Deepening





Limit = 0




Limit = 1



Limit = 2



Iterative Deepening



Iterative Deepening

- Space requirements: same as DFS.

Iterative Deepening

- Space requirements: same as DFS.
- Time requirements: would seem very expensive! But not much different from single BFS or DFS to depth d

Iterative Deepening

- Space requirements: same as DFS.
- Time requirements: would seem very expensive! But not much different from single BFS or DFS to depth d
- Reason: **Almost all work is in the final couple of layers.** E.g., binary tree: $1/2$ of the nodes are in the bottom layer. With $b = 10$, $9/10^{\text{th}}$ of the nodes in the final layer!
- So, repeated runs are on much smaller trees (i.e., exponentially smaller).

Iterative Deepening

Examples:

$b = 10$, $d = 5$, the number of nodes generated in a BFS:

$$b + b^2 + \dots + b^d = 10 + 100 + 1,000 + 10,000 + 100,000 = 111,110$$

For IDS:

$$(d)b + (d-1)b^2 + \dots + (1)b^d =$$

$$50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$


Cost of repeating the work at shallow depths is not prohibitive.

Cost of Iterative Deepening


- Space: $O(bd)$ as in DFS, time: $O(b^d)$
- Asymptotic ratio of the number of nodes generated by BFS and IDS:
 - $(b+1)/(b-1)$

b	ratio of IDS to DFS
2	3
3	2
5	1.5
10	1.2
25	1.08
100	1.02

Bidirectional Search




Bidirectional Search



- When is bidirectional search applicable?
 - Generate predecessors is easy

Bidirectional Search



- When is bidirectional search applicable?
 - Generate predecessors is easy
 - Goal state is clearly specified. (“no queen attacks another queen”)

Bidirectional Search

- Search forward from the start state and backward from the goal state simultaneously and stop when the two searches meet the middle

Bidirectional Search

- Search forward from the start state and backward from the goal state simultaneously and stop when the two searches meet the middle
- If branching factor = b from both directions, and solution exists at depth d , then need only $O(2b^{d/2}) = O(b^{d/2})$ steps.

Bidirectional Search

- Search forward from the start state and backward from the goal state simultaneously and stop when the two searches meet the middle
- If branching factor = b from both directions, and solution exists at depth d , then need only $O(2b^{d/2}) = O(b^{d/2})$ steps.
- Example: $b = 10$, $d = 6$ then BFS needs 1,111,110 nodes and bidirectional search needs only 2,220.

Limitations

- What are the problems of all the methods?

Limitations

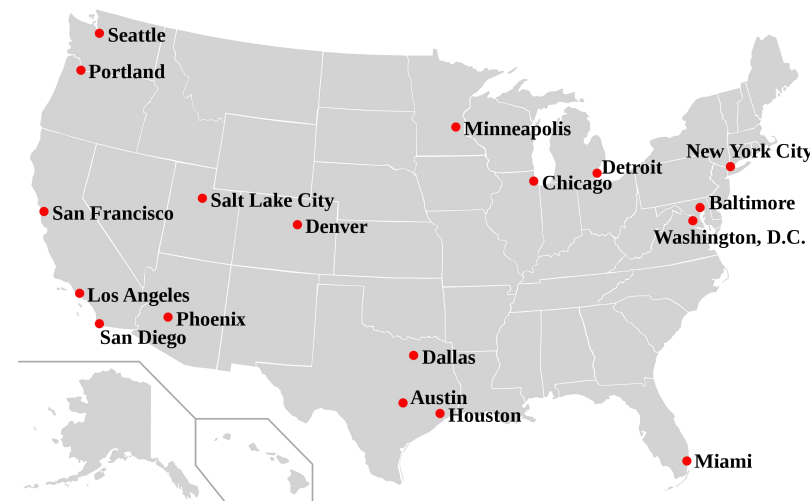
- What are the problems of all the methods? **Slow!**

Limitations

- What are the problems of all the methods? **Slow!**
- The search is blind in the sense that the information of the goal state is not used

Limitations

- What are the problems of all the methods? **Slow!**
- The search is blind in the sense that the information of the goal state is not used
- Informed search:
 - with the guidance of the goal state



Goals of This Lecture

- Understand state-space search
- Understand BFS, DFS, UCS and IDS
- Know the advantages and disadvantages of each search strategy

Reading: Sections 3.1 - 3.4, R&N.